

**INFO-H419 - Data Warehouse**

---

# Project 1 TPC-DS Benchmark Using DuckDB

---

**Group:**

Jintao Ma

**Professor**

Linhan Wang

Esteban Zimányi

Iyoha Peace Osamuyi

Hieu Nguyen Minh

November 2023

# Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
1.1	Overview . . . . .	1
1.2	Aim and Objectives . . . . .	1
1.3	Limitations . . . . .	2
<b>2</b>	<b>Data Warehouse and TPC-DS Benchmarking</b>	<b>3</b>
2.1	Data Warehouses . . . . .	3
2.2	TPC-DS Benchmark . . . . .	4
2.3	DuckDB . . . . .	5
2.3.1	Simple and portable . . . . .	5
2.3.2	Feature-rich . . . . .	6
2.3.3	Fast . . . . .	6
2.3.4	Free and Extensible . . . . .	6
2.4	Other Tools Used . . . . .	7

2.4.1	Python . . . . .	7
2.4.2	Code Editor: Microsoft Visual Code . . . . .	7
2.4.3	DBeaver . . . . .	7
2.4.4	Apache Jmeter . . . . .	8
<b>3</b>	<b>Benchmark Implementation Process</b>	<b>9</b>
3.1	Setting up the Project . . . . .	9
3.1.1	Hardware Specification . . . . .	9
3.2	Implementing TPC-DS on DuckDB . . . . .	10
3.2.1	Scaling and Database Population . . . . .	12
3.3	Implementation Process . . . . .	12
3.4	Setup . . . . .	13
3.4.1	DuckDB installation . . . . .	13
3.4.2	TPC-DS kit installation . . . . .	13
3.4.3	Apache Jmeter Installation . . . . .	13
3.5	Database schema creation . . . . .	14
3.6	Data generation . . . . .	15
3.7	Query generation . . . . .	15
3.7.1	Separating queries from the templates . . . . .	16
3.7.2	Query Modification . . . . .	16
3.7.3	Running Queries . . . . .	16

3.8 Load Test . . . . .	17
3.9 Power Test . . . . .	18
3.10 Throughput Test . . . . .	19
3.11 Data Maintenance Test . . . . .	22
3.12 Data Accessibility Test . . . . .	26
3.13 Scalability Test . . . . .	26
<b>4 Results and Discussion</b>	<b>30</b>
4.1 Load Test . . . . .	30
4.2 Power Test . . . . .	31
4.2.1 Load and Power Test Comparison . . . . .	32
4.2.2 Individual Scale Execution . . . . .	33
4.2.3 Execution Times in Different Scale Factors . . . . .	45
4.3 Throughput Test . . . . .	47
4.3.1 Comparison for Throughput, Power and Load Tests . . . . .	49
4.4 Data Maintenance Test . . . . .	50
4.5 Data Accessibility Test . . . . .	53
4.6 Scalability Test . . . . .	53
4.7 Optimization . . . . .	55
<b>5 Conclusion</b>	<b>72</b>
5.1 Conclusion . . . . .	72

## **Appendices** 75

A Appendix A . . . . .	75
A.1 Optimized Queries . . . . .	75

## Listings

3.1 DuckDB installation.	13
3.2 Database setup	14
3.3 Database setup	14
3.4 Data generation.	15
3.5 Generated data example.	15
3.6 Data generation.	15
3.7 Load Test Pseudocode	17
3.8 Python script for load test.	17
3.9 Power test pseudocode.	18
3.10 Power test Python script.	19
3.11 Throughput test pseudocode	20
3.12 Throughput test Python script.	20
3.13 Data maintenance test pseudocode.	23
3.14 Refresh Data generation.	23

3.15 Python script for maintenance test. . . . .	24
4.1 Query 95. . . . .	56
4.2 Optimized Query 95. . . . .	57
4.3 Query 67. . . . .	59
4.4 Optimized Query 67. . . . .	60
4.5 Query 23b. . . . .	62
4.6 Optimized Query 23b. . . . .	63
4.7 Query 4. . . . .	67
4.8 Optimized Query 14. . . . .	70
1 Optimized Query 11. . . . .	75
2 Optimized Query 14. . . . .	78
3 Optimized Query 14b. . . . .	81
4 Optimized Query 22. . . . .	84
5 Optimized Query 23. . . . .	85
6 Optimized Query 47. . . . .	88
7 Optimized Query 57. . . . .	89
8 Optimized Query 64. . . . .	91

## List of Figures

3.1 TPC-DS Benchmarking BPMN Diagram . . . . .	11
4.1 Load test results. . . . .	31
4.2 Power test results. . . . .	31
4.3 Load test and Power Test comparison(a). . . . .	32
4.4 Load test and Power Test comparison(b). . . . .	33
4.5 Power test result for scale factor 1. . . . .	34
4.6 Deleted Queries for scale factor 1. . . . .	35
4.7 After Deleting Queries for scale factor 1. . . . .	35
4.8 Power test result for scale factor 3. . . . .	37
4.9 Deleted Queries for scale factor 3. . . . .	38
4.10 After Deleting Queries for scale factor 3. . . . .	38
4.11 Power test result for scale factor 5. . . . .	40
4.12 Deleted Queries for scale factor 5. . . . .	41

4.13 After Deleting Queries for scale factor 5 . . . . .	41
4.14 Power test result for scale factor 7 . . . . .	43
4.15 Deleted Queries for scale factor 7 . . . . .	44
4.16 After Deleting Queries for scale factor 7 . . . . .	44
4.17 Comparison of power test results for the first 10 queries over 4 scale factors. . . . .	45
4.18 Comparison of power test results for the queries 10 to 20 over 4 scale factors. . . . .	46
4.19 Query 8 . . . . .	47
4.20 Query 19 . . . . .	48
4.21 Throughput 1 and 2. . . . .	48
4.22 Throughput test results. . . . .	49
4.23 Comparison for Throughput, Power and Load Tests Time for the 4 Scale Factors . . .	50
4.24 Data Maintenance test results. . . . .	51
4.25 Maintenance test results. . . . .	52
4.26 Maintenance test . . . . .	52
4.27 Data accessibility test results. . . . .	53
4.28 Time to perform the queries . . . . .	53
4.29 Time to perform the queries . . . . .	54
4.30 Time to perform the queries . . . . .	54
4.31 Statistics of Queries . . . . .	54
4.32 Statistics of Queries . . . . .	54

4.33 Statistics of Queries . . . . .	55
4.34 Optimization/Comparison of Origin and Optimization of query 95 . . . . .	58
4.35 Optimization/Comparison of Origin and Optimization of query 67 . . . . .	61
4.36 Optimization/Comparison of Origin and Optimization of query 23b . . . . .	66
4.37 Optimization/Comparison of Origin and Optimization of query 4 . . . . .	71

## **Abstract**

In this project we have implemented the TPC-DS Benchmark on DuckDB. It is a free and open-source database created by CWI.

In Chapter 1, we presented the general objectives and mission requirements of the project. In Chapter 2, we reviewed the basics of data warehousing and TPC-DS benchmarking and introduced the programming tools and databases we used. In Chapter 3, we described concretely the whole process of how we made the different tests executed in the benchmarking. This contains assumptions and code implementations. In Chapter 4, we used a lot of charts to compare, analyze, interpret, and completed some optimizations. In Chapter 5, we summarized our conclusions.

All experimental and analytical codes are in this github repository:

[https://github.com/woshimajintao/TPC\\_DS\\_for\\_DuckDB.git](https://github.com/woshimajintao/TPC_DS_for_DuckDB.git)

# 1

## Introduction

### 1.1 Overview

The Project entails performing TPC-DS Benchmarking on a selected Database Management System which include DuckDB, PostgreSQL, SparkSQL, MySQL etc. The Benchmarking is done with several scale factors which determine the size of the resulting data warehouse and then implement some metrics to evaluate the performance of the database. The project is carried out in groups of 4 persons and delivers a report in pdf file explaining the essential aspects of your implementation, and a zip file containing the code of your implementation, with all necessary instructions to be able to replicate your implementation. This report employs DuckDB as the preferred choice of Database Management System on which the TPC-DS benchmark is performed.

### 1.2 Aim and Objectives

The aim of this report is to implement and evaluate the TPC-DS benchmark on DuckDB, with the following objectives:

- Learning how to properly perform a benchmark efficiently.
- Evaluating the benchmarking performance and analyze the results.
- Optimizing queries that are complex and time consuming.
- Visualizing the result.

- Testing the performance with Apache Jmeter.

### 1.3 Limitations

The tool used has a limit on the resources and memory size, thereby preventing scaling higher benchmarks. Another Limitation was that the SQL Keyword "**ALTER**" hasn't been fully implemented to handle referential integrity constraints. Lastly, time was a limiting factor for us to thoroughly expose the objectives stated above.

## Data Warehouse and TPC-DS Benchmarking

### 2.1 Data Warehouses

In today's data-driven world, organizations are facing challenges of handling enormous information from various sources, including transactions, customers' behavior, web activities, etc. To derive meaningful insights from collected data, people in those organizations resort to decision support tools for analyzing the data. The whole process of data processing involves a lot of steps, from data extraction, data transformation, to data cleaning and data storage. In this regard, data are stored in a common repository where users can access and get immediate responses to complex queries. This repository is called a *data warehouse* [VZ14].

According to Vaisman and Zimányi [VZ14], a data warehouse can be defined as a collection of subject-oriented, integrated, nonvolatile, and time-varying data to support management decisions. These concepts can be broken down as:

- Data is subject-oriented when the data warehouse targets one or several subjects of analysis according to the analytical requirements of the organization.
- It is integrated when data from different sources (operational and external) are integrated together.
- It is nonvolatile in the sense that modification and removal of data is not allowed. This is because the data warehouse accumulates data for a long period of time, which cannot be changed.
- Finally, time-varying, because we are dealing with how the data has evolved over time.

A data warehouse is a central repository designed to store, organize, and manage data from different sources, making it available for data analysis and reporting. Unlike traditional databases, data warehouse are designed for efficient retrieval and processing of large data sets. To be more specific, data tuples (i.e. rows in tables) are used as read-only operations, such as querying and reporting, rather than modifying the data.

## 2.2 TPC-DS Benchmark

Benchmarking is a process of comparing and evaluating an organization's performance against an external standard or best practice model. It involves measuring and analyzing various aspects of an organization's processes, performance, and practices against those of other similar organizations, with the goal of identifying areas for improvement and setting goals for achieving better results [Mah21]. Benchmarking a database is the process of performing well defined tests on that particular database for the purpose of evaluating its performance [Sam21]. Benchmark test results facilitate means for cross platform comparisons of various database management systems by providing valuable information to database professionals on whether to utilise a particular database product. Hence, we are carrying and evaluating the performance of DuckDB against the TPC-DS Benchmarking standard.

**Transaction Processing Performance Council (TPC)** is a non-profit organization founded to define transaction processing and database benchmarks and to disseminate objective, verifiable TPC performance data to the industry<sup>1</sup>. In terms of data warehouses, TPC provide two benchmarks: TPC-DS and TPC-DI for the evaluation of decision support and data integration, respectively.

The TPC-DS is a decision support benchmark that models several generally applicable aspects of a decision support system, including queries and data maintenance. The benchmark provides a representative evaluation of the System Under Test's (SUT) performance as a general purpose decision support system [TPC21].

This TPC-DS benchmark depicts a decision support systems that:

- Examine large volumes of data;

---

<sup>1</sup>[https://en.wikipedia.org/wiki/Transaction\\_Processing\\_Performance\\_Council](https://en.wikipedia.org/wiki/Transaction_Processing_Performance_Council)

- Give answers to real-world business questions;
- Execute queries of various operational requirements and complexities (e.g., ad-hoc, reporting, iterative OLAP, data mining);
- Are characterized by high CPU and IO load;
- Are periodically synchronized with source OLTP databases through database maintenance functions.
- Run on “Big Data” solutions.

In addition, a benchmark result measures query response time in single user mode, query throughput in multi user mode and data maintenance performance for a given hardware, operating system, and data processing system configuration under a controlled, complex, multi-user decision support workload [TPC21].

## 2.3 DuckDB

DuckDB is an in-process SQL OLAP database management system.<sup>2</sup> The first open-source version of it was released by Hannes Mühleisen and Mark Raasveldt, who were researchers in the Database Architectures research group at Centrum Wiskunde & Informatica (CWI) in 2019.[RM19] It was the first purpose built in-process Online Analytical Processing(OLAP)-database management system.<sup>3</sup> It was designed to be fast, reliable and easy to use. There have been some important features of this tool:

### 2.3.1 Simple and portable

For DuckDB, there is no DBMS server software to install, update and maintain. It does not run as a separate process, but completely embedded within a host process. For the analytical use cases that DuckDB targets, this has the additional advantage of high-speed data transfer to and from the database. In some cases, DuckDB can process foreign data without copying. We can use

---

<sup>2</sup><https://duckdb.org/>

<sup>3</sup><https://www.cwi.nl/en/news/cwi-spin-off-company-duckdb-labs-provides-solutions-for-fast-database-analytics>

DuckDB without installing, updating, or maintaining separate DBMS server software. It doesn't run as an independent process but is entirely integrated within a host process. This approach provides added benefits for analytical use cases, notably enabling high-speed data transfer to and from the database.

### 2.3.2 Feature-rich

DuckDB enables users to run complex SQL queries and provides APIs for Java, C, C++, and more. It is also deeply integrated into Python and R, enabling users to conduct efficient interactive data analysis; thus, you can interact with DuckDB from your preferred programming language. There is also access to extra SQL keywords that make SQL queries easier to write, such as EXCLUDE, REPLACE, and ALL.

### 2.3.3 Fast

DuckDB can support the Online analytical processing (OLAP).It features a query execution engine that employs columnar vectorization. While a few queries may still be interpreted, a substantial batch of values, known as a "vector," is processed in a single operation. This approach significantly minimizes the overhead found in conventional systems like PostgreSQL, MySQL, or SQLite, which handle each row sequentially. The utilization of vectorized query execution significantly enhances performance in OLAP queries. They used an interactive SQL shell to demonstrate the capabilities of an analytical database in the browser[Koh+22].

### 2.3.4 Free and Extensible

Because the DuckDB is Open Source, all of the source code is freely available on GitHub. You can access it by click the link: <https://github.com/duckdb> They also invite developers from anyone to take issues.

## 2.4 Other Tools Used

We use the following tools to implement the benchmark:

### 2.4.1 Python

Python is an interpreted, general-purpose, interactive, object-oriented programming language. It incorporates modules, exceptions, dynamic typing, very high level dynamic data types, and classes. It supports multiple programming paradigms beyond object-oriented programming, such as procedural and functional programming. Python combines remarkable power with very clear syntax. It has interfaces to many system calls and libraries, as well as to various window systems, and is extensible in C or C++ [Pyt23]. In this project, we used python to write scripts for creating the database, loading of tables in our database, querying and performing various evaluation metrics. Also, Jupyter notebook with matplotlib was used to visualize our data. Python version 3.10 was used in this project.

### 2.4.2 Code Editor: Microsoft Visual Code

Visual Studio Code is a lightweight but powerful source code editor which runs on your desktop and is available for Windows, macOS and Linux. It comes with built-in support for JavaScript, TypeScript and Node.js and has a rich ecosystem of extensions for other languages and runtimes (such as C++, CSharp, Java, Python, PHP, Go, .NET) [Cod23]. It was useful in this project to write our python scripts and running dsdgen and dsqgen on command line.

### 2.4.3 DBeaver

DBeaver is a free cross-platform universal database tool for developers, database administrators, analysts, and everyone working with data. It supports all popular SQL databases like DuckDB, MySQL, MariaDB, PostgreSQL, SQLite, Apache Family, and more. It was used to test our queries and validating the output [Com23]

#### 2.4.4 Apache Jmeter

The Apache JMeter is pure Java open source software, which was first developed by Stefano Mazzocchi of the Apache Software Foundation, designed to load test functional behavior and measure performance. You can use JMeter to analyze and measure the performance of web applications or a variety of services. Performance Testing means testing a web application against heavy load, multiple and concurrent user traffic. JMeter originally is used for testing Web Application or FTP application. Nowadays, it is used for a functional test, database server test etc. In this Project, it was used to perform scalability testing on the TPC-DS dataset and Duckdb. [Ham23]

# 3

## Benchmark Implementation Process

### 3.1 Setting up the Project

Some members of our group used Windows OS, MacOS and Ubuntu OS so we installed DuckDB directly on our computers . The Apple MacBook was the main machine for the benchmark implementation as all the results were obtained from it, while the Hp Elitebook 840 G5 was used for performing the scalability test with Apache Jmeter.

#### 3.1.1 Hardware Specification

In order to carry out the benchmark experimentation, Apple Macbook and HP EliteBook 840 G5 were used. Table 3.1 shows the hardware specification of this computer:

Name	HP EliteBook 840 G5	Apple Macbook
Operating System	Ubuntu 22.04.3 LTS	MacOS 13.0
System Type	64 Bits	64 Bits
CPU	Intel® Core™ i7-8650U	Apple M1 Pro
CPU Frequency	1.90GHz × 8	3.2 GHz
RAM capacity	24 GB	16 GB
RAM type	DDR4	SDRAM
Hard disk capacity	512 GB	512 GB

Table 3.1: System Specifications.

## 3.2 Implementing TPC-DS on DuckDB

Figure 3.1 shows a business process model depicting a brief rundown of the implementation the TPC-DS benchmark on DuckDB.

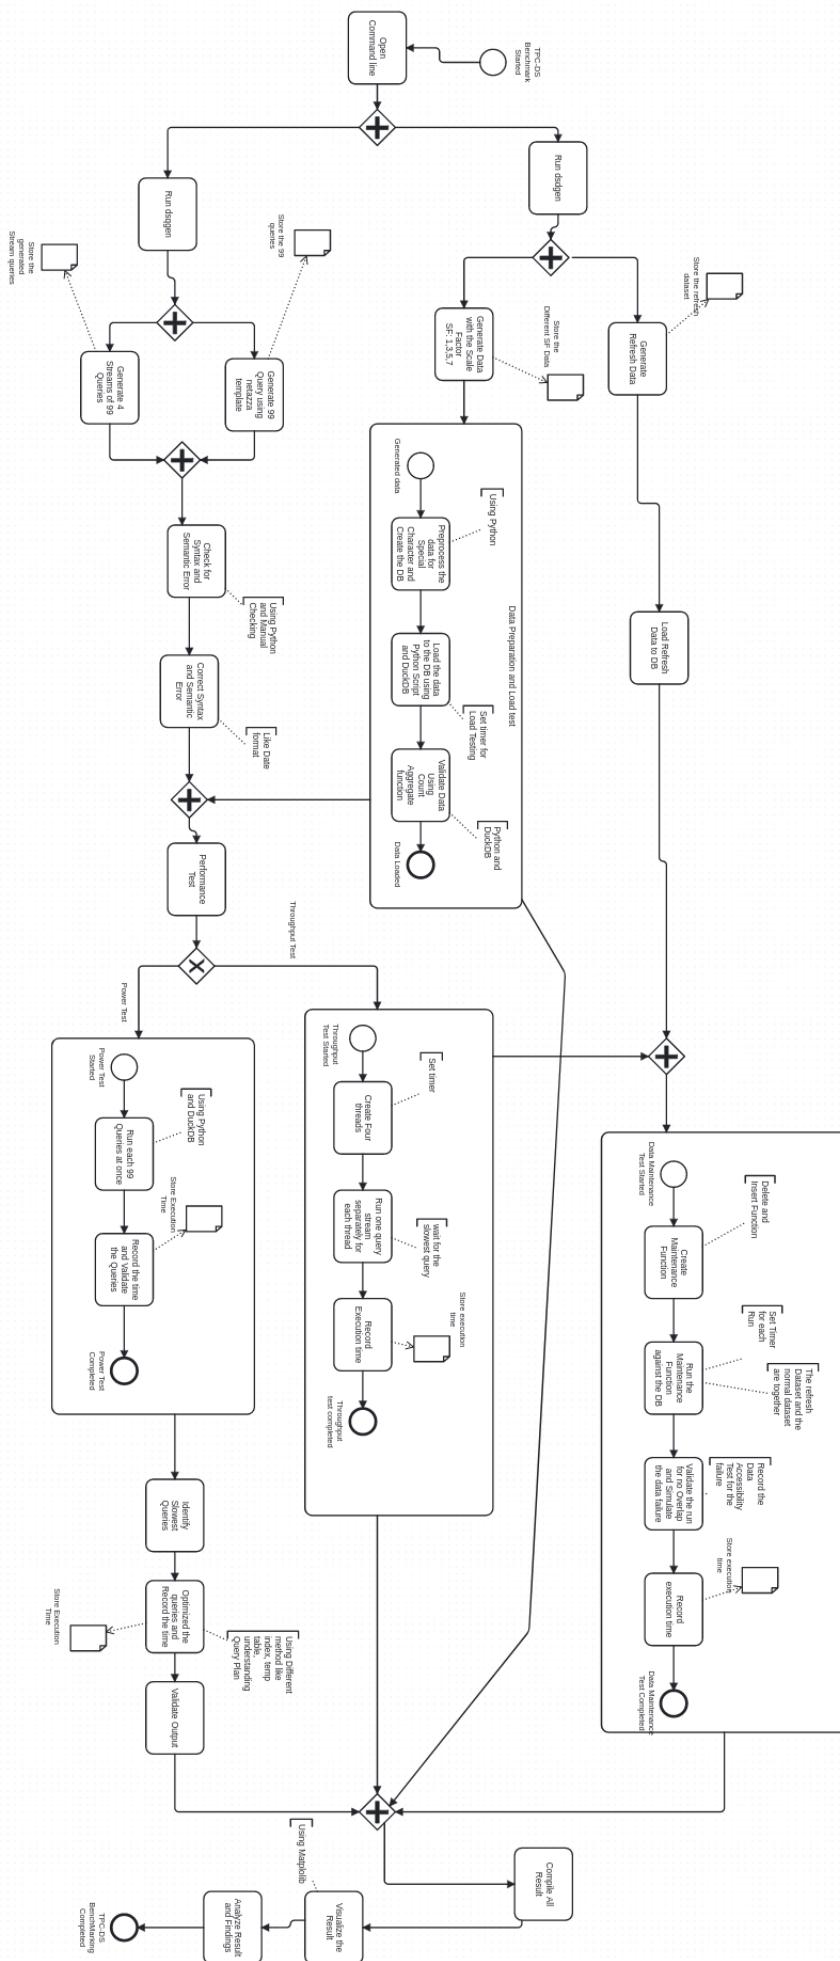


Figure 3.1: TPC-DS Benchmarking BPMN Diagram

From Figure 3.1, some of the process that was done include:

- Using dsdgen Utility from the TPC-DS tool kit to obtain load data for all our scale factor and also getting our refresh dataset too
- Generating 99 queries via the dsqgen using the Netezza SQL template and also four stream queries for our throughput test
- Matching the queries with DuckDB syntax to check for syntax mismatches and error.
- Using python scripts with DuckDB to perform all tests on the queries and compile all results.
- Discover slowest running queries in order to analyze and optimize them.
- Visualize the outputs and provide detailed reports.

### 3.2.1 Scaling and Database Population

The TPC-DS benchmark defines a set of discrete scaling factors based on the size of the raw data produced by dsdgen (from the TPC-DS toolkit). This raw data will be used to perform the benchmark [TPC21]. The standard set of scale factors defined for TPC-DS is 1TB, 3TB, 10TB, 30TB, 100TB, but due to storage, computing and memory limit , we opted for much smaller scale factors to perform our tests: 1GB, 3GB, 5GB, 7GB

## 3.3 Implementation Process

The implementation of the TPC-DS benchmarking were done in different stages, which include:

1. Database schema creation
2. Data generation
3. Query generaion
4. Tests for 4 scale factors

## 3.4 Setup

### 3.4.1 DuckDB installation

We use pip<sup>1</sup> - the package installer for Python - to install the DuckDB.

```
1 pip install duckdb==0.9.1
```

Listing 3.1: DuckDB installation.

### 3.4.2 TPC-DS kit installation

We install the TPC-DS kit from the tpcds- toolkit downloaded from TPC website<sup>2</sup>.

### 3.4.3 Apache Jmeter Installation

#### On Linux

Installing Apache JMeter on Ubuntu and other Linux Distributions requires java installed. You can either use JRE or JDK.

1. To install the default java on Ubuntu, type: \$ sudo apt install default-jre
2. Download Apache JMeter by visiting Apache JMeter website to download the latest version of JMeter.
3. Extract Apache jmeter Archive file from Downloads Folder
4. Run the Jmeter by first changing the directory as follows: \$ cd apache-jmeter-5.6.2/bin
5. Now run the JMeter using the following command: \$ ./jmeter

---

<sup>1</sup><https://duckdb.org/docs/installation/index>

<sup>2</sup><https://www.tpc.org/tpcds/>

## On Windows

Installing Apache JMeter on Windows requires java installed.

1. If Java is not install, Download and Install it <sup>3</sup>
2. Download Apache JMeter by visiting Apache JMeter website<sup>4</sup> to download the latest version of JMeter.
3. Unzip Apache jmeter Archive file from Downloads Folder
4. Go to the jmeter directory and the bin folder, double click on the jmeter.bat file

## 3.5 Database schema creation

The first step in the process is to connect to DuckDB and create a database:

```
1 import duckdb
2 con = duckdb.connect('sc_1.db') # SF-1
```

Listing 3.2: Database setup

Then we use the scripts in tpcds-kit/tools/tpcds.sql to create the database scheme. Below is an example of table creation in the scripts:

```
1 con.sql(
2 "create table dbgen_version
3 (
4     dv_version          varchar(16)      ,
5     dv_create_date      date            ,
6     dv_create_time      time            ,
7     dv_cmdline_args    varchar(200)
8 );"
9 )
```

Listing 3.3: Database setup

---

<sup>3</sup><https://www.oracle.com/java/technologies/downloads/>

<sup>4</sup>[https://jmeter.apache.org/download\\_jmeter.cgi](https://jmeter.apache.org/download_jmeter.cgi)

## 3.6 Data generation

We generated data for evaluation with the TPC-DS program dsdgen. At this step, we are in the tpcds-kit/tools folder. We ran the following command lines to generate data in 4 scale factors and store them in the generated-data folder.

```

1 ./dsdgen -scale 1 -dir "../../generated-data/scale_1" -delimiter "," # SF 1
2 ./dsdgen -scale 3 -dir "../../generated-data/scale_3" -delimiter "," # SF 3
3 ./dsdgen -scale 5 -dir "../../generated-data/scale_5" -delimiter "," # SF 5
4 ./dsdgen -scale 7 -dir "../../generated-data/sale_7" -delimiter "," # SF 7

```

Listing 3.4: Data generation.

After this stage, 25 .csv files are generated regardless of the scale factor. They only differ in size for each scale. We show an example of a row in call\_center.csv in :

```

1 1|AAAAAAAAABAAAAAAA|1998-01-01|||2450952|NY Metro|large|2|1138|8AM-4PM|Bob
    Belcher|6|More than other authori|Shared others could not count fully
    dollars. New members ca|Julius Tran|3|pri|6|cally|730|Ash Hill|Boulevard|
    Suite 0|Midway|Williamson County|TN|31904|United States|-5|0.11|

```

Listing 3.5: Generated data example.

As can be seen from Listing 3.5, data fields are separated by the " | " delimiter. A null value is represented by " || ", and we leave it by default because DuckDB is able to identify it and produce a NULL value.

Also, We generated refresh data for all scale factors (1,3,5,7) using dsdgen for separate refresh runs, so that we can implement the data maintenance test. Refresh runs must be executed using refresh data sets in the order generated by dsdgen.

## 3.7 Query generation

We generated data for evaluation with the TPC-DS program dsqgen. We are still in the tpcds-kit/tools folder and ran the following command lines to generate 99 queries for SF 1 (we do similarly for other scales) and store them in the generated-queries folder.

```
1 dsqgen -directory "../query_templates/" -input "../query_templates/templates
 .lst" -dialect netezza -scale 1 -output_dir "../../generated-queries"
```

Listing 3.6: Data generation.

This command generates a file query\_0.sql that contains 99 SQL queries with netezza dialect.

### 3.7.1 Separating queries from the templates

We have developed a python script called split\_queries.py. This script basically splits the generated queries created by the dsqgen called query\_0.sql and puts them in 99 different files, named query\_n.sql, where n is from 1, ..., 99.

### 3.7.2 Query Modification

After generating the queries, we had to fix syntax in some SQL files so that we could generate queries that can work well in our DuckDB Database (i.e. uses DuckDB SQL dialect). In particular, we edited SQL files for queries 5, 12, 16, 20, 21, 30, 32, 37, 40, 58, 77, 80, 82, 92, 94, 95, 98. Here are the specific changes we made:

- Adding intervals to the cast(as date) function– to identify the interval of days.
- Changing the reserved keyword used as a variable name e.g in query 77.
- Fixing some ambiguous variable names in query 30 and 58.

### 3.7.3 Running Queries

DuckDB is an **embedded, in-process**, relational, OnLine Analytical Processing (OLAP) DataBase Management System (DBMS). This means that the DBMS features are running from within the application you're trying to access from instead of an external process your application connects to. In our case , we use Python API to run the queries and also we used DBeaver to test the queries, modify the queries and execute them.

## 3.8 Load Test

The generated data was loaded into the tables of DuckDB to measure how good the system does this job. This is called **Load Test**. This data loading procedure has to be done automatically, not involving manual interaction so as to measure the time accurately. **The Database Load Time** is the difference between **Load Start Time** and **Load End Time**, where

- Load Start Time is defined as the timestamp when the first character is read from any of the flat files (dat or csv files)
- Load End Time is defined as the timestamp taken when the DuckDB Database is fully populated.

The pseudo code for data loading and load test is as follows:

```

1 Assumption:
2     The schema defined in tpcds.sql in official tools has been loaded to
3     the database.
4
4 Algorithm:
5     set timer start
6     load data into database
7     set timer end

```

Listing 3.7: Load Test Pseudocode

We execute the following Python script to perform the load test:

```

1 import duckdb
2 import os
3 import time
4
5 sc = '1' # scale factor
6 # specify the path for dataset
7 csv_folder_name = 'generated_data/sc_'+ sc
8 csv_list = os.listdir(csv_folder_name)
9
10 # set the timer start
11 start_time = time.time()
12 for csv_file in csv_list:

```

```

13     table_name = csv_file[:-4]
14
15     # load data into tables
16
17     copy_command = "COPY "+table_name+" FROM "+"+" +csv_folder_name+ "/" +
18         format(csv_file, "")+""
19
20     con.sql(copy_command)
21
22     # set the timer end
23
24     end_time = time.time()
25
26     time_taken = end_time - start_time
27
28     print(time_taken)

```

Listing 3.8: Python script for load test.

## 3.9 Power Test

The Power Test measures the ability of the system to process a sequence of queries in the least amount of time in a single stream. This Test is executed immediately following the load test. Only one query can be executed at a time during the power test. The **elapsed time of the Power test** is the difference between **Power test start time** and **Power test end time**, where

- Power Test Start Time, which is the timestamp that must be taken before the execution of the first query of Stream 0.
- Power Test End Time, which is the timestamp that must be taken right after the execution of the last query of Stream 0.

Below is the pseudocode and the Python scripts for the power test.

```

1 Assumption:
2     - The Load Test has already been executed.
3
4 Algorithm:
5     set timer start
6     for every query in 99 generated queries:
7         set timer start
8         run the query
9         set timer end
10    set timer end

```

Listing 3.9: Power test pseudocode.

```

1 import duckdb
2
3 import os
4 import fnmatch
5 import time
6 import pandas as pd
7 import pickle
8
9 sc = "1" # scale factor
10 start_time = time.time()
11 conn = duckdb.connect("sc_"+sc+".db")
12
13 list_time = []
14 # read and run all the queries
15 for root, dirnames, filenames in os.walk("generated-queries/queries_0"):
16     for filename in fnmatch.filter(filenames, '*.sql'):
17         t = time.time()
18         sql = open(root+ "/" + filename, 'r')
19         sql = sql.read()
20
21         # run all the queries
22         result = conn.execute(sql)
23         close = time.time()
24         list_time[filename] = close - t
25
26 # Calculate the elapsed time
27 end_time= time.time()
28 elapsed_time = end_time -start_time

```

Listing 3.10: Power test Python script.

## 3.10 Throughput Test

The Throughput Tests measure the ability of the system to process the most queries in the least amount of time with multiple users. It measures the Database's multi-user functionalities. There are two Throughput tests during the benchmarking, Throughput Test 1 immediately follows the Power Test while Throughput Test 2 immediately follows Data Maintenance Test 1. During the test, only one query shall be active on any of the sessions at any point of time.

The **elapsed time of Throughput Test 1** is the difference between **Throughput Test 1 Start Time** and **Throughput Test 1 End Time**.

- Throughput Test 1 Start Time: timestamp taken before the first query of the first user starts.
- Throughput Test 1 End Time: timestamp taken after the last query of the last user finishes.

The **elapsed time of Throughput Test 2** is the difference between **Throughput Test 2 Start Time** and **Throughput Test 2 End Time**

- Throughput Test 2 Start Time is defined as a timestamp taken before the first query of the first user starts immediately after Data Maintenance Test 1 End Time.
- Throughput Test 2 End Time, is defined as the timestamp taken after the last query of the last user finishes.

To accomplish this, We used 4 different threads having access to the database at the same time to simulate different users. Also, 4 query streams are generated to simulate different concurrent requests. Below are the pseudocode and the Python script for the throughput test.

```

1 Assumption:
2   - 4 query streams are generated according to the specified order of
      queries defined in the documentation.
3
4 Algorithm:
5   create 4 threads
6   set timer start
7   for each thread
8     run one query stream separately
9     wait for the slowest query
10    set timer end

```

Listing 3.11: Throughput test pseudocode

```

1 import duckdb
2 import os
3 import fnmatch
4 import time
5 import threading

```

```
6
7 sc = "1" # scale factor
8 db_connection = duckdb.connect("sc_"+sc+".db")
9
10 def execute_query(query, file_name):
11     global db_connection # Use the shared database connection
12     start_time = time.time()
13
14     cursor = db_connection.cursor()
15     cursor.execute(query)
16
17     # Calculate the elapsed time
18     end_time = time.time()
19     elapsed_time = end_time - start_time
20
21 def main():
22     TP_test_start_time_1 = time.time()
23     queries = []
24     threads = []
25
26     for root, dirnames, filenames in os.walk('data/throughput_queries'):
27         for filename in fnmatch.filter(filenames, '*.sql'):
28             sql_file = open(os.path.join(root, filename), 'r')
29             query = sql_file.read()
30             queries.append((query, filename))
31
32     # create threads and assign query stream to threads
33     for query, file_name in queries:
34         thread = threading.Thread(target=execute_query, args=(query,
35         file_name))
36         threads.append(thread)
37         thread.start()
38
39     # Wait for all threads to complete
40     for thread in threads:
41         thread.join()
42
43     # calculate the time for throughput test
44     TP_test_end_time_1 = time.time()
        TP_test_time_1 = TP_test_end_time_1 - TP_test_start_time_1
```

```

45     output = f'Throughput test time = {TP_test_time_1}\n'
46
47     with open("throughput_test/output.txt", 'w+') as f:
48         f.write(output)
49
50 if __name__ == "__main__":
51     main()

```

Listing 3.12: Throughput test Python script.

## 3.11 Data Maintenance Test

The Data Maintenance Tests is another benchmark that measures the ability to perform desired data changes to the TPC-DS data set. Data Maintenance Test 1 immediately follows Throughput Test 1 and Data Maintenance Test 2 immediately follows Throughput Test 2.

To achieve this purpose, refresh data sets are generated by dsdgen, and different maintenance functions are applied to the database:

- Fact insert data maintenance: new facts are inserted into the database.
- Fact delete data maintenance: some facts are deleted from the database, using date filters.
- Inventory delete data maintenance: some data is deleted from the inventory table, using date filter

The Durable Medium failure required as part of the Data Accessibility Test must be triggered during Data Maintenance Test 1 and also during the Data Maintenance Test 2.

The elapsed time of **Data Maintenance Test 1** is the difference between:

- **Data Maintenance Test 1 Start Time**, defined as the starting timestamp of the first refresh run in Data Maintenance Test 1 and
- **Data Maintenance Test 1 End Time**, defined as the ending timestamp, of the last refresh run in Data Maintenance Test 1.

The elapsed time of **Data Maintenance Test 2** is the difference between:

- **Data Maintenance Test 2 Start Time**, defined as the starting timestamp of the first refresh run in Data Maintenance Test 2 and
- **Data Maintenance Test 2 End Time**, defined as the ending timestamp, of the last refresh run in Data Maintenance Test 2.

```

1 Assumption:
2   - Refresh dataset are generated.
3   - All data maintenance functions are implemented.
4
5 Algorithm:
6   set timer start
7   for every refresh run:
8     load the refresh dataset
9     for every delete function:
10       set timer start
11       run the query
12       set timer end
13     for every insert function:
14       set timer start
15       run the query
16       set timer start
17     set timer end

```

Listing 3.13: Data maintenance test pseudocode.

## Refresh Dataset Generation

During the Data Maintenance Test, the dataset is also generated by dsdgen but different from the dataset we used in previous tests.

```

1 dsdgen -scale 1 -dir "../../generated-data/scale_1" -delimiter "," update 1
2 dsdgen -scale 3 -dir "../../generated-data/scale_3" -delimiter "," update 1
3 dsdgen -scale 5 -dir "../../generated-data/scale_5" -delimiter "," update 1
4 dsdgen -scale 7 -dir "../../generated-data/sale_7" -delimiter "," update 1

```

Listing 3.14: Refresh Data generation.

We use the refresh dataset to populate the schema defined in TPC-DS official tools.

## Maintenance Functions

In order to test the maintenance of DuckDB, we implemented 11 functions following the TPC-DS documentation. The functions are listed in the Table 3.2.

Table 3.2: Data Maintenance Functions

ID	Function	Type	Target Table	Source Schema Tables
1	LF_CR	Insert	catalog_returns	s_catalog_returns
2	LF_CS	Insert	catalog_sales	s_catalog_order, s_catalog_lineitem
3	LF_I	Insert	inventory	s_inventory
4	LF_SR	Insert	store_returns	s_store_returns
5	LF_SS	Insert	store_sales	s_purchase, s_purchase_lineitem
6	LF_WR	Insert	web_returns	s_web_returns
7	LF_WS	Insert	web_sales	s_web_order, s_web_order_lineitem
8	DF_CS	Delete	catalog_sales, catalog_returns	-
9	DF_SS	Delete	store_sales, store_returns	-
10	DF_WS	Delete	web_sales, web_returns	-
11	DF_I	Delete	inventory	-

Below is the Python script for the data maintenance test.

```

1 import os
2 import time
3 import duckdb
4 import threading
5
6 sc = "1" # scale factor
7 con = duckdb.connect("sc_"+sc+".db")
8
9 # simulate the database failure for the data accessibility test
10 def simulate_duckdb_failure():
11
12     # Record the start time before triggering the failure
13
14     failure_start_time = time.time()
15
16     # Simulate a failure event by stopping DuckDB
17     print("Simulating DuckDB failure during Data Maintenance Test 1...")
18
19     time.sleep(20)
20
21     # Record the end time after triggering the failure

```

```
22     failure_end_time = time.time()
23     print(f"DuckDB failure time: {failure_end_time - failure_start_time:.2f} seconds")
24
25 # Function to run all maintenance functions in a folder
26 def run_sql_queries_in_folder(folder_path):
27     global con
28     if folder_path.endswith("2"):
29         run = 2
30     else:
31         run = 1
32
33     # Get a list of SQL files in the folder
34     py_files = [file for file in sorted(os.listdir(folder_path)) if file.endswith(".py")]
35     time_dict = {}
36
37     # Loop through the SQL files and execute them
38     for py_file in py_files:
39         with open(os.path.join(folder_path, py_file), "r") as sql_script:
40             t=time.time()
41             py_query = sql_script.read()
42
43             try:
44                 exec(py_query)
45                 print(f"successfully executed! {py_file} with time {time.time()-t} seconds")
46                 time_dict[py_file] = time.time() - t
47             except Exception as e:
48                 print(f"problem executing {py_file}! : {e}")
49                 ttime = time.time()
50
51                 # Simulate DuckDB failure in a separate thread
52                 failure_thread = threading.Thread(target=simulate_duckdb_failure)
53                 failure_thread.start()
54
55                 # Wait for the failure to complete
56                 failure_thread.join()
57                 print("Durable media failure completed.")
```

```

58         #Add end time
59         tftime= time.time()
60         time_dict["failure"] = tftime-ttime
61         print(f'completed: {tftime-ttime}')
62     return time_dict
63
64 # List all subdirectories in the root directory
65 b_time= time.time()
66 subdirectories = [d for d in os.listdir(refresh-data) if os.path.isdir(os.
67     path.join(refresh-data, d))]
68
69 # record the result for data maintenance test
70 result_run_1 = run_sql_queries_in_folder("maintenance_functions")
71 result_run_2 = run_sql_queries_in_folder("maintenance_functions_2")
72 e_time= time.time()
73 print(f"This is the end of the Data maintenance query: {e_time-b_time}")

```

Listing 3.15: Python script for maintenance test.

## 3.12 Data Accessibility Test

The Data Accessibility Test is performed by causing the failure of a Durable Media during the execution of the first Data Maintenance Test. In our benchmarking Project, we caused an exception by introducing a file with wrong syntax and errors so as to simulate the data failure process.

The Data Accessibility Test is successful if all in-flight data maintenance functions, subsequent queries and data maintenance functions complete successfully after the above durable media have failed.

## 3.13 Scalability Test

Scalability Testing measures performance of a system or network when the number of user requests are scaled up or down. The purpose of Scalability testing is to ensure that the system can handle projected increase in user traffic, data volume, transaction counts frequency, etc. It tests the system's ability to meet the growing needs. For our Project we wanted to know how the

database can scale up when the user increases and at what point will the system become unresponsive. To achieve this, we had to implement the TPC-DS dataset and DuckDB database on Apache Jmeter

## Using Apache JMeter

To test your database, you need to

- Create a test plan
- Execute the test plan

### Create a Test Plan

A basic test plan consists of the following elements:

- Download DuckDB Connector
- Thread Group
- JDBC Connection Configuration
- JDBC Request
- JDBC Listener

JMeter has a default test plan, which you can use if you don't want to create multiple test plans. However, there's no elements in the default test plan, so we need to add the missing elements. Also, add a jar path in the library to the duckdb\_jdbc jar file you downloaded.

### DuckDB JDBC Connector

We use the DuckDB database to create our Database. First, you have to connect the DuckDB database with JMeter using the DuckDB JDBC connector. First, download the duckdb\_jdbc-0.9.1.jar

from the Maven website<sup>5</sup>. Copy “duckdb\_jdbc-0.9.1.jar” file (Version may differ) from the folder and paste it into the “lib” folder of JMeter and Restart the JMeter so that it will work with the library provided by DuckDB connector.

## Create Thread Group

The Thread Group is a set of threads that performs a test scenario. In this screen, we can set the number of users and other similar settings to simulate the user requests. We right-click on the TestPlan and then we select the Add->Threads (Users)->Thread Group

## Add JDBC Connection Configuration

This is used to create a valid database connection. You add JDBC Connection Configuration by following the steps : "Add > Config Element > JDBC Connection Configuration"

Then configure your DuckDB Database with this:

**Database URL:** jdbc:duckdb:/home/pce/Pictures/duckdb\_terminal/scale\_1.db

where "/home/pce/Pictures/duckdb\_terminal"- directory of where the DB is created and scale\_1.db is the name of the DB

**Database Class:** org.duckdb.DuckDBDriver

## JDBC Request

JDBC request element helps to define a SQL query that will be executed by the test user(s). To work with data from the database, You need to Add JDBC Request. So add JDBC Request By Clicking Add > Sampler > JDBC Request we set the Variable Name of Pool declared in JDBC Connection Configuration parameter. This parameter value has to be same as the JDBC connection pool name

## JDBC Listener

Before starting our test, we need a Listener that helps to monitor and analyze the result of the test. For this test, we will use the View Results Tree, of listeners. We right-click on the JDBC Request

---

<sup>5</sup>[https://repo1.maven.org/maven2/org/duckdb/duckdb\\_jdbc/0.9.1/](https://repo1.maven.org/maven2/org/duckdb/duckdb_jdbc/0.9.1/)

and select Add->Listener->View Results Tree to monitor the detailed result of the test. At the same time, we add a Summary Report that helps to monitor a summarized result of the test. To add a summary report, we right-click on the JDBC Request and select Add->Listener->Summary Report to monitor the detailed result of the test.

Test your database performance by clicking the **Run button on the menu bar**.

# 4

## Results and Discussion

As mentioned in the previous chapter, the benchmark was performed on a local machine, with a total of four different scale factors (1, 3, 5 and 7) and our scale factors increased linearly. In this chapter, we will be focusing on a few different criteria, namely, discussing the test result, the evolution of the run times as the scale factor of the database increases and making comparison between optimized and original query performances.

### 4.1 Load Test

We ran the load test script for 4 scale factors and show the results in Figure 4.1. The charts provided the loading time for 4 scale factors. The bar chart presents the information about the Load Test Time in seconds for four scale factors. As can be seen from the chart there is a significant difference between Load Test Time for each scale factor. The time increases from 11.41 seconds for SF 1 factor up to 73.3 seconds for SF 7. The Load Test indicator increased by 7 times. This is reasonable considering the data size of 7 GB and 1 GB. The line chart also shows that the loading time increased almost linearly as the data was increased from 1GB to 7GB.

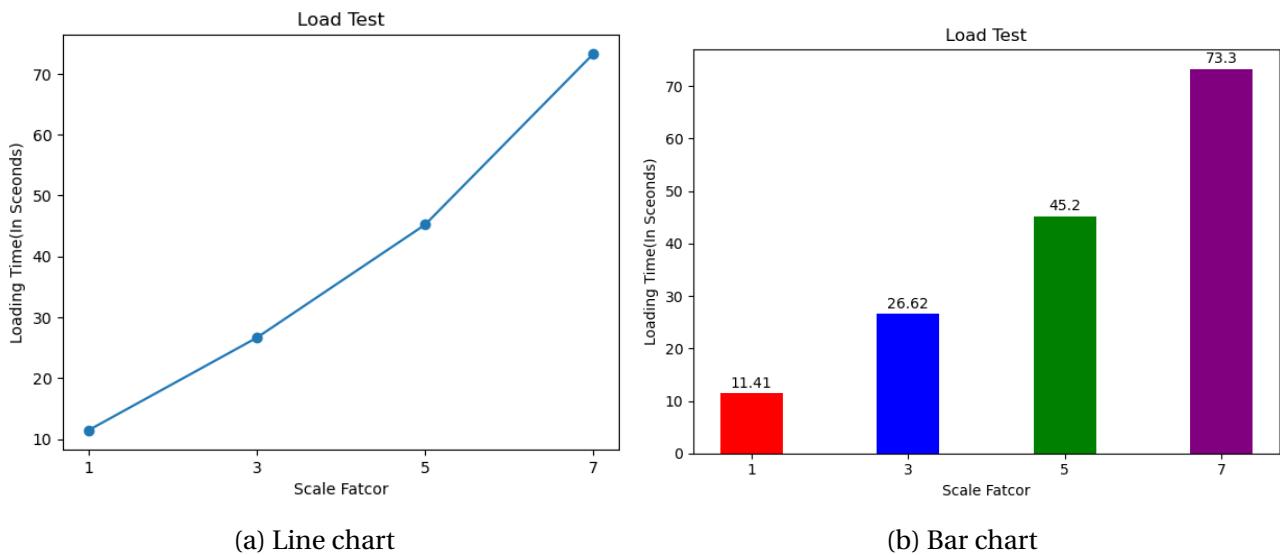


Figure 4.1: Load test results.

## 4.2 Power Test

The results of the Power test can be found in Figure 4.2. We can see the Power Time for all scale factors in the figure. From the bar chart, we can find the time increasing from 4.04 seconds for SF-1, 8.46 seconds for SF-3 to 18.32 seconds for SF-7. From the line chart, we found the increase was really close to be linear, meaning that the system was scaling nearly optimally in terms of query processing.

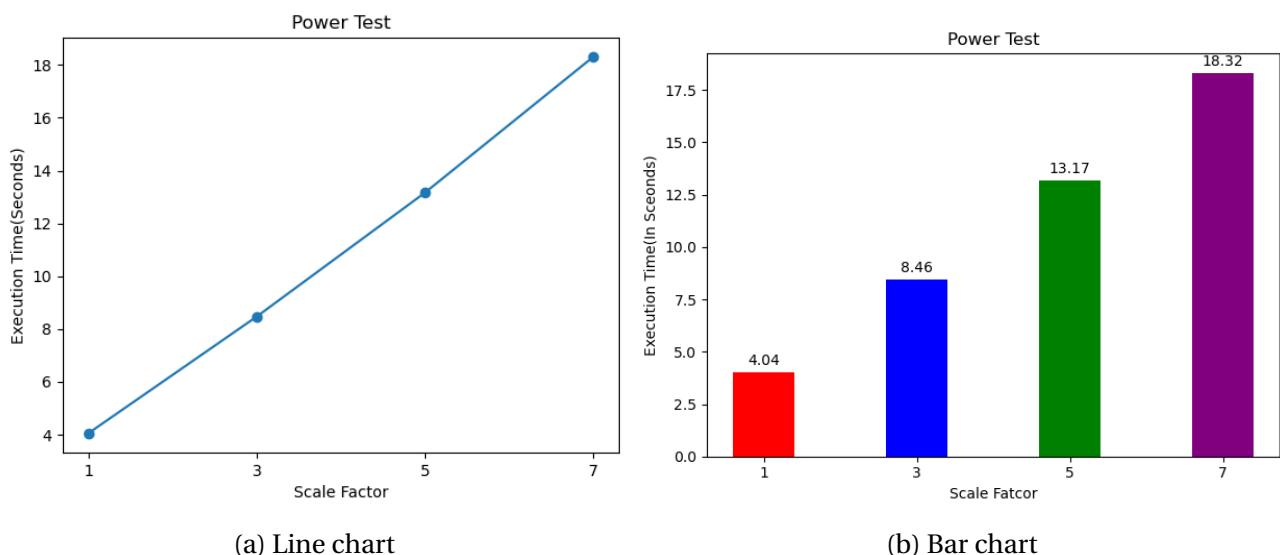


Figure 4.2: Power test results.

#### 4.2.1 Load and Power Test Comparison

It can be observed that Power and Load Test are close to being linear. But critically looking at it from the fig 4.3 and 4.4 below, We can clearly see how the load test presents a higher increase than the power test. DuckDB was able to process the query faster because it is an analytical database optimized for read operations and even faster when processed with good hardware, in our case Apple Macbook Pro.

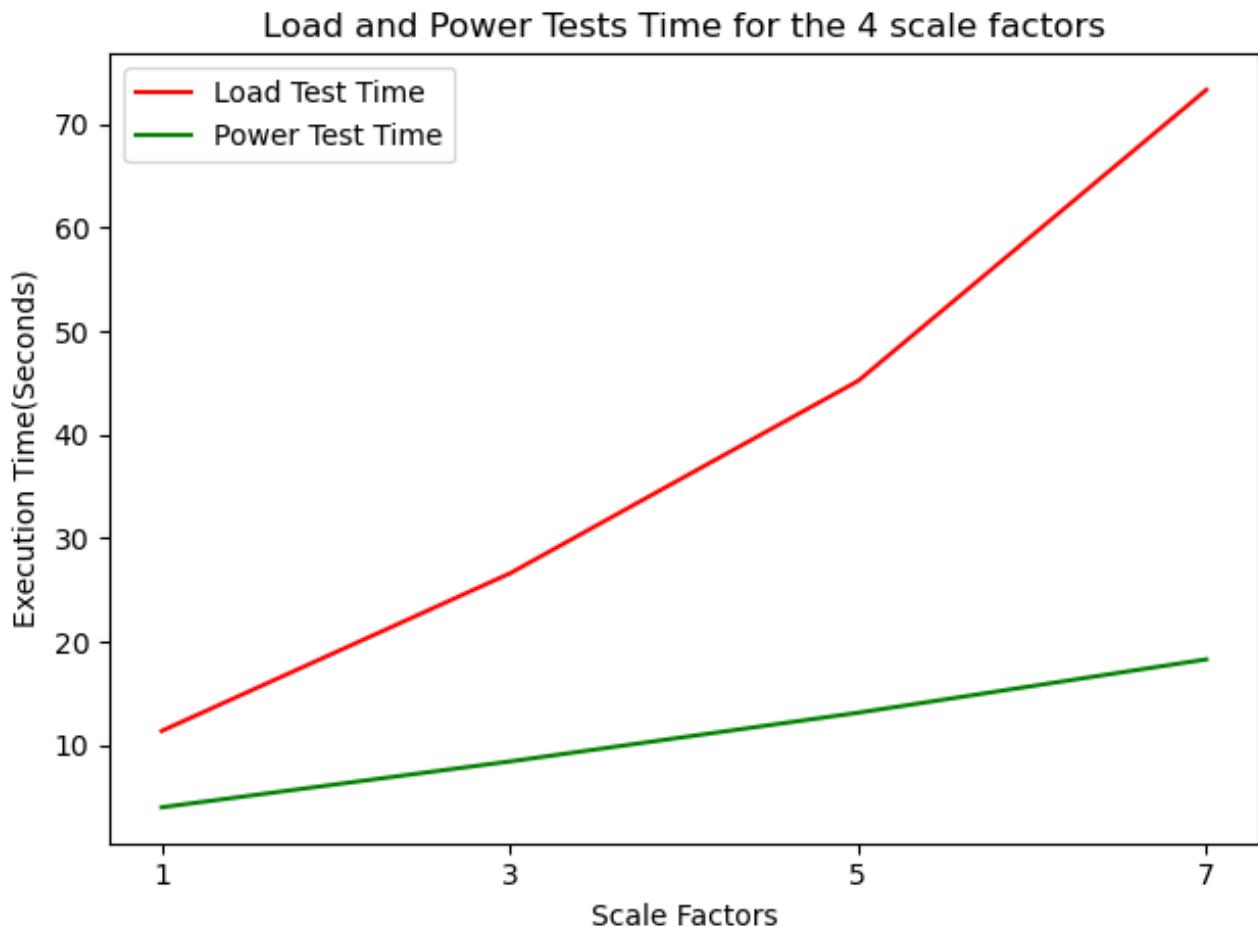


Figure 4.3: Load test and Power Test comparison(a).

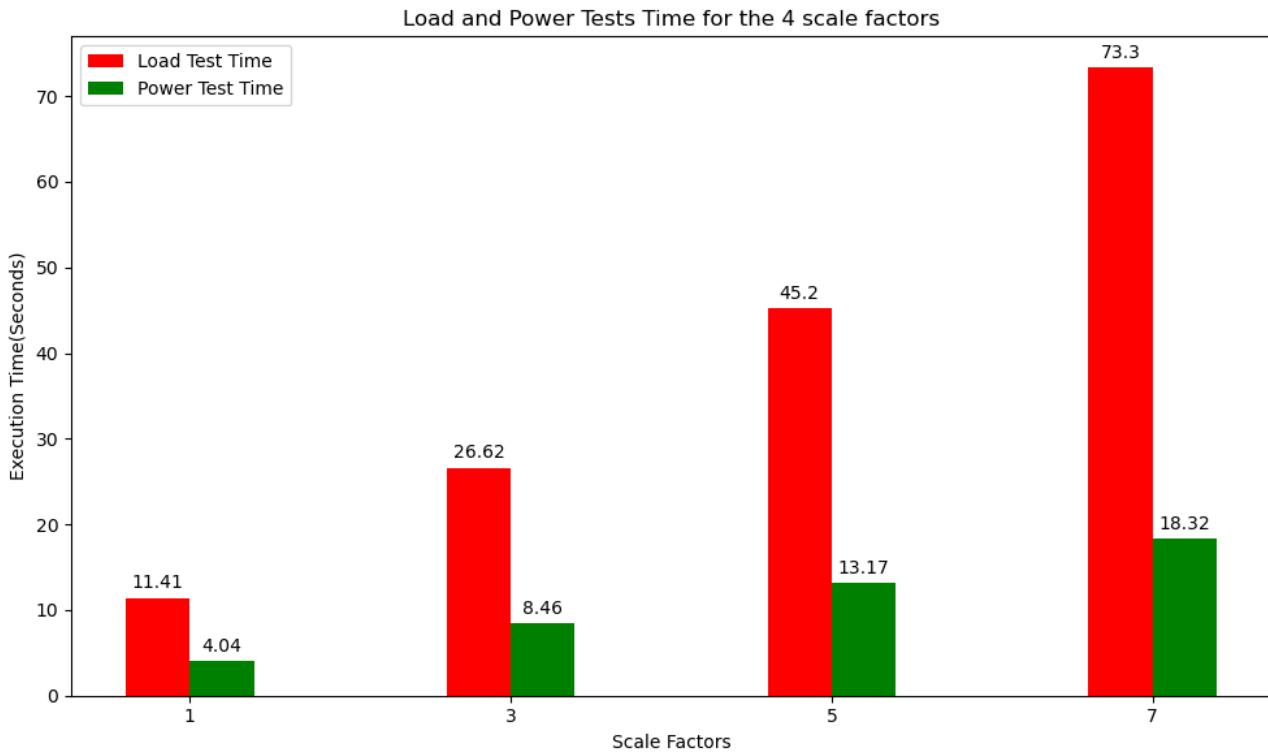


Figure 4.4: Load test and Power Test comparison(b).

#### 4.2.2 Individual Scale Execution

We analyzed the execution times for each query and each scale factor. The power test results for 4 scales are shown in Figures 4.5, 4.8, 4.11, and 4.14. Figure 4.17 compares execution time for the first 10 queries over 4 scale factors. Some queries with large running time are 23, 23b, 67, and 95. Besides, the query running time increases when the data size is bigger.

##### Scale 1

In Figure 4.5, it showed how much time it took to run the 102 queries on scale factor 1. Query 23 is the longest outlier with around 0.19935 seconds of execution times. Query 4,14,14b,22,23,23b,47,57,64, 67,69,95 are longer than other queries. The average times is about 0.03949 seconds for all the queries. So in Figure 4.7, we plotted the execution times of all the queries except for query 4,14,14b,22,23,23b, 47,57,64,67,69,95 to see their values more clearly. Query 11 is another longest outlier with around 0.10838 seconds. And the average times is about 0.02603 seconds after removing the queries.

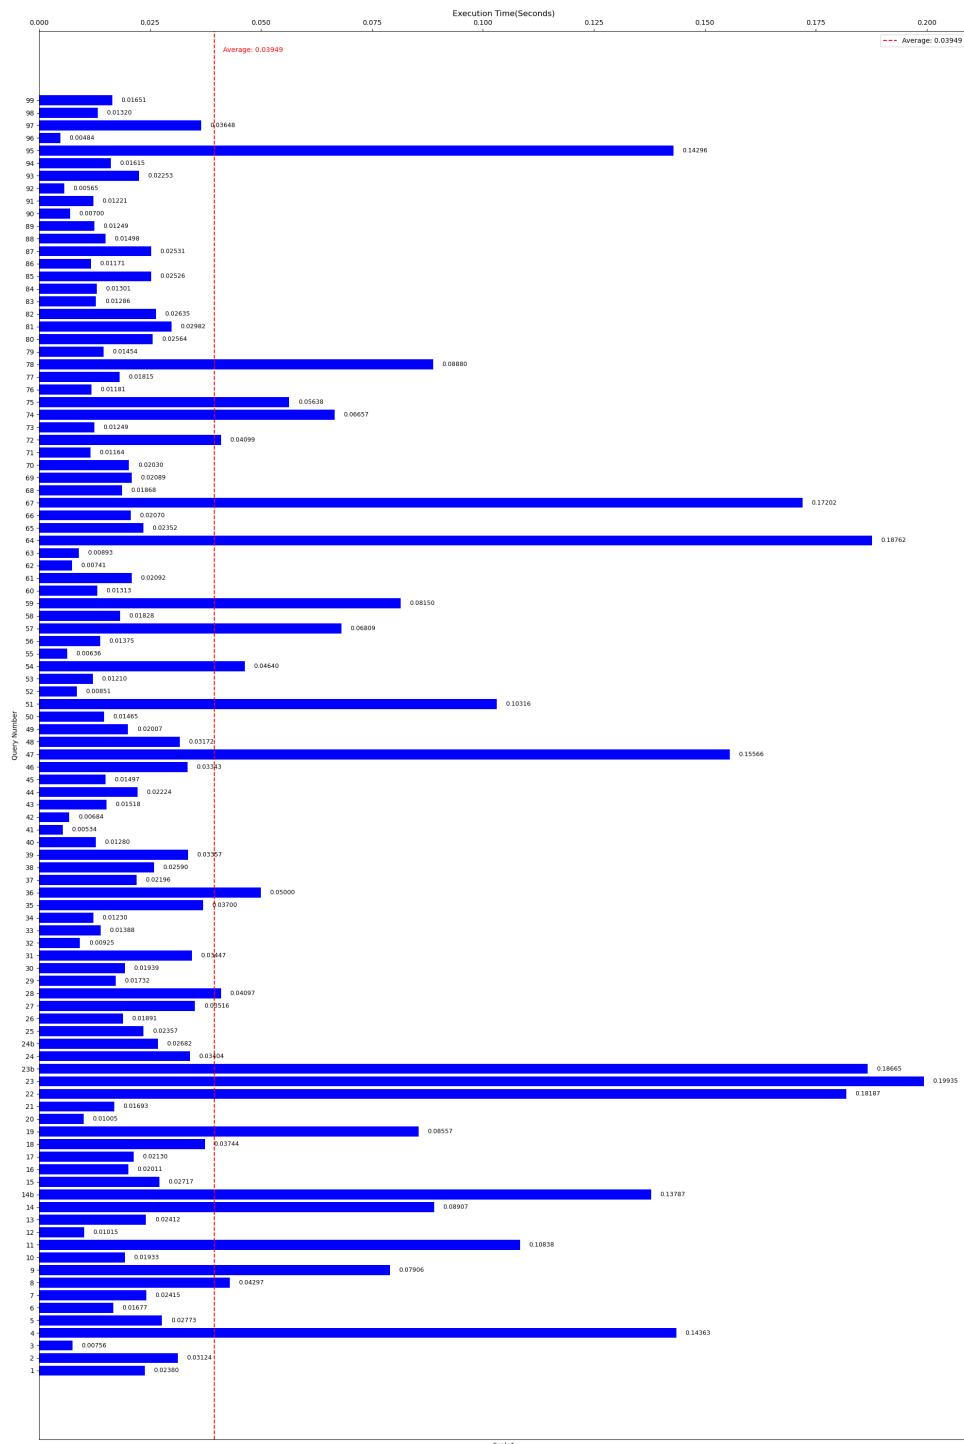


Figure 4.5: Power test result for scale factor 1.

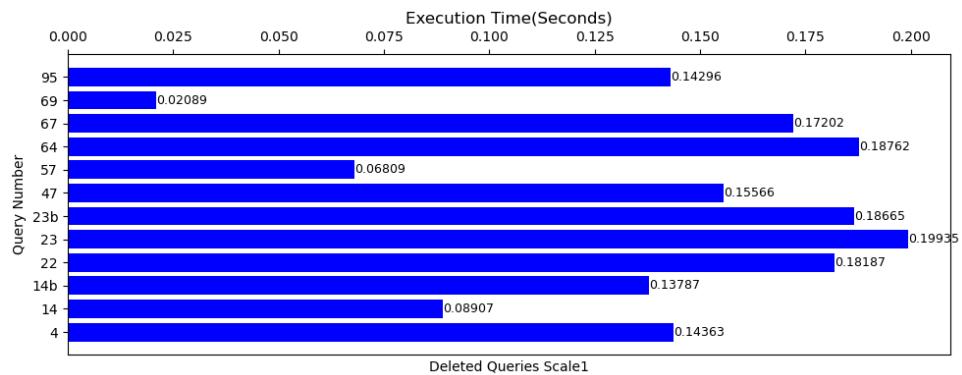


Figure 4.6: Deleted Queries for scale factor 1.

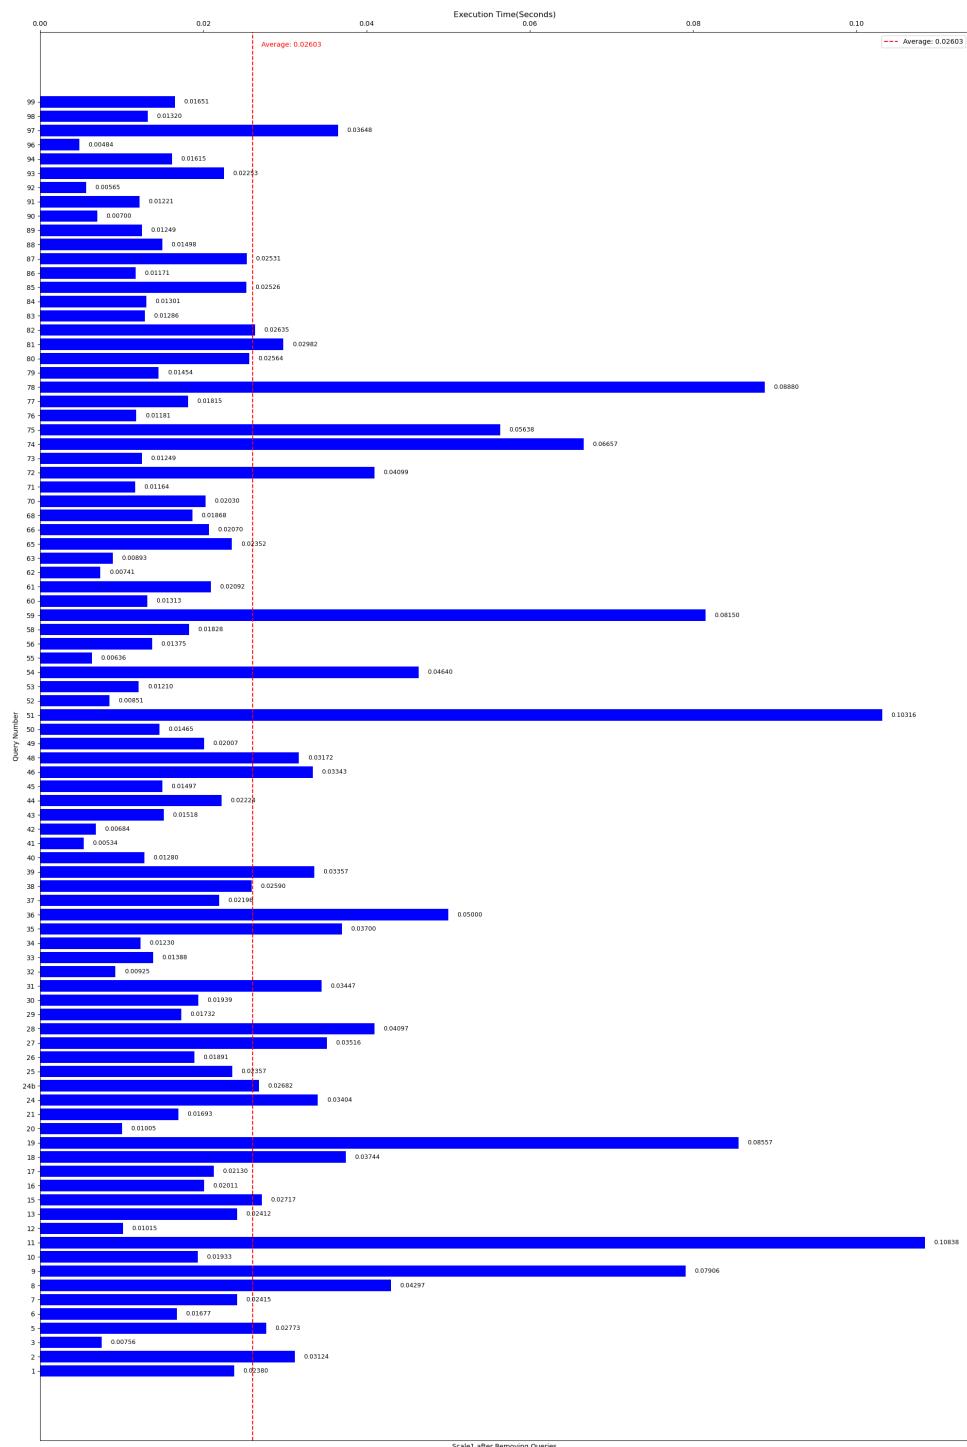


Figure 4.7: After Deleting Queries for scale factor 1.

**Scale 3**

In Figure 4.5, it showed how much time it took to run the 102 queries on scale factor 3. Query 23b is the longest outlier with around 0.52983 seconds of execution times. Query 4,14,14b,22,23,23b,47,57, 64,67,69,95 are longer than other queries. The average times is about 0.08282 seconds for all the queries. So in Figure 4.10, we plotted the execution times of all the queries except for query 4,14,14b,22,23,23b,47,57,64,67,69,95 to see their values more clearly. Query 19 is another longest outlier with around 0.26662 seconds. And the average times is about 0.05186 seconds after removing the queries.

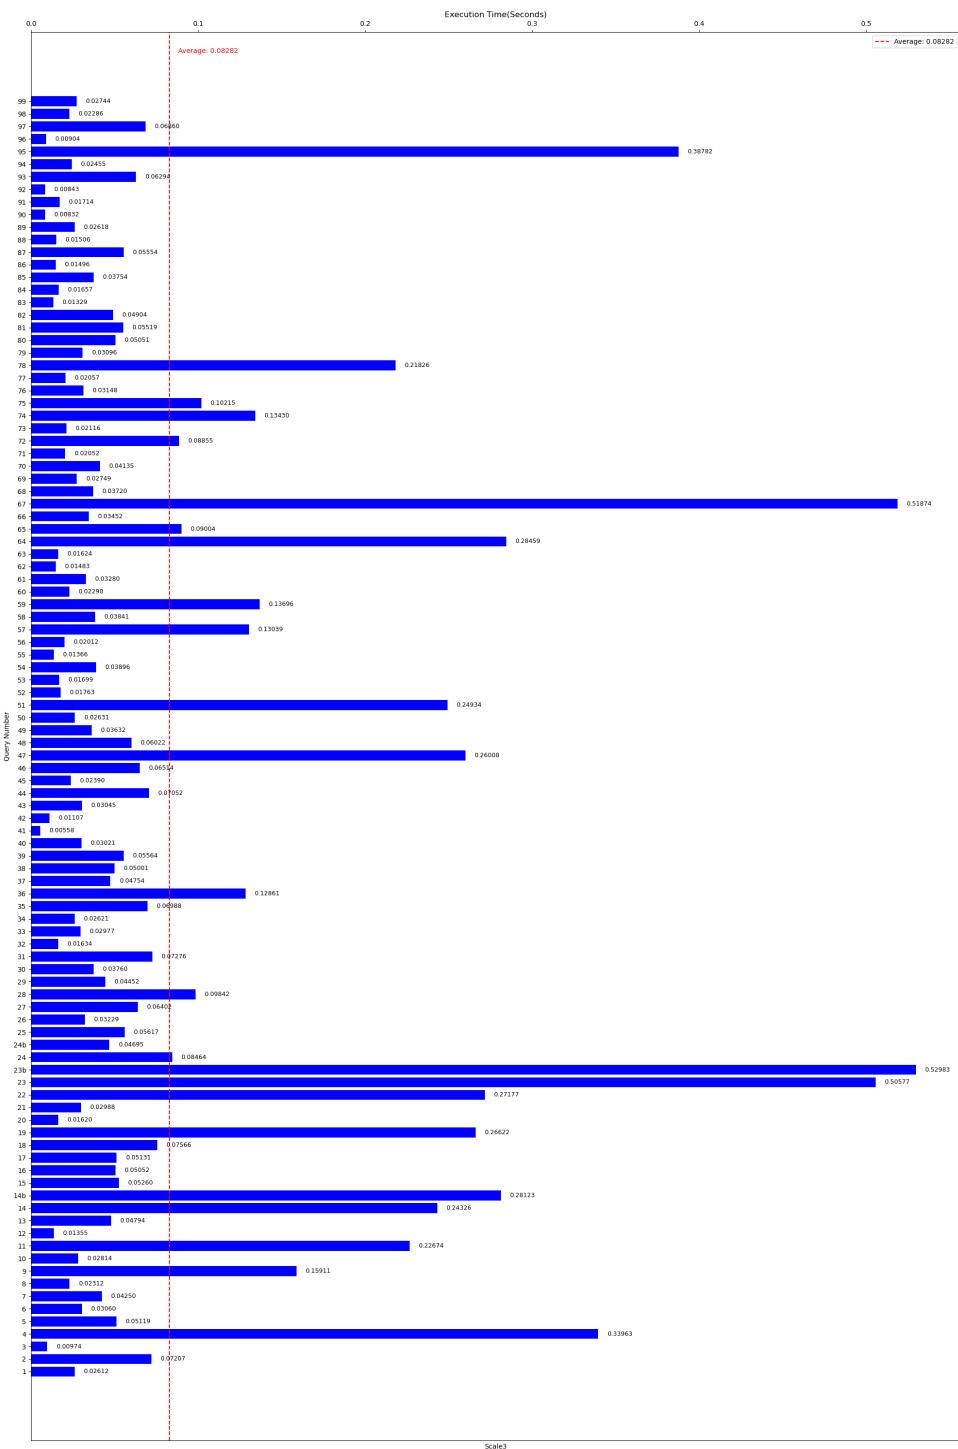


Figure 4.8: Power test result for scale factor 3.

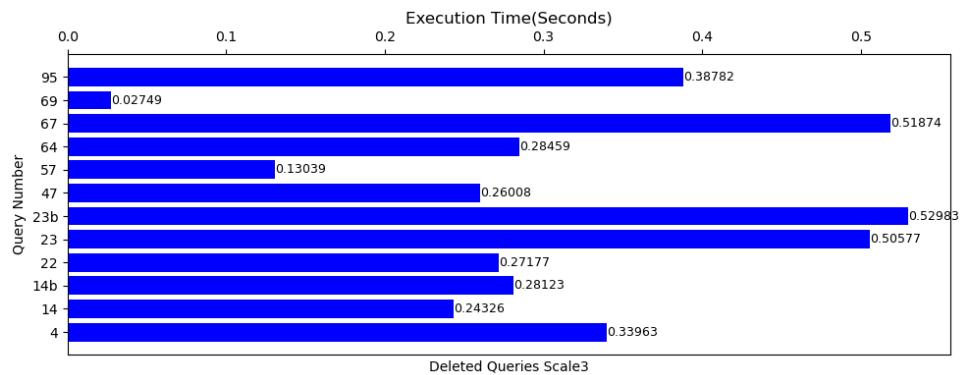


Figure 4.9: Deleted Queries for scale factor 3.

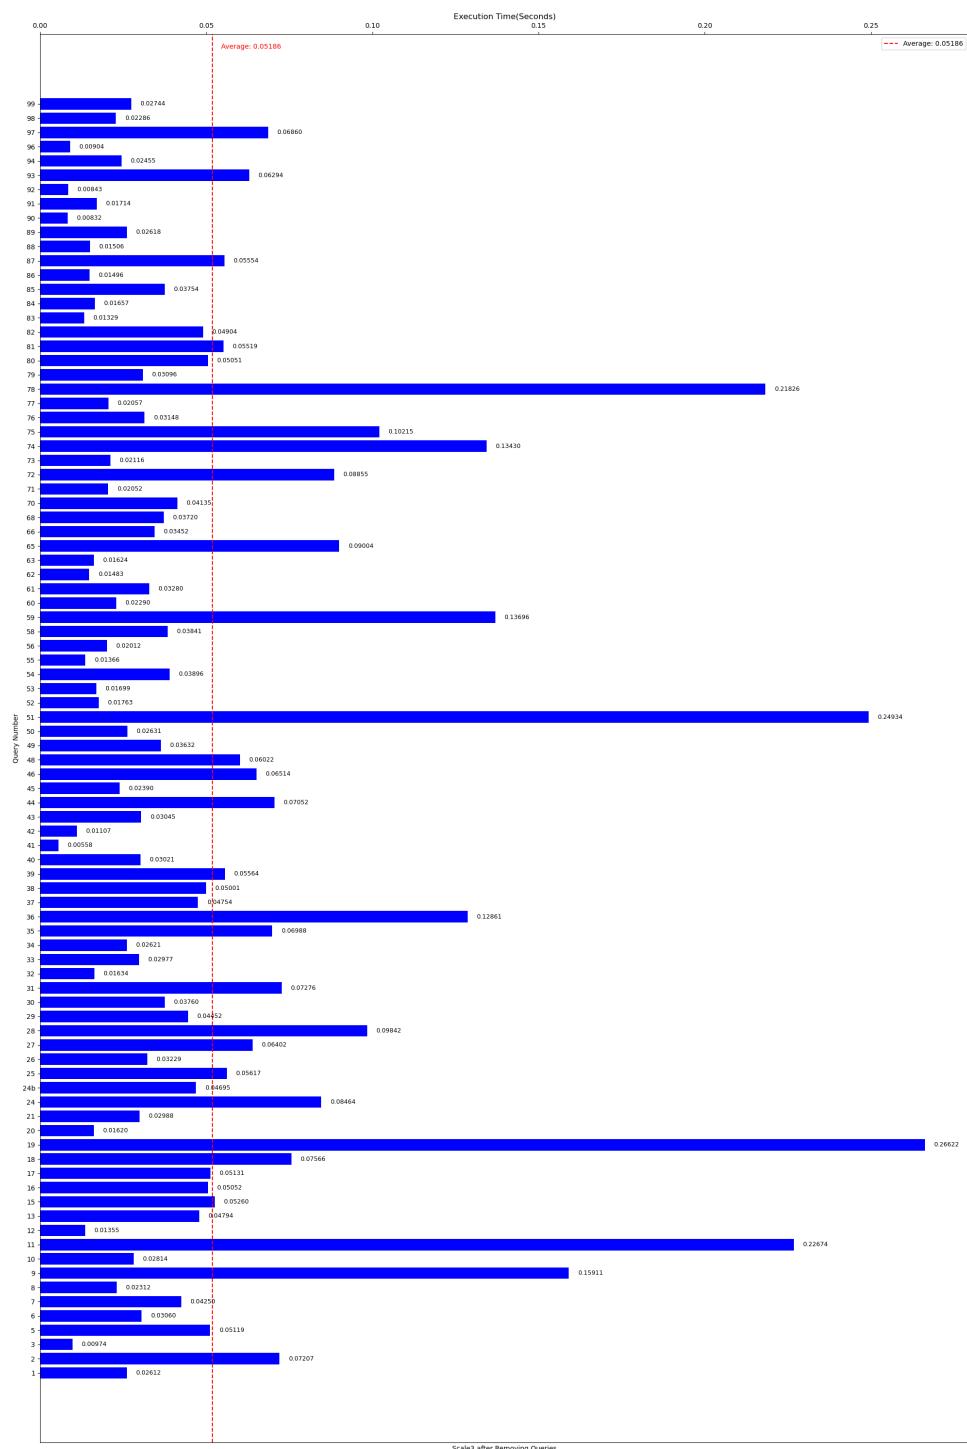


Figure 4.10: After Deleting Queries for scale factor 3.

**Scale 5**

In Figure 4.11, it showed how much time it took to run the 102 queries on scale factor 5. Query 95 is the longest outlier with around 1.12474 seconds of execution times. Query 4,14,14b,22,23,23b,47,57,64, 67,69,95 are longer than other queries. The average times is about 0.12894 seconds for all the queries. So in Figure 4.13, we plotted the execution times of all the queries except for query 4,14,14b,22,23,23b,47,57, 64,67,69,95 to see their values more clearly. Query 51 is another longest outlier with around 0.39623 seconds. And the average times is about 0.07611 seconds after removing the queries.

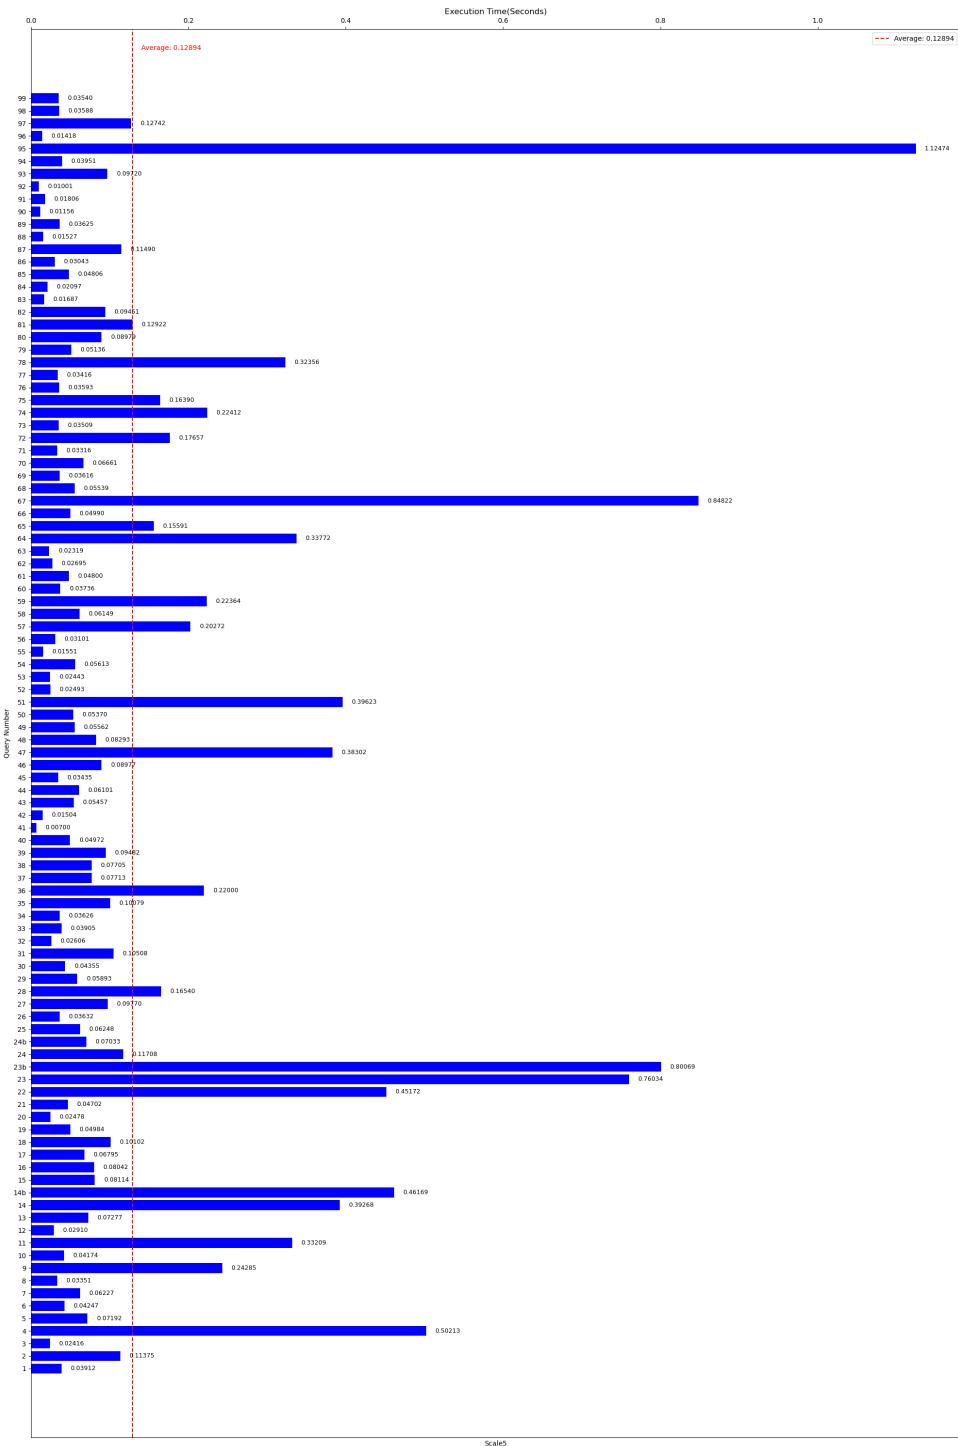


Figure 4.11: Power test result for scale factor 5.

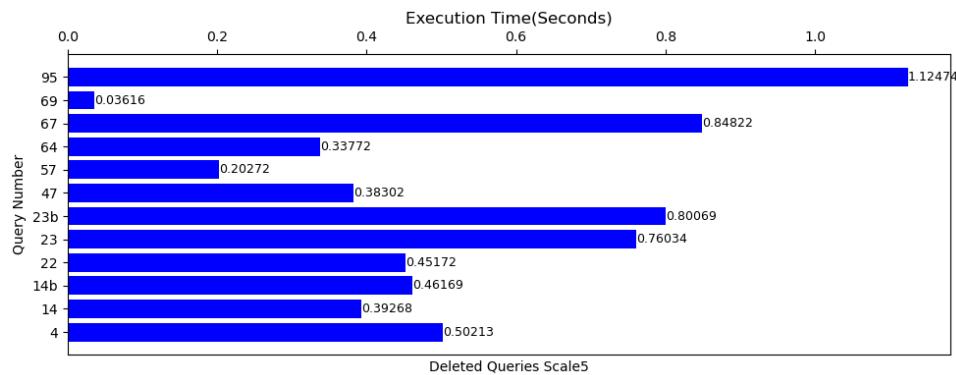


Figure 4.12: Deleted Queries for scale factor 5.

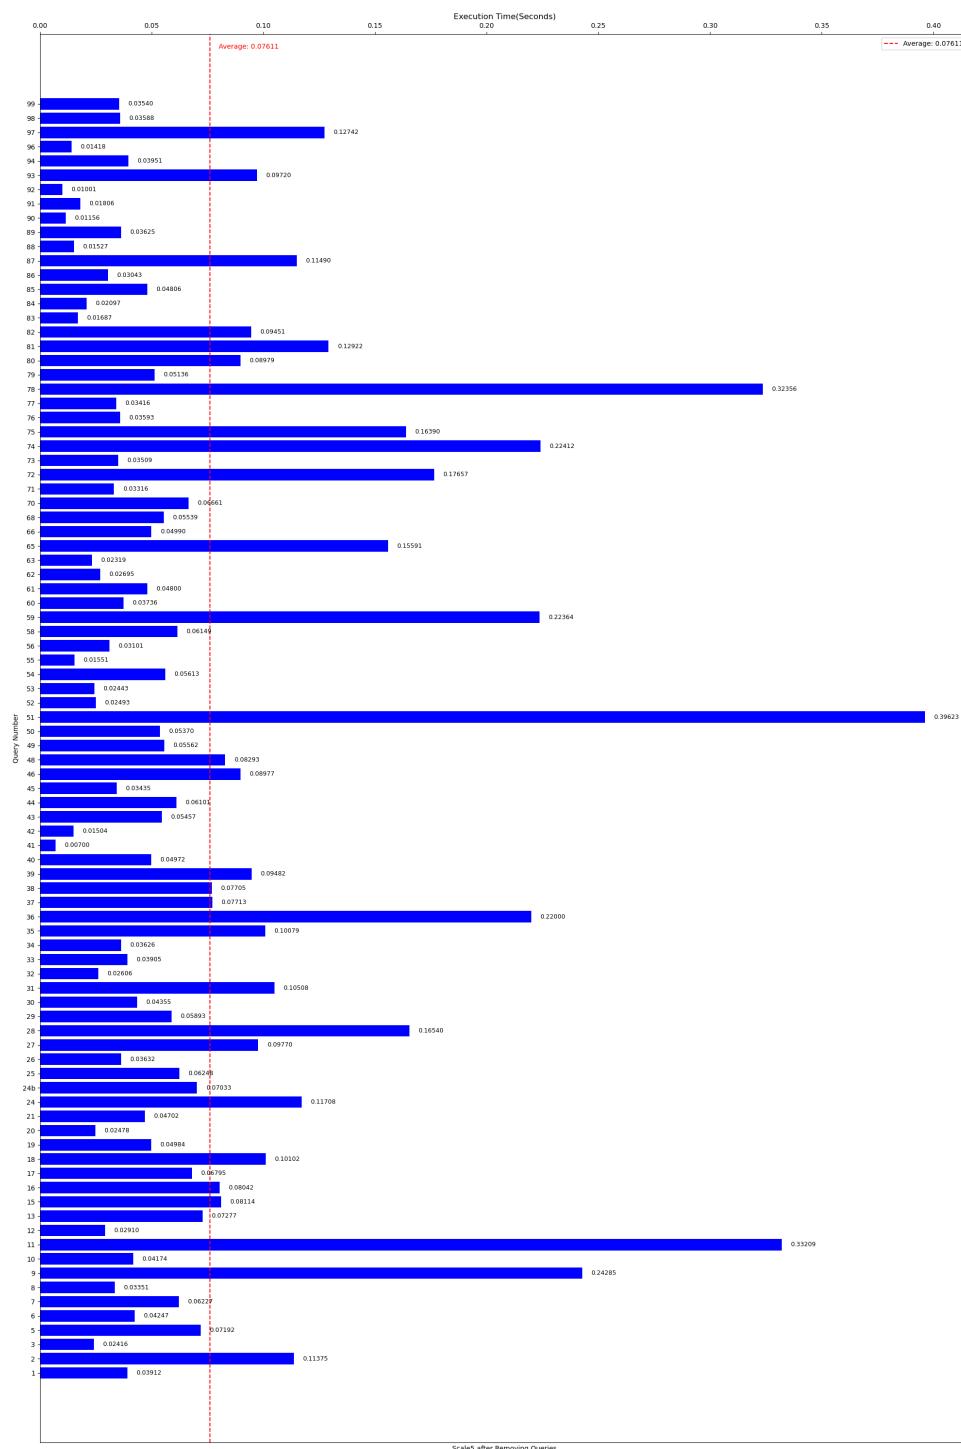


Figure 4.13: After Deleting Queries for scale factor 5.

**Scale 7**

In Figure 4.14, it showed how much time it took to run the 102 queries on scale factor 7. Query 95 is the longest outlier with around 1.77944 seconds of execution times. Query 4,14,14b,22,23,23b,47,57, 64,67,69,95 are longer than other queries. The average times is about 0.12894 seconds for all the queries. So in Figure 4.16, we plotted the execution times of all the queries except for query 4,14,14b,22,23,23b,47,57,64,67,69,95 to see their values more clearly. Query 51 is another longest outlier with around 0.52974 seconds. And the average times is about 0.09893 seconds after removing the queries.

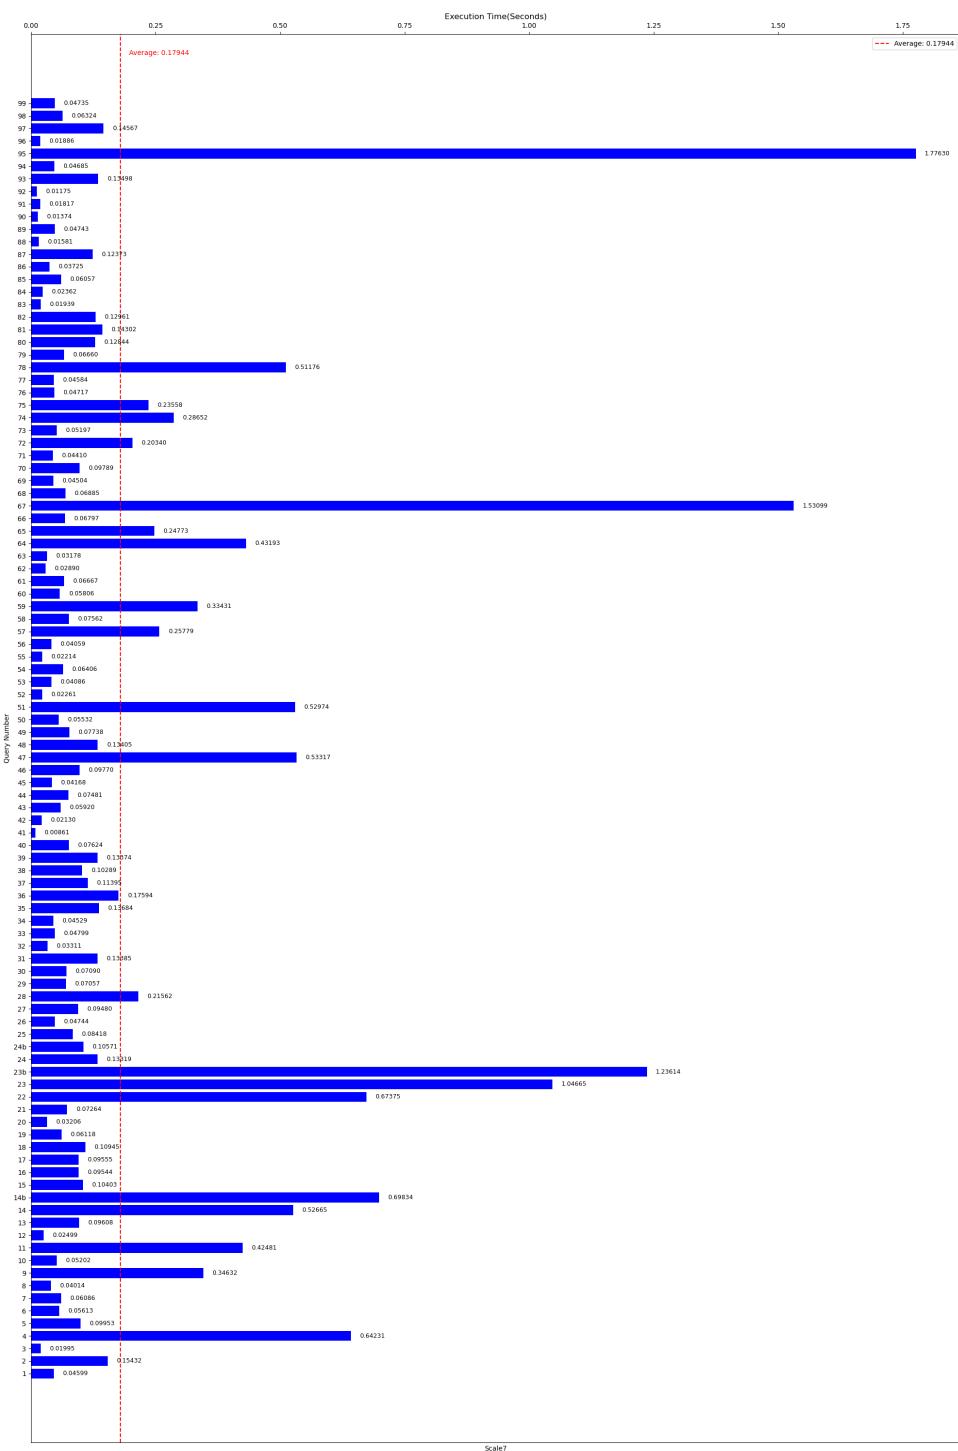


Figure 4.14: Power test result for scale factor 7.



Figure 4.15: Deleted Queries for scale factor 7.

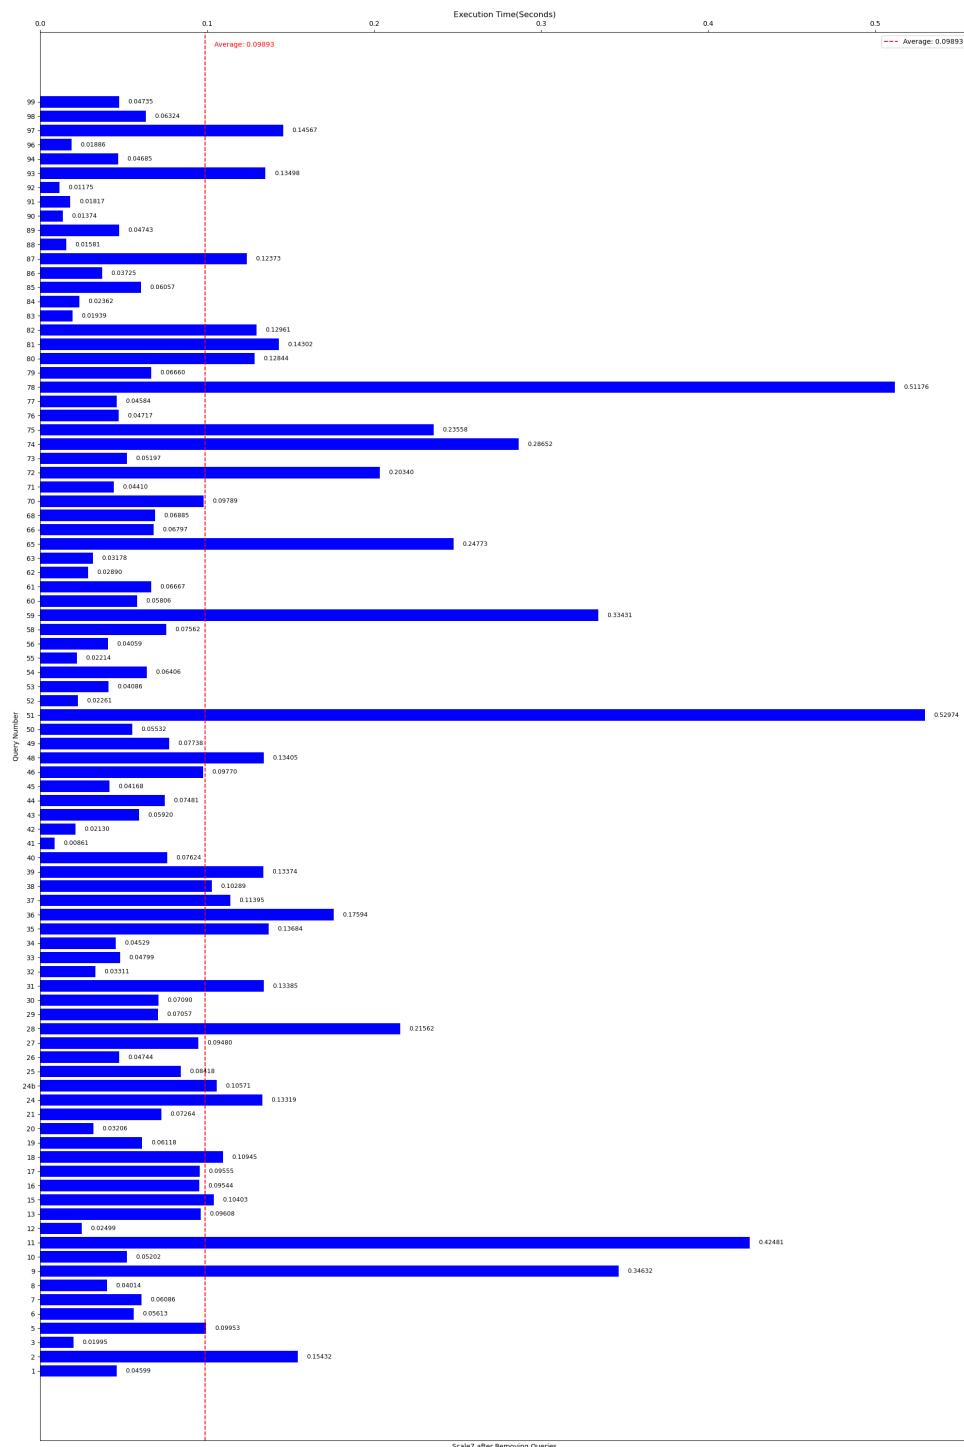


Figure 4.16: After Deleting Queries for scale factor 7.

### 4.2.3 Execution Times in Different Scale Factors

To better demonstrate the relationship between execution times and scale factors, we chose to draw a comparative bar chart of the first ten queries generated. As can be seen from the figure, the execution time of each query grows along with the size factor. And it is worth noting that this growth is close to linear. This is also very much in line with our expectations. Because We thought to choose the scale factors to be 1, 3, 5 and 7 and our scale factors increase linearly. We tried to plotted additional bar graphs for the other groups, and almost all of them showed this linear increase.

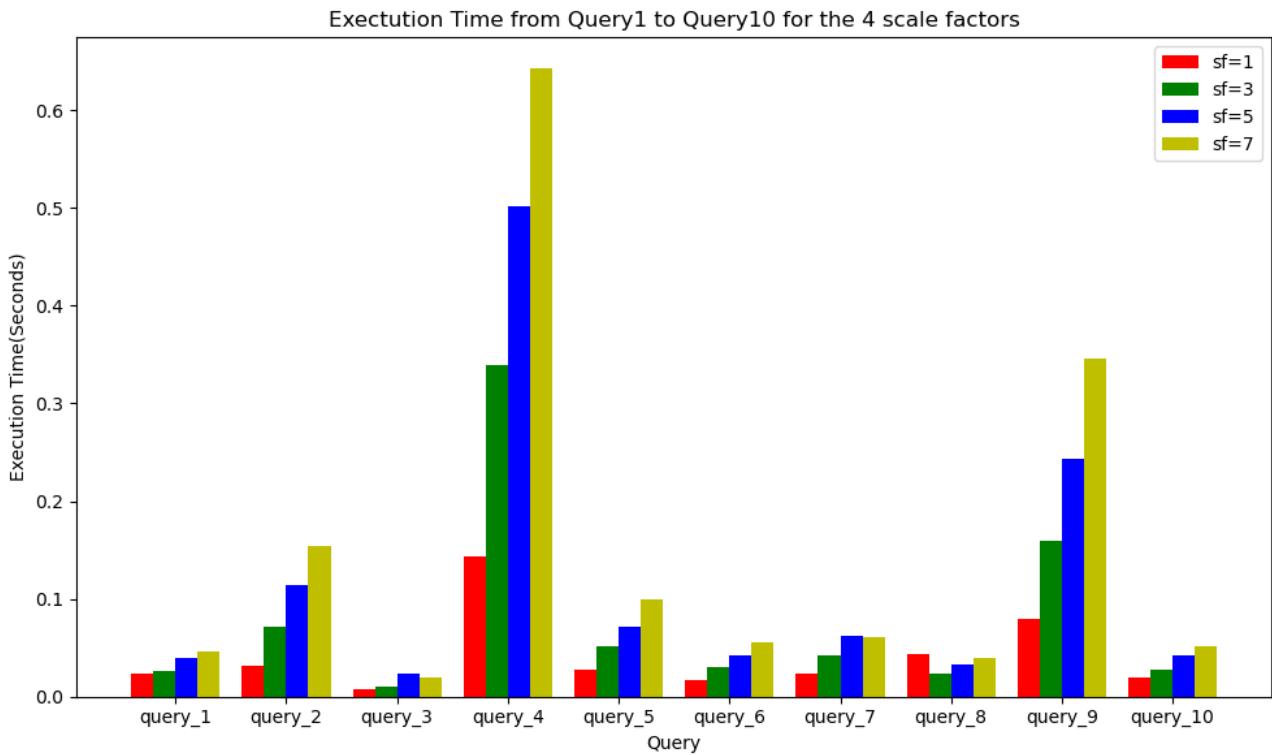


Figure 4.17: Comparison of power test results for the first 10 queries over 4 scale factors.

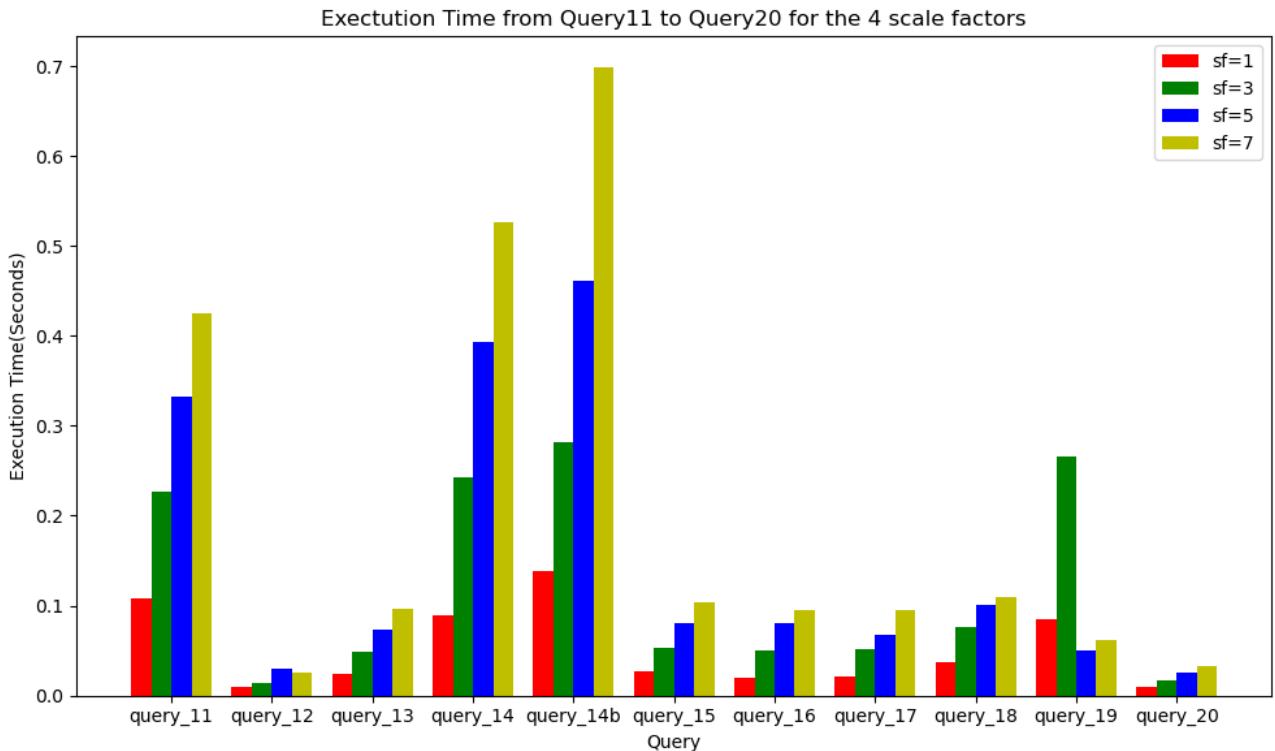


Figure 4.18: Comparison of power test results for the queries 10 to 20 over 4 scale factors.

In addition to this, we encountered a number of queries that were different from the linear growth case. These outliers represent roughly 10% of the overall queries.

We thought the performance of this test does not necessarily increase linearly for several possible reasons:

- Complexity of Queries: TPC-DS queries have varying degrees of complexity. Some queries may involve a lot of data processing and computation, while others may be relatively simple. This means that the execution time may vary significantly from one query to another.
- Resource Competition: In a database system, various resources such as CPU, memory, and storage can become bottlenecks. Some queries may place a higher demand on a particular resource, while other resources may be relatively idle.
- Data and Index Layout: The physical layout of the data and index structure can greatly affect query performance. For example, for a large number of table scanning operations, good data localization can improve performance.
- System Concurrency: Although Power Test evaluates the performance of a single query, there

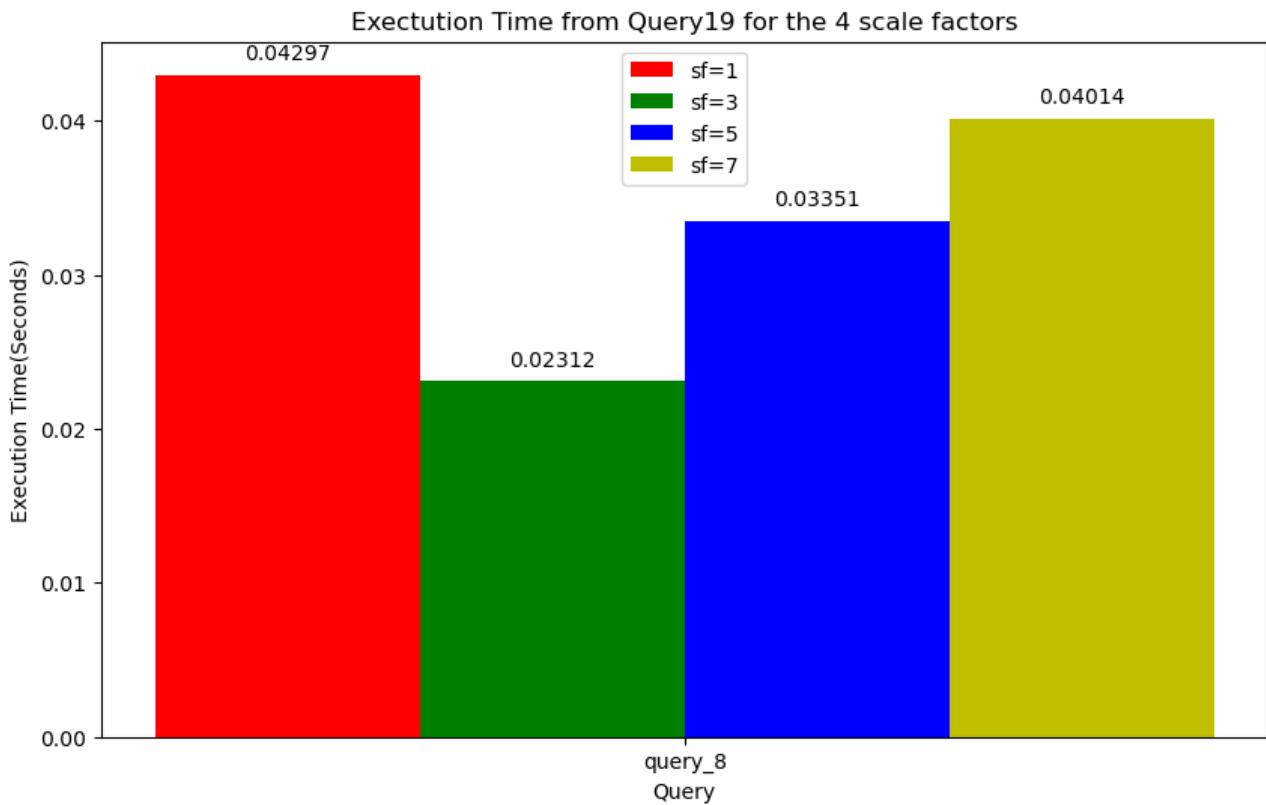


Figure 4.19: Query 8

may be other background tasks and operations in the system that may affect query performance.

In summary, due to these and other possible factors, the performance of TPC-DS Power Test is not guaranteed to be linear. Different queries, data layouts, system configurations, and resource constraints can all lead to non-linear variations in performance.

### 4.3 Throughput Test

Figure 4.21 below shows the results for our implementation of the Throughput Test. During the throughput testing phase, we executed two different refresh run twice. The throughput test 1 was done after the power test while The throughput test 2 was after data maintenance test. Figure 4.21 compares the execution time of throughput 1 and 2 over 4 scale factors. The results of both throughput tests were very close and both showed linear growth. This might indicate that the system is good in multi-user scenarios. Also, after performing the data maintenance test 1, we observed the throughput test maintained same linear trend, showing that even during deletion

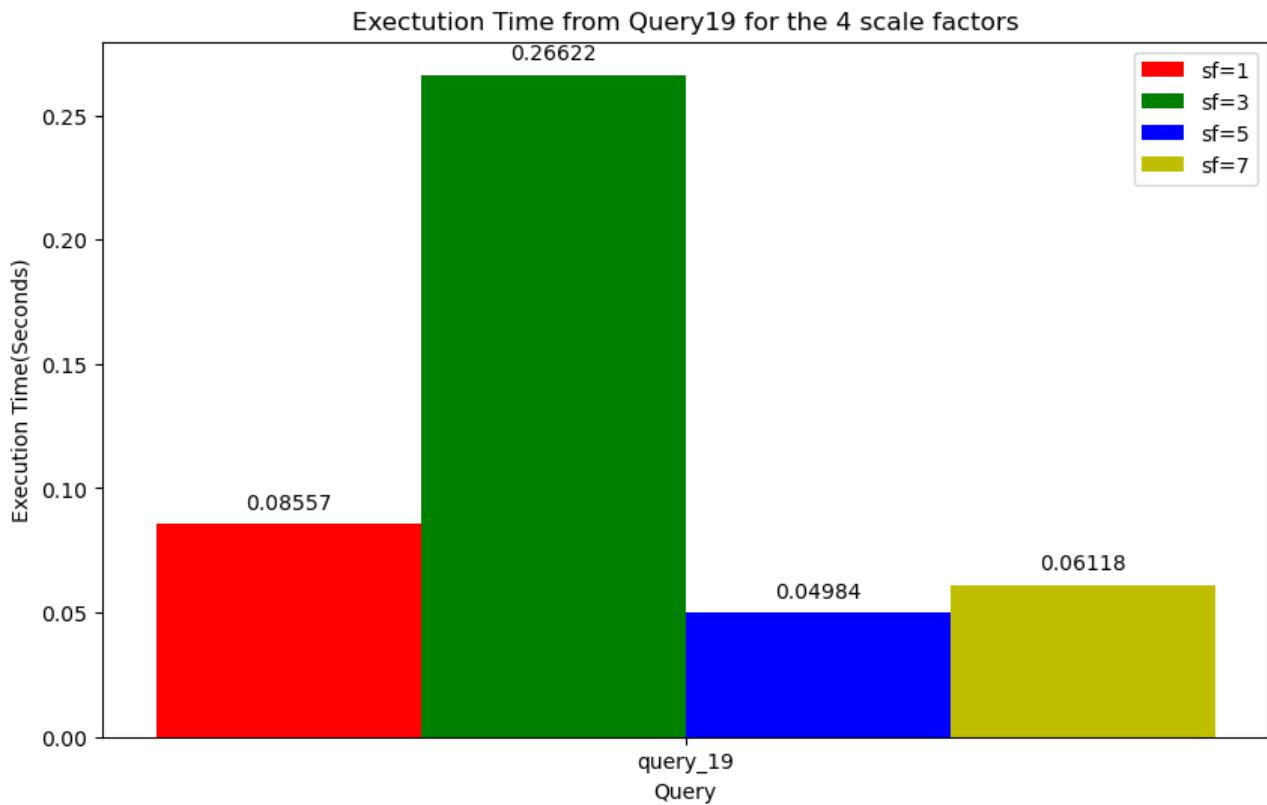


Figure 4.20: Query 19

and insertion, the system can still perform well. We also used 4 different threads having access to the database at the same time to simulate different 4 users. We had named it 'A','B','C','D'.Figure 4.22 shows the throughput test results in 4 scale factors.

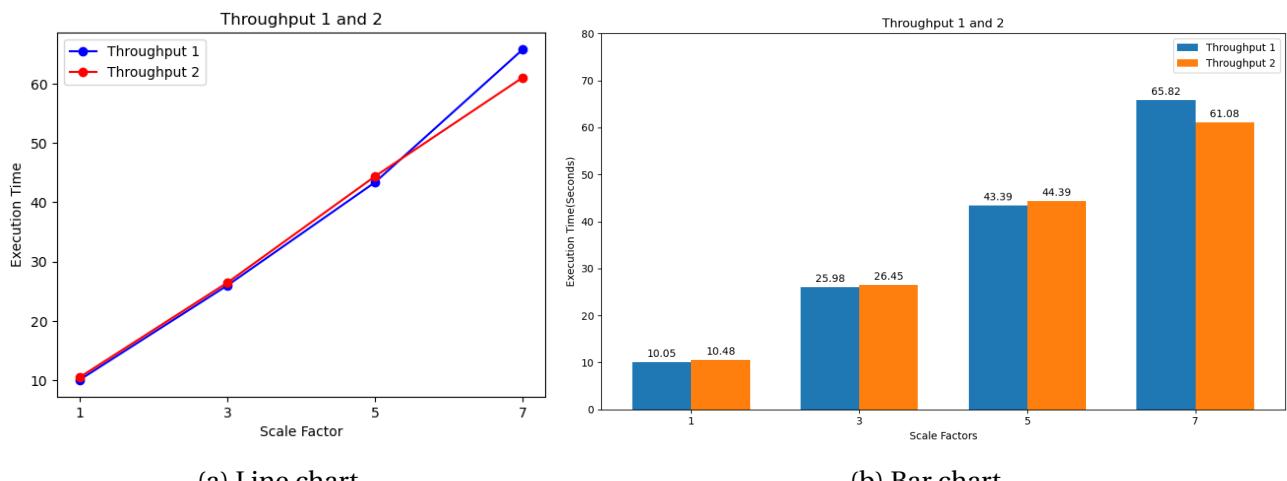


Figure 4.21: Throughput 1 and 2.

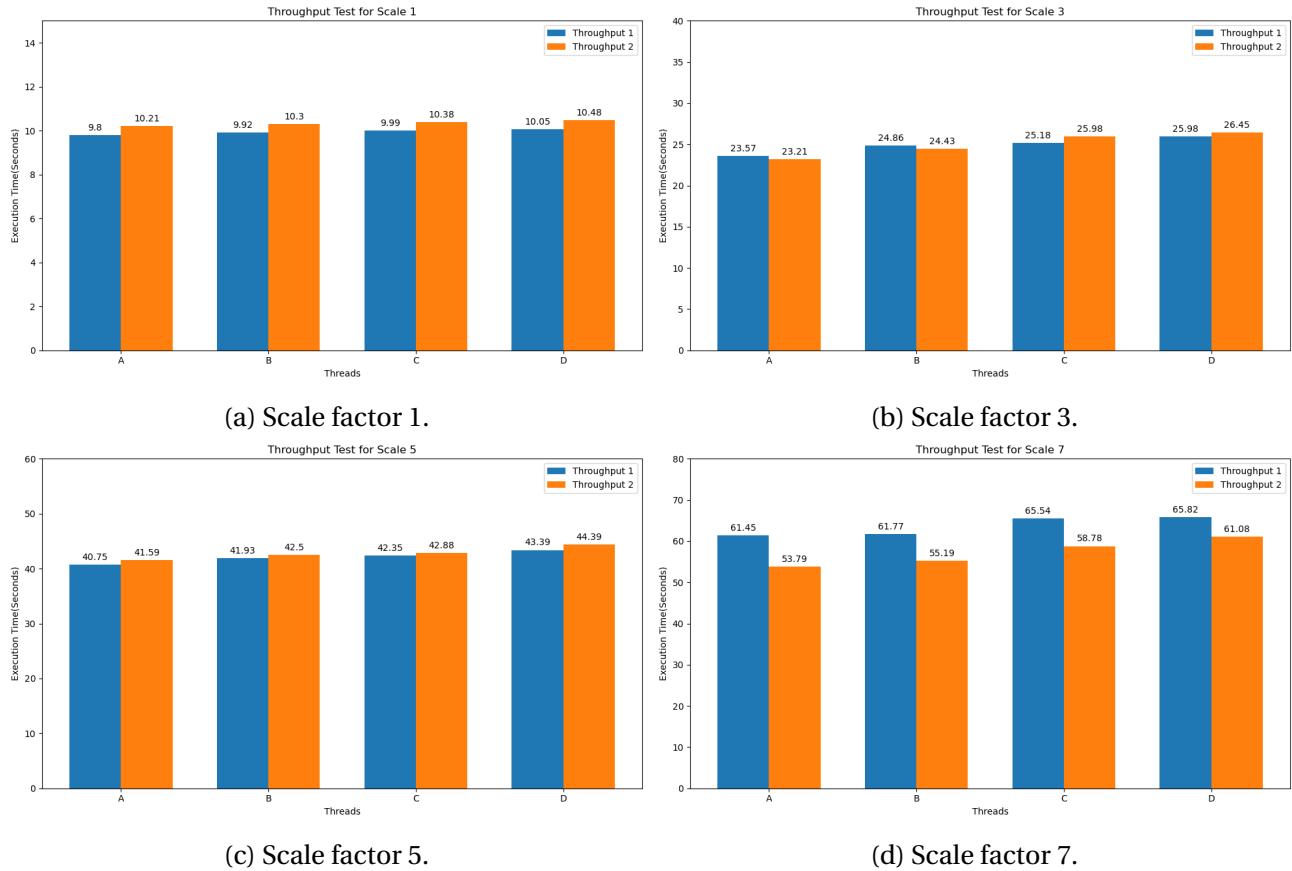


Figure 4.22: Throughput test results.

### 4.3.1 Comparison for Throughput, Power and Load Tests

In Figure 4.23, We examined the execution times for two Throughput Tests and compared them to the Power and Load Tests. At scale factor 1, the duration of the Throughput Test was about twice that of the Power Test, as anticipated. For higher scale factors, however, the Throughput Test's execution time increased more than the Power Test's. This discrepancy is evident in the bar chart, showing a duration increase that's over double but under four times. Such a trend implies our system not only excels in sequential run efficiency but also effectively manages the Throughput Test. We thought the efficiency of multi-user operations needed to be improved.

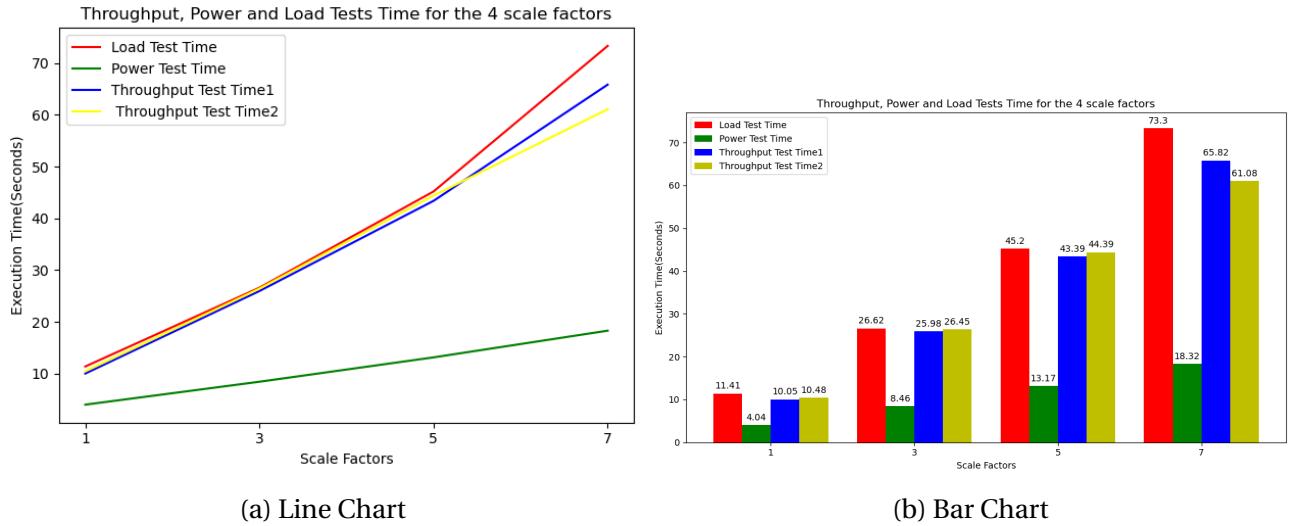


Figure 4.23: Comparison for Throughput, Power and Load Tests Time for the 4 Scale Factors

## 4.4 Data Maintenance Test

Figure 4.24, and 4.25 and 4.26 show the results for the data maintenance test 1, data maintenance test 2 and their different run respectively. It shows the execution time for each maintenance function for all scale factors. It can be noted that the deletion at the store sale fact table and insertion at the inventory fact table took more time because of the data contained in them.

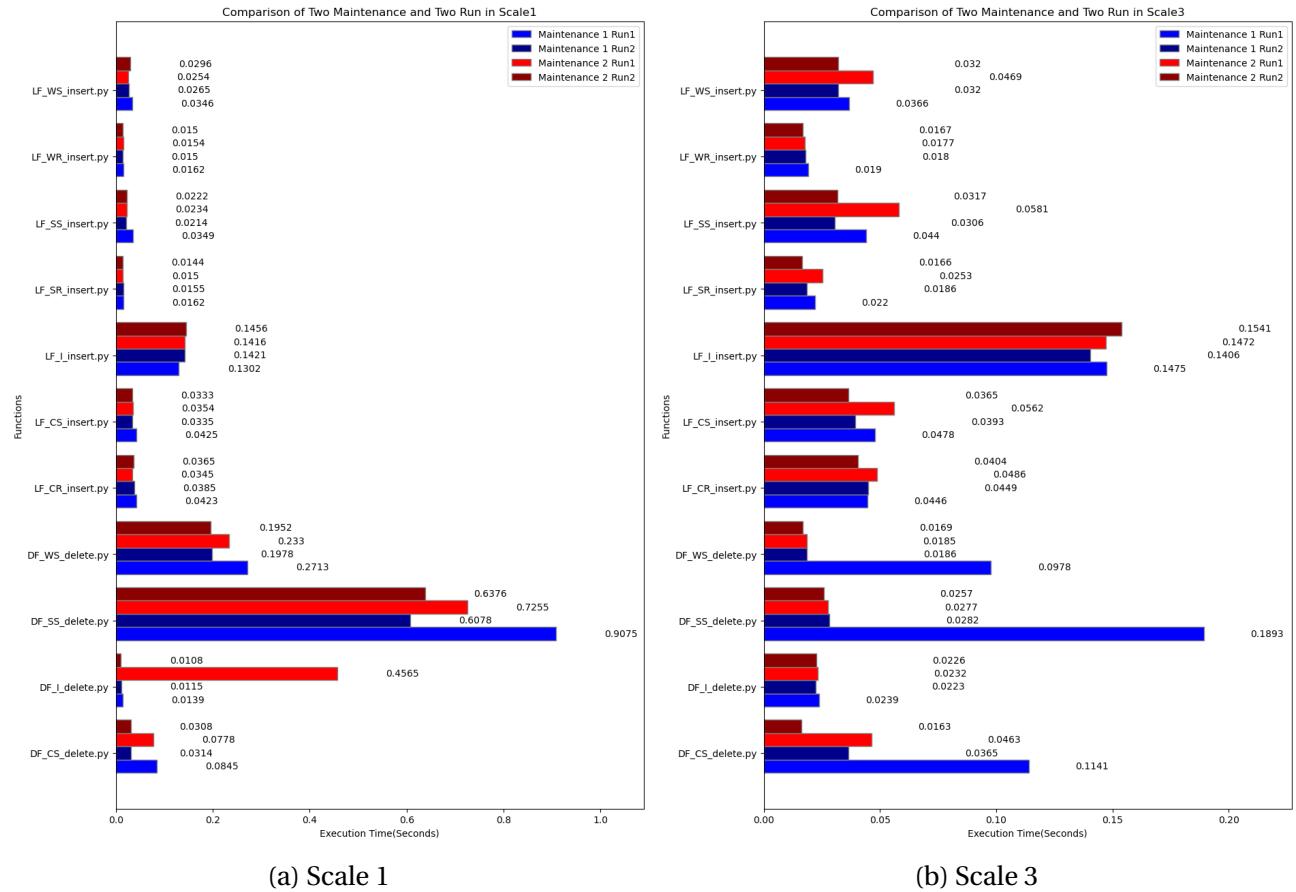


Figure 4.24: Data Maintenance test results.

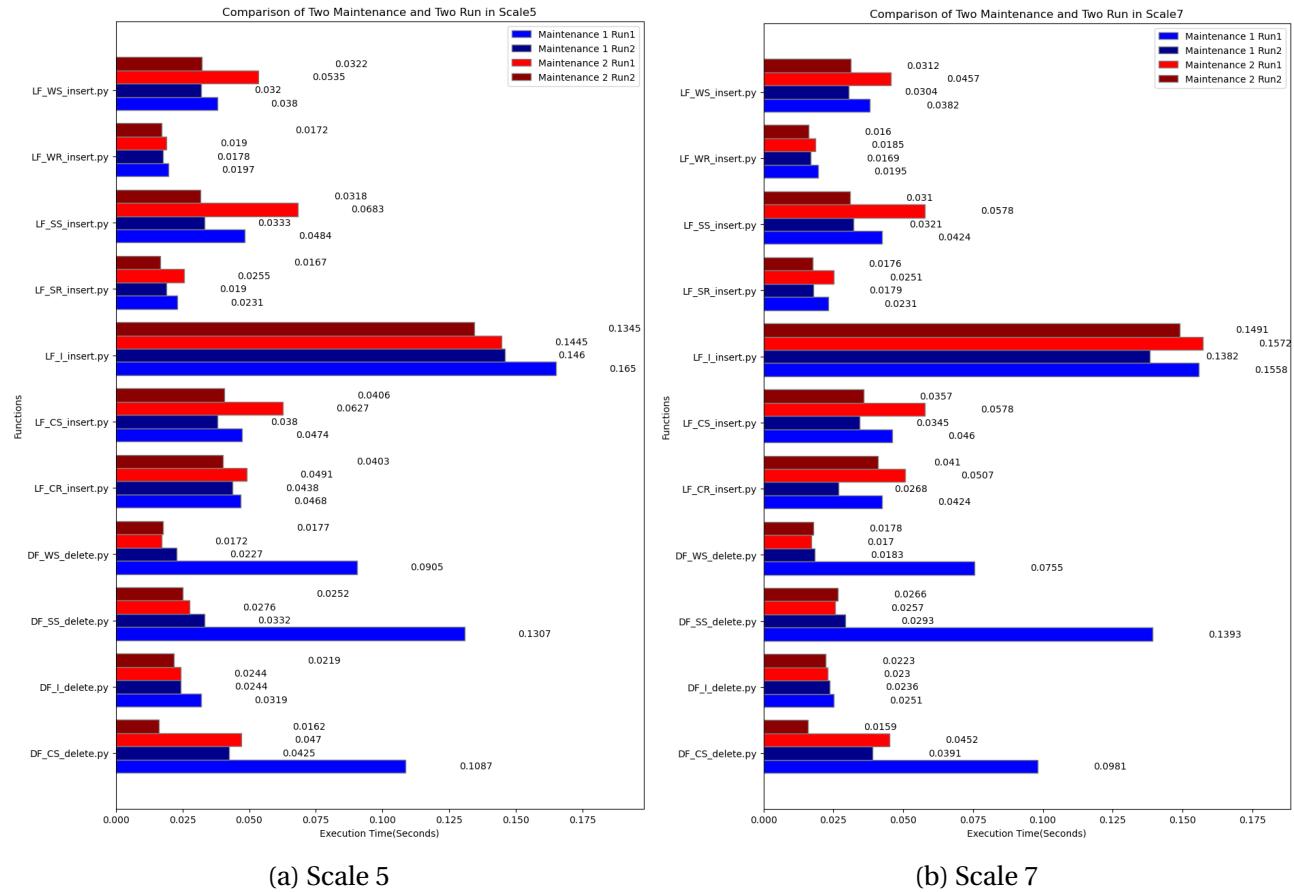


Figure 4.25: Maintenance test results.

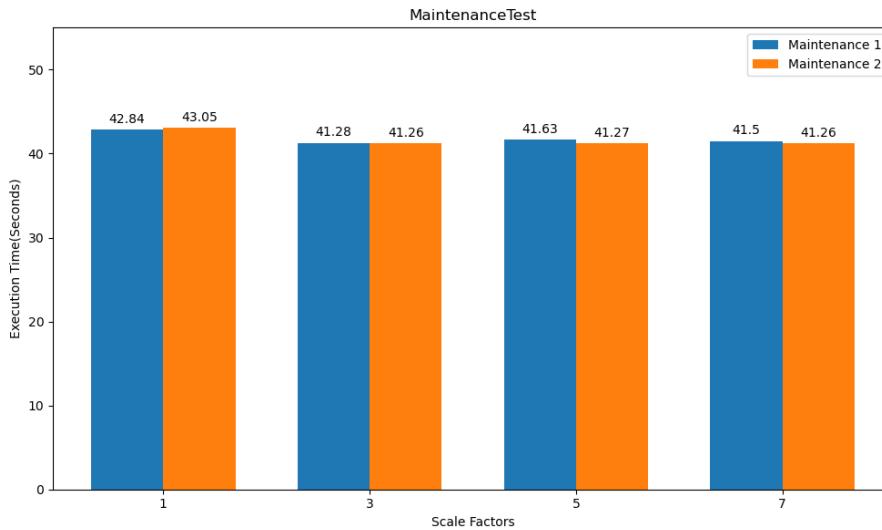


Figure 4.26: Maintenance test

## 4.5 Data Accessibility Test

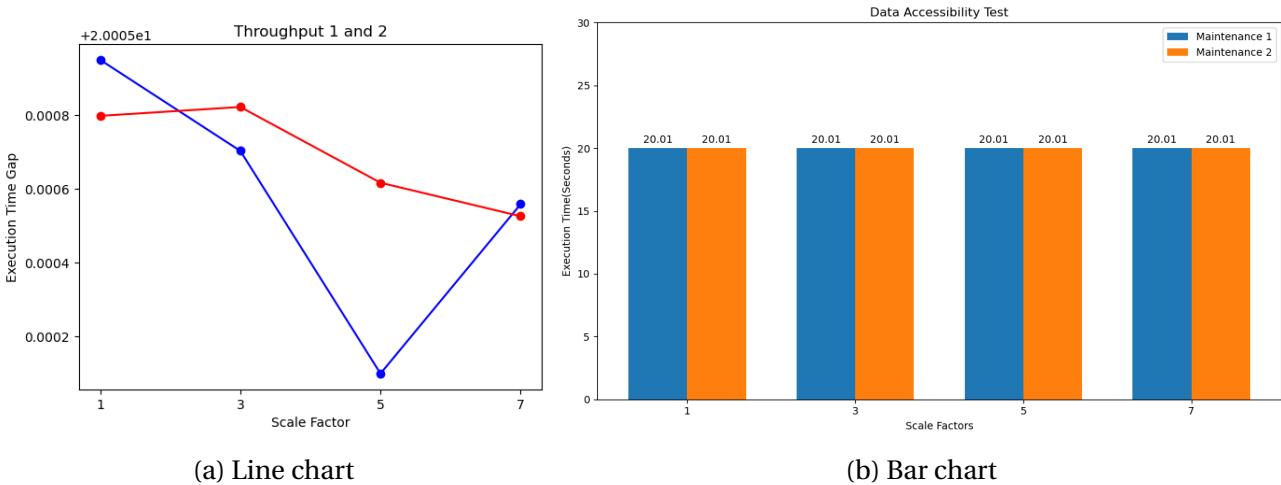


Figure 4.27: Data accessibility test results.

S

## 4.6 Scalability Test

We used Apache jmeter to perform scalability testing on the DuckDB database by increasing the number of users (called thread) and then performing the 99 queries and also running the whole stream query with their stream unique identification (query\_0,query\_1,query\_2,query\_3). It was able to perform well on scale factor 1 as we increased the user from 1 to 20 but when we tried to increase it to 50 users, the system became unresponsive. We did the same for scale factor 3, it became unresponsive for 10 users running the whole queries at once. Same thing applied for Scale factor 5 and 7, they were not able to run with 10 users. We were able to identify the slowest execution queries by filtering the time, they include query 14b, 23b, 23,67, 95 and 14. The following figures below shows the time for the test and the response time for the different scale factor and users:

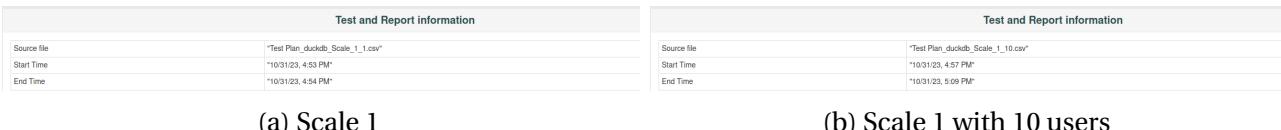


Figure 4.28: Time to perform the queries

Test and Report information													
Source file	"Test Plan_duckdb_Scale_1_20.csv"												
Start Time	"10/31/23, 5:33 PM"												
End Time	"10/31/23, 6:02 PM"												

(a) Scale 1 with 20 users

(b) Scale 3

Figure 4.29: Time to perform the queries

Test and Report information													
Source file	"Test Plan_duckdb_scale_5.csv"												
Start Time	"10/31/23, 10:56 PM"												
End Time	"10/31/23, 11:00 PM"												

(a) Scale 5

(b) Scale 7

Figure 4.30: Time to perform the queries

Statistics																	
Requests Label	#	Executions				Response Times (ms)				Throughput				Network (kB/sec)			
		#Samples	#Fails	#Error %	#	Average	#	Min	#	Median	#	90th pct	#	Transactions	# Received	# Sent	#
Total	109	3	2.75%	681.27	6	14955	87.00	693.00	888.50	14943.60	1.45	11.76	0.00				
JDBC All_query_0	1	0	0.00%	14955.00	14955	14955.00	14955.00	14955.00	14955.00	14955.00	0.07	0.35	0.00				
JDBC All_query_3	1	0	0.00%	14841.00	14841	14841.00	14841.00	14841.00	14841.00	14841.00	0.07	0.08	0.00				
JDBC All_query_4	1	0	0.00%	14841.00	14841	14841.00	14841.00	14841.00	14841.00	14841.00	0.07	0.08	0.00				
JDBC All_query_5	1	0	0.00%	14218.00	14218	14218.00	14218.00	14218.00	14218.00	14218.00	0.07	0.19	0.00				
JDBC Query_1	1	0	0.00%	940.00	940	940.00	940.00	940.00	940.00	940.00	0.07	0.35	0.00				
JDBC Query_14b	1	0	0.00%	940.00	940	940.00	940.00	940.00	940.00	940.00	0.07	0.35	0.00				
JDBC Query_23b	1	0	0.00%	825.00	825	825.00	825.00	825.00	825.00	825.00	0.07	0.02	0.00				
JDBC Query_29	1	0	0.00%	807.00	807	807.00	807.00	807.00	807.00	807.00	0.14	0.14	0.00				
JDBC 67	1	0	0.00%	726.00	726	726.00	726.00	726.00	726.00	726.00	0.13	0.27	0.00				
JDBC 95	1	0	0.00%	657.00	657	657.00	657.00	657.00	657.00	657.00	0.14	0.10	0.00				
JDBC 96	1	0	0.00%	657.00	657	657.00	657.00	657.00	657.00	657.00	0.14	0.81	0.00				
JDBC 97	1	0	0.00%	650.00	650	650.00	650.00	650.00	650.00	650.00	0.14	0.40	0.00				
JDBC 47	1	0	0.00%	506.00	506	506.00	506.00	506.00	506.00	506.00	0.18	19.33	0.00				
JDBC 11	1	0	0.00%	451.00	451	451.00	451.00	451.00	451.00	451.00	0.22	6.69	0.00				
JDBC Query_32	1	0	0.00%	404.00	404	404.00	404.00	404.00	404.00	404.00	0.24	12.19	0.00				
JDBC Show Table	1	0	0.00%	360.00	360	360.00	360.00	360.00	360.00	360.00	0.27	1.65	0.00				
JDBC 20	1	0	0.00%	341.00	341	341.00	341.00	341.00	341.00	341.00	0.27	11.45	0.00				
JDBC 87	1	0	0.00%	331.00	331	331.00	331.00	331.00	331.00	331.00	0.28	37.05	0.00				
JDBC Query_7	1	0	0.00%	281.00	281	281.00	281.00	281.00	281.00	281.00	0.35	14.14	0.00				
JDBC 79	1	0	0.00%	273.00	273	273.00	273.00	273.00	273.00	273.00	0.36	20.71	0.00				
JDBC 79	1	0	0.00%	253.00	253	253.00	253.00	253.00	253.00	253.00	0.35	62.22	0.00				
JDBC Query_19	1	0	0.00%	228.00	228	228.00	228.00	228.00	228.00	228.00	0.49	21.88	0.00				
JDBC 72	1	0	0.00%	217.00	217	217.00	217.00	217.00	217.00	217.00	0.46	54.60	0.00				
JDBC 75	1	0	0.00%	210.00	210	210.00	210.00	210.00	210.00	210.00	0.46	47.40	0.00				

(a) Scale 1 with 20 users

(b) Scale 1 with 10 users

Figure 4.31: Statistics of Queries

Statistics																	
Requests Label	#	Executions				Response Times (ms)				Throughput				Network (kB/sec)			
		#Samples	#Fails	#Error %	#	Average	#	Min	#	Median	#	90th pct	#	Transactions	# Received	# Sent	#
Total	2180	60	2.75%	1397.56	4	42652	1342.00	1676.70	36555.00	34926.50	1.27	10.33	0.00				
JDBC All_query_0	20	0	0.00%	34174.25	288059	42652	540261.50	395912.00	424782.45	42652	0.03	0.15	0.00				
JDBC All_query_1	20	0	0.00%	332740.00	285014	403452	534422.00	382055.40	403452	403452	0.03	0.09	0.00				
JDBC All_query_2	20	0	0.00%	330740.00	285014	403452	534422.00	382055.40	403452	403452	0.03	0.09	0.00				
JDBC All_query_3	20	0	0.00%	287773.40	162549	308503	295251.00	341106.10	362835.00	308503.00	0.04	0.05	0.00				
JDBC Query_29	20	0	0.00%	20105.15	19129	20105.15	19129	19129	19129	19129	0.17	0.02	0.00				
JDBC Query_23b	20	0	0.00%	18107.00	10324	18107.00	18107.00	18107.00	18107.00	18107.00	0.15	0.00	0.00				
JDBC Query_67	20	0	0.00%	1608.00	11667	1608.00	1608.00	1608.00	1608.00	1608.00	0.14	0.92	0.00				
JDBC 95	20	0	0.00%	1545.00	10563	19996	13204.00	19452.00	19875.15	19996.00	0.17	0.65	0.00				
JDBC Query_14b	20	0	0.00%	2182.80	6954	2081.00	25533.00	31310.80	35696.05	25533.00	0.24	0.81	0.00				
JDBC 4	20	0	0.00%	2184.35	5148	31150.20	27500.00	30500.30	31164.20	31164.20	0.29	0.15	0.00				
JDBC 64	20	0	0.00%	17035.25	9250	25448	17800.00	21500.00	25448.00	17800.00	0.14	0.05	0.00				
JDBC 47	20	0	0.00%	1665.55	7031	29707	17420.00	21747.10	28893.15	29707.00	0.15	1.50	0.00				
JDBC Query_14	20	0	0.00%	14605.95	4421	23104	16830.00	22449.00	23076.00	231010.00	0.32	1.98	0.00				
JDBC Query_11	20	0	0.00%	13895.70	3336	22469	14277.00	19378.80	22164.90	22469.00	0.38	1.14	0.00				
JDBC 51	20	0	0.00%	12895.30	6963	19996	13204.00	19452.00	19875.15	19996.00	0.17	0.65	0.00				
JDBC 57	20	0	0.00%	12432.00	3909	16577	9568.00	15351.60	16454.35	16577.00	0.16	1.93	0.00				
JDBC 58	20	0	0.00%	12300.00	3909	16577	9568.00	15351.60	16454.35	16577.00	0.17	1.93	0.00				
JDBC Query_19	20	0	0.00%	1222.05	2830	11075	6300.00	8047.00	11075.00	11245.00	0.24	1.22	0.00				
JDBC 59	20	0	0.00%	120139	2884	10401	5425.00	8490.00	102335.00	10401	0.29	1.24	0.00				
JDBC 75	20	0	0.00%	4201.45	2028	7171	3916.00	4775.90	7154.40	7171.00	0.14	0.72	0.00				
JDBC 72	20	0	0.00%	4205.00	2201	8103	4097.00	5050.00	8103.00	8103.00	0.14	1.72	0.00	</td			

Statistics																									
Requests	Label	Executions					Response Times (ms)					Throughput			Network (Mbps)										
		#	iSamples	#	FAIL	%	Error %	#	Average	#	Min	#	Max	#	Median	#	90th pct	#	99th pct	#	Transactions	#	Received	#	Sent
Total	109	2	1.82%	1874.76	15	39199	148.00	1567.00	17944.00	39180.00	0.51	7.01	0.00												
JDBC_Query_41	1	0	0.00%	15.00	15.00	15.00	0.00%	15.00	15.00	15.00	15.00	0.00	15.00	15.00	15.00	15.00	15.00	15.00	15.00	15.00	15.00	15.00	15.00	15.00	
JDBC_90	1	0	0.00%	20.00	20	20.00	20.00	20.00	20.00	20.00	20.00	0.00	20.00	20.00	20.00	20.00	20.00	20.00	20.00	20.00	20.00	20.00	20.00	20.00	
JDBC_42	1	0	0.00%	26.00	26	26.00	26.00	26.00	26.00	26.00	26.00	0.00	26.00	26.00	26.00	26.00	26.00	26.00	26.00	26.00	26.00	26.00	26.00	26.00	
JDBC_Query_27	1	0	0.00%	27.00	27	27.00	27.00	27.00	27.00	27.00	27.00	0.00	27.00	27.00	27.00	27.00	27.00	27.00	27.00	27.00	27.00	27.00	27.00	27.00	
JDBC_Query_42	1	0	0.00%	27.00	27	27.00	27.00	27.00	27.00	27.00	27.00	0.00	27.00	27.00	27.00	27.00	27.00	27.00	27.00	27.00	27.00	27.00	27.00	27.00	
JDBC_95	1	0	0.00%	28.00	28	28.00	28.00	28.00	28.00	28.00	28.00	0.00	28.00	28.00	28.00	28.00	28.00	28.00	28.00	28.00	28.00	28.00	28.00	28.00	
JDBC_88	1	0	0.00%	37.00	37	37.00	37.00	37.00	37.00	37.00	37.00	0.00	37.00	37.00	37.00	37.00	37.00	37.00	37.00	37.00	37.00	37.00	37.00	37.00	
JDBC_86	1	0	0.00%	37.00	37	37.00	37.00	37.00	37.00	37.00	37.00	0.00	37.00	37.00	37.00	37.00	37.00	37.00	37.00	37.00	37.00	37.00	37.00	37.00	
JDBC_85	1	0	0.00%	37.00	37	37.00	37.00	37.00	37.00	37.00	37.00	0.00	37.00	37.00	37.00	37.00	37.00	37.00	37.00	37.00	37.00	37.00	37.00	37.00	
JDBC_91	1	0	0.00%	40.00	40	40.00	40.00	40.00	40.00	40.00	40.00	0.00	40.00	40.00	40.00	40.00	40.00	40.00	40.00	40.00	40.00	40.00	40.00	40.00	
JDBC_54	1	0	0.00%	42.00	42	42.00	42.00	42.00	42.00	42.00	42.00	0.00	42.00	42.00	42.00	42.00	42.00	42.00	42.00	42.00	42.00	42.00	42.00	42.00	
JDBC_Query_20	1	0	0.00%	44.00	44	44.00	44.00	44.00	44.00	44.00	44.00	0.00	44.00	44.00	44.00	44.00	44.00	44.00	44.00	44.00	44.00	44.00	44.00	44.00	
Insertion	1	1	100.00%	45.00	45	45.00	45.00	45.00	45.00	45.00	45.00	0.00	45.00	45.00	45.00	45.00	45.00	45.00	45.00	45.00	45.00	45.00	45.00	45.00	
JDBC_03	1	0	0.00%	46.00	46	46.00	46.00	46.00	46.00	46.00	46.00	0.00	46.00	46.00	46.00	46.00	46.00	46.00	46.00	46.00	46.00	46.00	46.00	46.00	
JDBC_03	1	0	0.00%	47.00	47	47.00	47.00	47.00	47.00	47.00	47.00	0.00	47.00	47.00	47.00	47.00	47.00	47.00	47.00	47.00	47.00	47.00	47.00	47.00	
JDBC_52	1	0	0.00%	50.00	50	50.00	50.00	50.00	50.00	50.00	50.00	0.00	50.00	50.00	50.00	50.00	50.00	50.00	50.00	50.00	50.00	50.00	50.00	50.00	
JDBC_62	1	0	0.00%	52.00	52	52.00	52.00	52.00	52.00	52.00	52.00	0.00	52.00	52.00	52.00	52.00	52.00	52.00	52.00	52.00	52.00	52.00	52.00	52.00	
JDBC_84	1	0	0.00%	53.00	53	53.00	53.00	53.00	53.00	53.00	53.00	0.00	53.00	53.00	53.00	53.00	53.00	53.00	53.00	53.00	53.00	53.00	53.00	53.00	
JDBC_84	1	0	0.00%	53.00	53	53.00	53.00	53.00	53.00	53.00	53.00	0.00	53.00	53.00	53.00	53.00	53.00	53.00	53.00	53.00	53.00	53.00	53.00	53.00	
JDBC_Query_12	1	0	0.00%	60.00	60	60.00	60.00	60.00	60.00	60.00	60.00	0.00	60.00	60.00	60.00	60.00	60.00	60.00	60.00	60.00	60.00	60.00	60.00	60.00	
JDBC_53	1	0	0.00%	61.00	61	61.00	61.00	61.00	61.00	61.00	61.00	0.00	61.00	61.00	61.00	61.00	61.00	61.00	61.00	61.00	61.00	61.00	61.00	61.00	
JDBC_Query_33	1	0	0.00%	65.00	65	65.00	65.00	65.00	65.00	65.00	65.00	0.00	65.00	65.00	65.00	65.00	65.00	65.00	65.00	65.00	65.00	65.00	65.00	65.00	
JDBC_71	1	0	0.00%	66.00	66	66.00	66.00	66.00	66.00	66.00	66.00	0.00	66.00	66.00	66.00	66.00	66.00	66.00	66.00	66.00	66.00	66.00	66.00	66.00	
JDBC_73	1	0	0.00%	66.00	66	66.00	66.00	66.00	66.00	66.00	66.00	0.00	66.00	66.00	66.00	66.00	66.00	66.00	66.00	66.00	66.00	66.00	66.00	66.00	

(a) Scale 5

(b) Scale 7

Figure 4.33: Statistics of Queries

## 4.7 Optimization

Considering the results from the low performance queries, optimization were made. The optimized queries are detailed in Appendix A. The optimization were done using the following:

- The Query plans were studied to analyze possible points of optimization.
- Indexes were added for fast retrieval of table tuples.
- Creating of TEMPORARY table to store and manipulate results

### Query 95

Query 95 that has the question, *Produce a count of web sales and total shipping cost and net profit in a given 60 day period to customers in a given state from a named web site for returned orders shipped from more than one warehouse.*

Qualification Substitution Parameters:

- STATE.01=IL
- MONTH.01=2
- YEAR.01=1999

The original Query was:

```

1      with ws_wh as
2 (select ws1.ws_order_number,ws1.ws_warehouse_sk wh1,ws2.ws_warehouse_sk wh2
3 from web_sales ws1,web_sales ws2
4 where ws1.ws_order_number = ws2.ws_order_number
5   and ws1.ws_warehouse_sk <> ws2.ws_warehouse_sk)
6 select
7   count(distinct ws_order_number) as "order count"
8 ,sum(ws_ext_ship_cost) as "total shipping cost"
9 ,sum(ws_net_profit) as "total net profit"
10 from
11   web_sales ws1
12 ,date_dim
13 ,customer_address
14 ,web_site
15 where
16   d_date between '1999-02-01' and
17         (cast('1999-02-01' as date) + interval 60 days)
18 and ws1.ws_ship_date_sk = d_date_sk
19 and ws1.ws_ship_addr_sk = ca_address_sk
20 and ca_state = 'IL'
21 and ws1.ws_web_site_sk = web_site_sk
22 and web_company_name = 'pri'
23 and ws1.ws_order_number in (select ws_order_number
24                           from ws_wh)
25 and ws1.ws_order_number in (select wr_order_number
26                           from web_returns,ws_wh
27                           where wr_order_number = ws_wh.ws_order_number)
28 order by count(distinct ws_order_number)
29 limit 100;

```

Listing 4.1: Query 95.

This query involves joins on several tables (web\_sales, date\_dim, customer\_address and web\_site), and filtering conditions to retrieve count order of web sales and total shipping cost and net profit. The following was done to optimize this query

- creating indexes,
- using CTEs, and

- optimizing the structure of the SQL statement

```

1   CREATE INDEX if not exists idx_web_sales_order_number ON web_sales(
2     ws_order_number);
3
4 -- Create an index on web_returns to improve join performance
5 CREATE INDEX if not exists idx_web_returns_order_number ON web_returns(
6   wr_order_number);
7
8 -- Use CTEs for better readability and performance
9 WITH OrderWarehouses AS (
10   SELECT ws1.ws_order_number, ws1.ws_warehouse_sk AS wh1, ws2.
11     ws_warehouse_sk AS wh2
12   FROM web_sales ws1
13   JOIN web_sales ws2 ON ws1.ws_order_number = ws2.ws_order_number
14   WHERE ws1.ws_warehouse_sk <> ws2.ws_warehouse_sk
15 ),
16 FilteredOrders AS (
17   SELECT DISTINCT ws_order_number
18   FROM web_sales
19   WHERE ws_order_number IN (SELECT ws_order_number FROM OrderWarehouses)
20 ),
21 FilteredReturns AS (
22   SELECT DISTINCT wr_order_number
23   FROM web_returns wr
24   JOIN OrderWarehouses ow ON wr.wr_order_number = ow.ws_order_number
25 ),
26 DateRange AS (
27   SELECT d_date_sk
28   FROM date_dim
29   WHERE d_date BETWEEN '1999-02-01' AND ('1999-02-01'::DATE + INTERVAL '60
30   days')
31 )
32
33 -- Retrieve results
34
35 SELECT
36   COUNT(DISTINCT ws.ws_order_number) AS "order count",
37   SUM(ws.ws_ext_ship_cost) AS "total shipping cost",
38   SUM(ws.ws_net_profit) AS "total net profit"
39
40 FROM web_sales ws
41 JOIN DateRange dr ON ws.ws_ship_date_sk = dr.d_date_sk
42 JOIN customer_address ca ON ws.ws_ship_addr_sk = ca.ca_address_sk

```

```

36 JOIN web_site ws2 ON ws.ws_web_site_sk = ws2.web_site_sk
37 WHERE ca.ca_state = 'IL'
38     AND ws2.web_company_name = 'pri'
39     AND ws.ws_order_number IN (SELECT * FROM FilteredOrders)
40     AND ws.ws_order_number IN (SELECT * FROM FilteredReturns)
41 GROUP BY dr.d_date_sk
42 ORDER BY "order count" DESC
43 LIMIT 100;

```

Listing 4.2: Optimized Query 95.

As a result, testing these changes in SF1, the Query 95 went from 0.1328871250152588 seconds to 0.0005030632019042969 seconds, SF-3 from 0.48744726181030273 seconds to 0.0005521774291992188 seconds, SF-5 from 0.7921462059020996 seconds to 0.0007989406585693359 seconds and SF-7 from 1.0165119171142578 seconds to 0.0007801055908203125 seconds. The figure below shows the result. They were so far apart by orders of magnitude that we have to use a logarithmic scale on the y-axis.

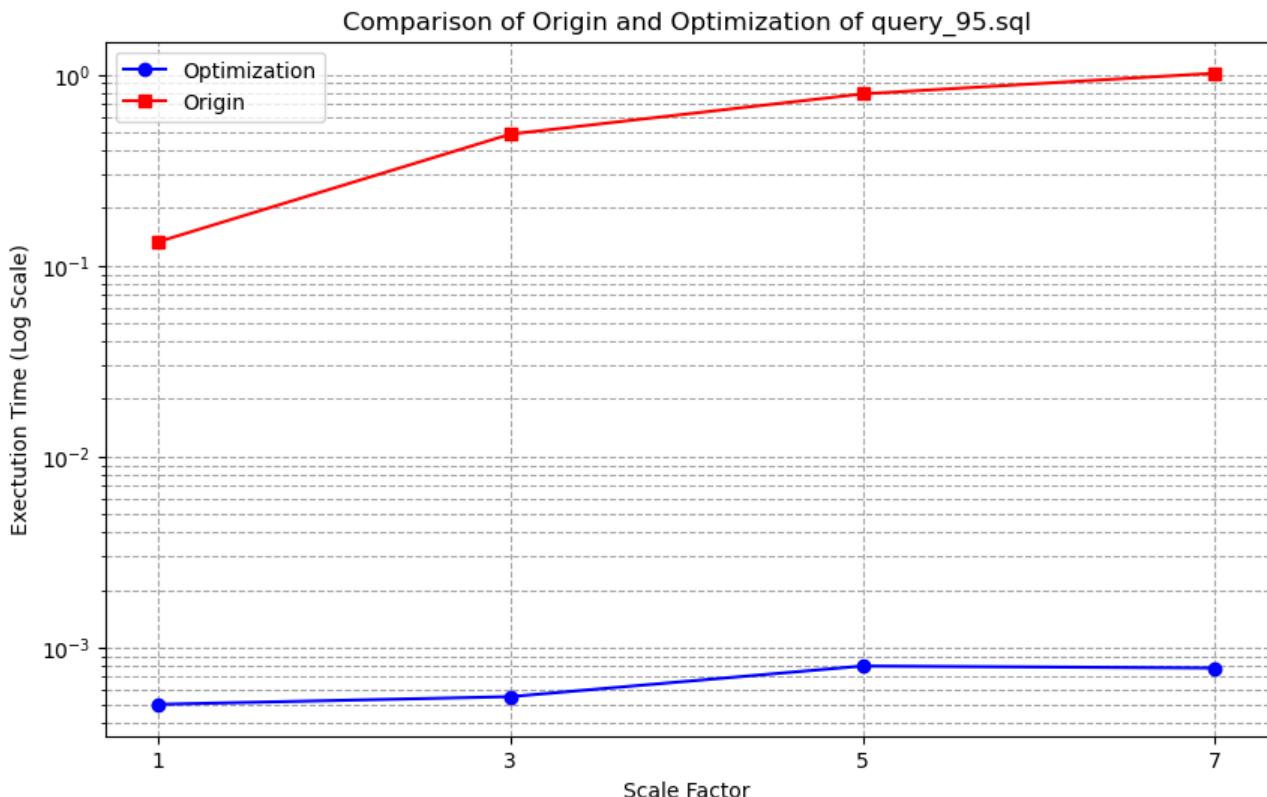


Figure 4.34: Optimization/Comparison of Origin and Optimization of query 95

## Query 67

Query 67 has the question *Find top stores for each category based on store sales in a specific year.*

Qualification Substitution Parameters:

- DMS.01 = 1200

The original Query was:

```

1      select  *
2 from (select i_category
3           ,i_class
4           ,i_brand
5           ,i_product_name
6           ,d_year
7           ,d_qoy
8           ,d_moy
9           ,s_store_id
10          ,sumsales
11          ,rank() over (partition by i_category order by sumsales desc) rk
12 from (select i_category
13           ,i_class
14           ,i_brand
15           ,i_product_name
16           ,d_year
17           ,d_qoy
18           ,d_moy
19           ,s_store_id
20           ,sum(coalesce(ss_sales_price*ss_quantity,0)) sumsales
21     from store_sales
22           ,date_dim
23           ,store
24           ,item
25   where ss_sold_date_sk=d_date_sk
26     and ss_item_sk=i_item_sk
27     and ss_store_sk = s_store_sk
28     and d_month_seq between 1200 and 1200+11
29   group by rollup(i_category, i_class, i_brand, i_product_name, d_year,
30             d_qoy, d_moy,s_store_id))dw1) dw2
30 where rk <= 100

```

```

31 order by i_category
32     ,i_class
33     ,i_brand
34     ,i_product_name
35     ,d_year
36     ,d_qoy
37     ,d_moy
38     ,s_store_id
39     ,sumsales
40     ,rk
41 limit 100;

```

Listing 4.3: Query 67.

This query was optimized by the use of indexing. Index were added on items, date\_dim, store\_sales and store.

```

1 --Indexes for store_sales
2 CREATE INDEX idx_store_sales_sold_date ON store_sales(ss_sold_date_sk);
3 CREATE INDEX idx_store_sales_item ON store_sales(ss_item_sk);
4 CREATE INDEX idx_store_sales_store ON store_sales(ss_store_sk);

5
6 --Indexes for item
7 CREATE INDEX idx_item_item_sk ON item(i_item_sk);

8
9 --Indexes for store
10 CREATE INDEX idx_store_store_sk ON store(s_store_sk);

11
12 --Indexes for date_dim
13 CREATE INDEX idx_date_dim_date_sk ON date_dim(d_date_sk);

```

Listing 4.4: Optimized Query 67.

The result is shown below:

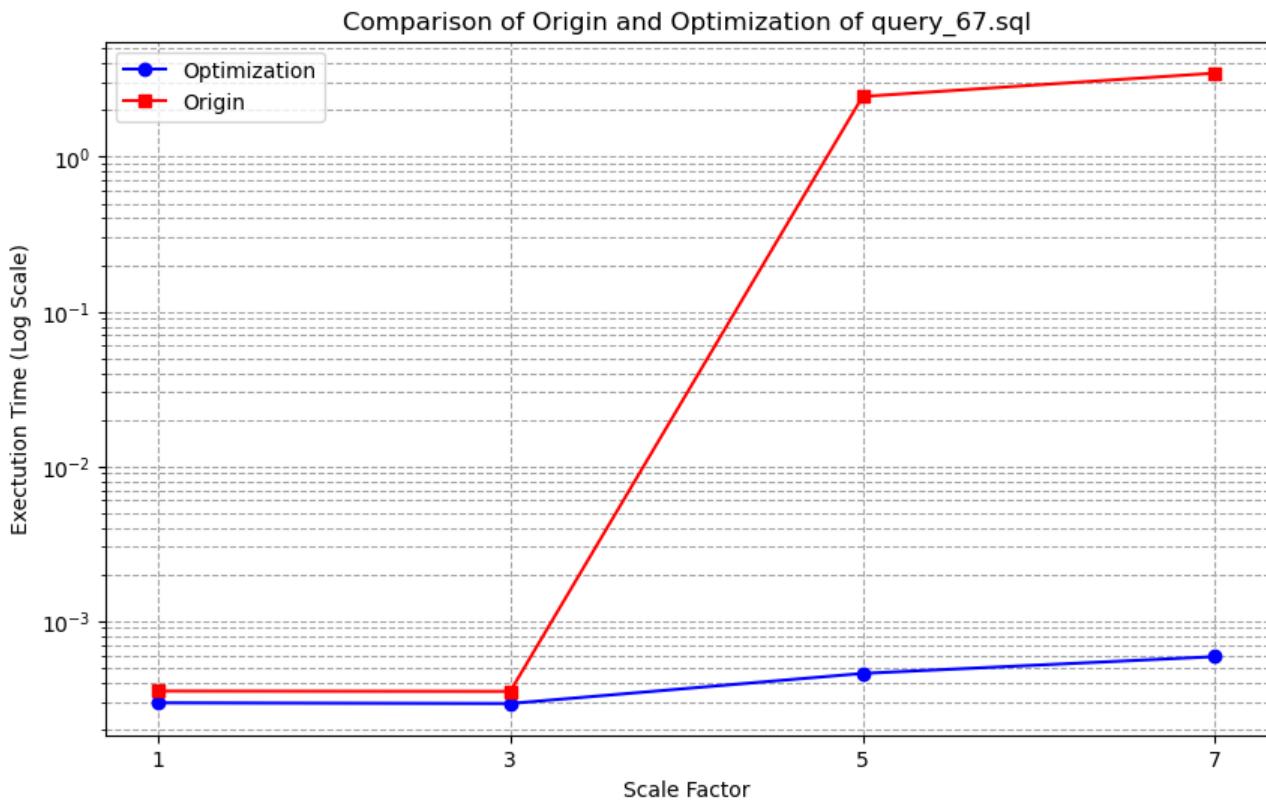


Figure 4.35: Optimization/Comparison of Origin and Optimization of query 67

### Query 23b

Query 23 was divided into two part because of the question it has to answer. This query contains multiple, related iterations: has the question *Find frequently sold items that were sold more than 4 times on any day during four consecutive years through the store sales channel. Compute the maximum store sales made by any given customer in a period of four consecutive years (same as above). Compute the best store customers that are in the 5th percentile of sales. Finally, compute the total web and catalog sales in a particular month made by our best store customers buying our most frequent store items.*

Qualification Substitution Parameters:

- MONTH.01 = 2
- YEAR.01 = 2000
- TOPPERCENT=50

Original Query for 23b

```
1 with frequent_ss_items as
2 (select substr(i_item_desc,1,30) itemdesc,i_item_sk item_sk,d_date solddate,
3   count(*) cnt
4   from store_sales
5   ,date_dim
6   ,item
7   where ss_sold_date_sk = d_date_sk
8   and ss_item_sk = i_item_sk
9   and d_year in (2000,2000+1,2000+2,2000+3)
10  group by substr(i_item_desc,1,30),i_item_sk,d_date
11  having count(*) >4),
12 max_store_sales as
13 (select max(csales) tpcds_cmax
14  from (select c_customer_sk,sum(ss_quantity*ss_sales_price) csales
15    from store_sales
16    ,customer
17    ,date_dim
18   where ss_customer_sk = c_customer_sk
19   and ss_sold_date_sk = d_date_sk
20   and d_year in (2000,2000+1,2000+2,2000+3)
21   group by c_customer_sk)),
22 best_ss_customer as
23 (select c_customer_sk,sum(ss_quantity*ss_sales_price) ssales
24  from store_sales
25  ,customer
26 where ss_customer_sk = c_customer_sk
27 group by c_customer_sk
28 having sum(ss_quantity*ss_sales_price) > (50/100.0) * (select
29 *
30 from
31 max_store_sales))
32 select sum(sales)
33 from (select cs_quantity*cs_list_price sales
34   from catalog_sales
35   ,date_dim
36   where d_year = 2000
37   and d_moy = 2
38   and cs_sold_date_sk = d_date_sk
39   and cs_item_sk in (select item_sk from frequent_ss_items)
40   and cs_bill_customer_sk in (select c_customer_sk from
```

```

    best_ss_customer)

40   union all

41     select ws_quantity*ws_list_price sales
42
43       from web_sales
44
45         ,date_dim
46
47       where d_year = 2000
48
49       and d_moy = 2
50
51       and ws_sold_date_sk = d_date_sk
52
53       and ws_item_sk in (select item_sk from frequent_ss_items)
54
55       and ws_bill_customer_sk in (select c_customer_sk from
56
57         best_ss_customer))
58
59 limit 100;

```

Listing 4.5: Query 23b.

The query was optimized by the use a TEMPORARY table instead of using the Common Table Expressions and then index was added on the temporary table, as seen below:

```

1 CREATE TEMPORARY TABLE frequent_ss_items AS
2
3   SELECT
4
5     SUBSTR(i_item_desc, 1, 30) AS itemdesc,
6
7     i_item_sk AS item_sk,
8
9     d_date AS solddate,
10
11    COUNT(*) AS cnt
12
13  FROM
14
15    store_sales
16
17  JOIN
18
19    date_dim ON ss_sold_date_sk = d_date_sk
20
21  JOIN
22
23    item ON ss_item_sk = i_item_sk
24
25  WHERE
26
27    d_year IN (2000, 2001, 2002, 2003)
28
29  GROUP BY
30
31    SUBSTR(i_item_desc, 1, 30), i_item_sk, d_date
32
33  HAVING COUNT(*) > 4;
34
35
36  -- Create an index on the temporary table
37
38  CREATE INDEX idx_frequent_items ON frequent_ss_items (item_sk);
39
40
41  -- Create a temporary table for maximum store sales
42
43  CREATE TEMPORARY TABLE max_store_sales AS

```

```
24 SELECT MAX(csales) AS tpcds_cmax
25 FROM (
26     SELECT
27         c_customer_sk ,
28         SUM(ss_quantity * ss_sales_price) AS csales
29     FROM
30         store_sales
31     JOIN
32         customer ON ss_customer_sk = c_customer_sk
33     JOIN
34         date_dim ON ss_sold_date_sk = d_date_sk
35     WHERE
36         d_year IN (2000, 2001, 2002, 2003)
37     GROUP BY
38         c_customer_sk
39 );
40
41 -- Create a temporary table for the best store sales customers
42 CREATE TEMPORARY TABLE best_ss_customer AS
43 SELECT
44     c_customer_sk ,
45     SUM(ss_quantity * ss_sales_price) AS ssales
46 FROM
47     store_sales
48 JOIN
49     customer ON ss_customer_sk = c_customer_sk
50 GROUP BY
51     c_customer_sk
52 HAVING
53     SUM(ss_quantity * ss_sales_price) > (50/100.0) * (SELECT tpcds_cmax FROM
54     max_store_sales);
55
56 -- Finally, run your main query using the temporary tables
57 SELECT
58     c_last_name ,
59     c_first_name ,
60     sales
61 FROM (
62     SELECT
63         c_last_name ,
```

```
63      c_first_name ,
64      SUM(cs_quantity * cs_list_price) AS sales
65  FROM
66    catalog_sales
67  JOIN
68    customer ON cs_bill_customer_sk = c_customer_sk
69  JOIN
70    date_dim ON cs_sold_date_sk = d_date_sk
71 WHERE
72   d_year = 2000
73 AND
74   d_moy = 2
75 AND
76   cs_item_sk IN (SELECT item_sk FROM frequent_ss_items)
77 AND
78   cs_bill_customer_sk IN (SELECT c_customer_sk FROM best_ss_customer)
79 AND
80   cs_bill_customer_sk = c_customer_sk
81 GROUP BY
82   c_last_name , c_first_name
83 UNION ALL
84 SELECT
85   c_last_name ,
86   c_first_name ,
87   SUM(ws_quantity * ws_list_price) AS sales
88  FROM
89    web_sales
90  JOIN
91    customer ON ws_bill_customer_sk = c_customer_sk
92  JOIN
93    date_dim ON ws_sold_date_sk = d_date_sk
94 WHERE
95   d_year = 2000
96 AND
97   d_moy = 2
98 AND
99   ws_item_sk IN (SELECT item_sk FROM frequent_ss_items)
100 AND
101   ws_bill_customer_sk IN (SELECT c_customer_sk FROM best_ss_customer)
102 AND
```

```

103     ws_bill_customer_sk = c_customer_sk
104
105     GROUP BY
106         c_last_name, c_first_name
107 ) AS final_results
108 ORDER BY
109     c_last_name, c_first_name, sales
110 LIMIT 100;

```

Listing 4.6: Optimized Query 23b.

The result is visualized below:

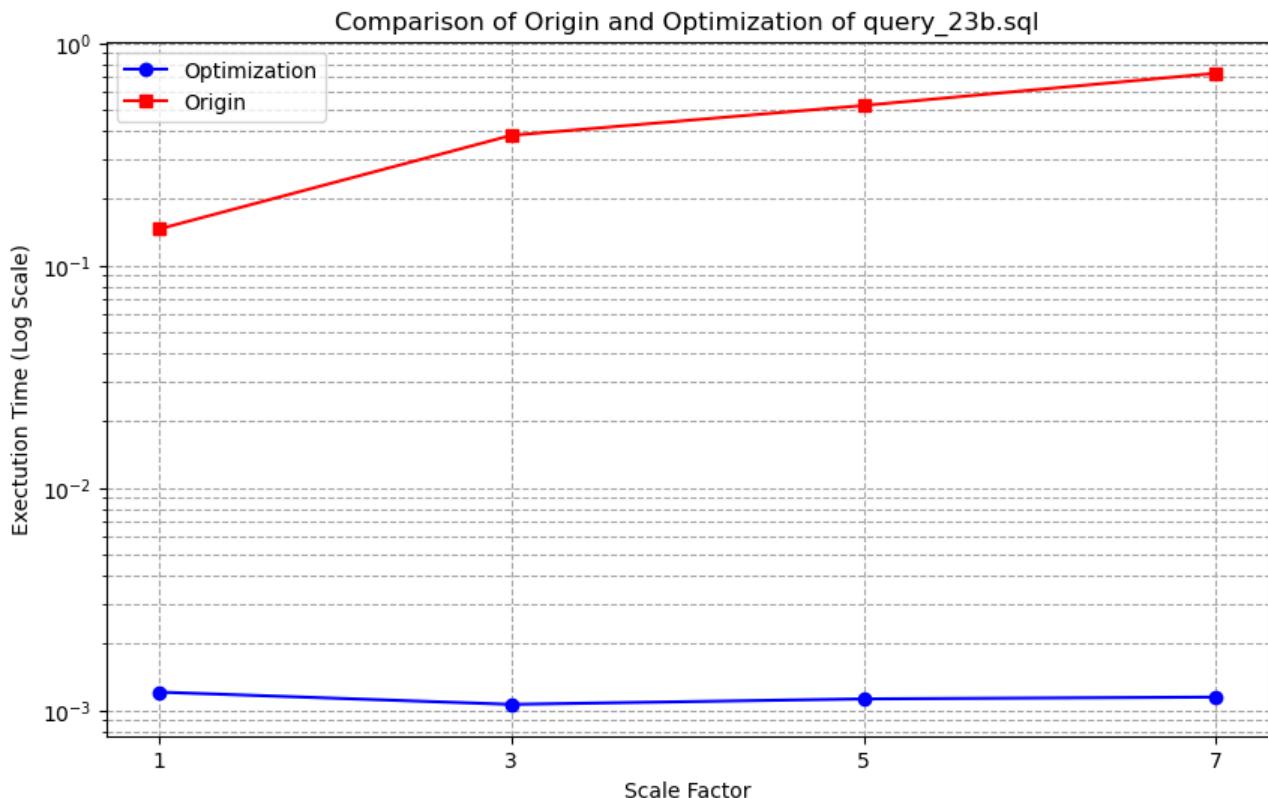


Figure 4.36: Optimization/Comparison of Origin and Optimization of query 23b

## Query 4

Query 4 answer the question, *Find customers who spend more money via catalog than in stores. Identify preferred customers and their country of origin.*

Qualification Substitution Parameters:

- YEAR.01=2001

- • SELECTONE.01=t\_s\_secyear.customer\_preferred\_cust\_flag

The Original Query is :

```

1      with year_total as (
2 select c_customer_id customer_id
3       ,c_first_name customer_first_name
4       ,c_last_name customer_last_name
5       ,c_preferred_cust_flag customer_preferred_cust_flag
6       ,c_birth_country customer_birth_country
7       ,c_login customer_login
8       ,c_email_address customer_email_address
9       ,d_year dyear
10      ,sum(((ss_ext_list_price-ss_ext_wholesale_cost-ss_ext_discount_amt) +
11             ss_ext_sales_price)/2) year_total
12      , 's' sale_type
13
14 from customer
15       ,store_sales
16       ,date_dim
17 where c_customer_sk = ss_customer_sk
18   and ss_sold_date_sk = d_date_sk
19 group by c_customer_id
20       ,c_first_name
21       ,c_last_name
22       ,c_preferred_cust_flag
23       ,c_birth_country
24       ,c_login
25       ,c_email_address
26       ,d_year
27
28 union all
29
30 select c_customer_id customer_id
31       ,c_first_name customer_first_name
32       ,c_last_name customer_last_name
33       ,c_preferred_cust_flag customer_preferred_cust_flag
34       ,c_birth_country customer_birth_country
35       ,c_login customer_login
36       ,c_email_address customer_email_address
37       ,d_year dyear
38      ,sum(((cs_ext_list_price-cs_ext_wholesale_cost-cs_ext_discount_amt) +
39             cs_ext_sales_price)/2) ) year_total
40      , 'c' sale_type

```

```
36 from customer
37     ,catalog_sales
38     ,date_dim
39 where c_customer_sk = cs_bill_customer_sk
40     and cs_sold_date_sk = d_date_sk
41 group by c_customer_id
42         ,c_first_name
43         ,c_last_name
44         ,c_preferred_cust_flag
45         ,c_birth_country
46         ,c_login
47         ,c_email_address
48         ,d_year
49 union all
50 select c_customer_id customer_id
51     ,c_first_name customer_first_name
52     ,c_last_name customer_last_name
53     ,c_preferred_cust_flag customer_preferred_cust_flag
54     ,c_birth_country customer_birth_country
55     ,c_login customer_login
56     ,c_email_address customer_email_address
57     ,d_year dyear
58     ,sum((((ws_ext_list_price-ws_ext_wholesale_cost-ws_ext_discount_amt) +
59 ws_ext_sales_price)/2) ) year_total
60     , 'W' sale_type
61 from customer
62     ,web_sales
63     ,date_dim
64 where c_customer_sk = ws_bill_customer_sk
65     and ws_sold_date_sk = d_date_sk
66 group by c_customer_id
67         ,c_first_name
68         ,c_last_name
69         ,c_preferred_cust_flag
70         ,c_birth_country
71         ,c_login
72         ,c_email_address
73         ,d_year
74     )
75 select
```

```
75          t_s_secyear.customer_id  
76          ,t_s_secyear.customer_first_name  
77          ,t_s_secyear.customer_last_name  
78          ,t_s_secyear.customer_preferred_cust_flag  
79 from year_total t_s_firstyear  
80     ,year_total t_s_secyear  
81     ,year_total t_c_firstyear  
82     ,year_total t_c_secyear  
83     ,year_total t_w_firstyear  
84     ,year_total t_w_secyear  
85 where t_s_secyear.customer_id = t_s_firstyear.customer_id  
86   and t_s_firstyear.customer_id = t_c_secyear.customer_id  
87   and t_s_firstyear.customer_id = t_c_firstyear.customer_id  
88   and t_s_firstyear.customer_id = t_w_firstyear.customer_id  
89   and t_s_firstyear.customer_id = t_w_secyear.customer_id  
90   and t_s_firstyear.sale_type = 's'  
91   and t_c_firstyear.sale_type = 'c'  
92   and t_w_firstyear.sale_type = 'w'  
93   and t_s_secyear.sale_type = 's'  
94   and t_c_secyear.sale_type = 'c'  
95   and t_w_secyear.sale_type = 'w'  
96   and t_s_firstyear.dyear = 2001  
97   and t_s_secyear.dyear = 2001+1  
98   and t_c_firstyear.dyear = 2001  
99   and t_c_secyear.dyear = 2001+1  
100  and t_w_firstyear.dyear = 2001  
101  and t_w_secyear.dyear = 2001+1  
102  and t_s_firstyear.year_total > 0  
103  and t_c_firstyear.year_total > 0  
104  and t_w_firstyear.year_total > 0  
105  and case when t_c_firstyear.year_total > 0 then t_c_secyear.year_total /  
106    t_c_firstyear.year_total else null end  
107      > case when t_s_firstyear.year_total > 0 then t_s_secyear.  
108        year_total / t_s_firstyear.year_total else null end  
109  and case when t_c_firstyear.year_total > 0 then t_c_secyear.year_total /  
110    t_c_firstyear.year_total else null end  
111      > case when t_w_firstyear.year_total > 0 then t_w_secyear.  
112        year_total / t_w_firstyear.year_total else null end  
113 order by t_s_secyear.customer_id  
114           ,t_s_secyear.customer_first_name
```

```

111      ,t_s_secyear.customer_last_name
112      ,t_s_secyear.customer_preferred_cust_flag
113 limit 100;

```

Listing 4.7: Query 4.

It was Optimized by the use of Indexing as shown below:

```

1 CREATE INDEX idx_customer_c_customer_sk ON customer (c_customer_sk);
2 CREATE INDEX idx_customer_c_customer_id ON customer (c_customer_id);
3 CREATE INDEX idx_store_sales_ss_customer_sk ON store_sales (ss_customer_sk);
4 CREATE INDEX idx_store_sales_ss_sold_date_sk ON store_sales (ss_sold_date_sk
   );
5 CREATE INDEX idx_date_dim_d_date_sk ON date_dim (d_date_sk);
6 CREATE INDEX idx_catalog_sales_cs_bill_customer_sk ON catalog_sales (
   cs_bill_customer_sk);
7 CREATE INDEX idx_catalog_sales_cs_sold_date_sk ON catalog_sales (
   cs_sold_date_sk);
8 CREATE INDEX idx_web_sales_ws_bill_customer_sk ON web_sales (
   ws_bill_customer_sk);
9 CREATE INDEX idx_web_sales_ws_sold_date_sk ON web_sales (ws_sold_date_sk);

```

Listing 4.8: Optimized Query 14.

The resulting visual is shown below:

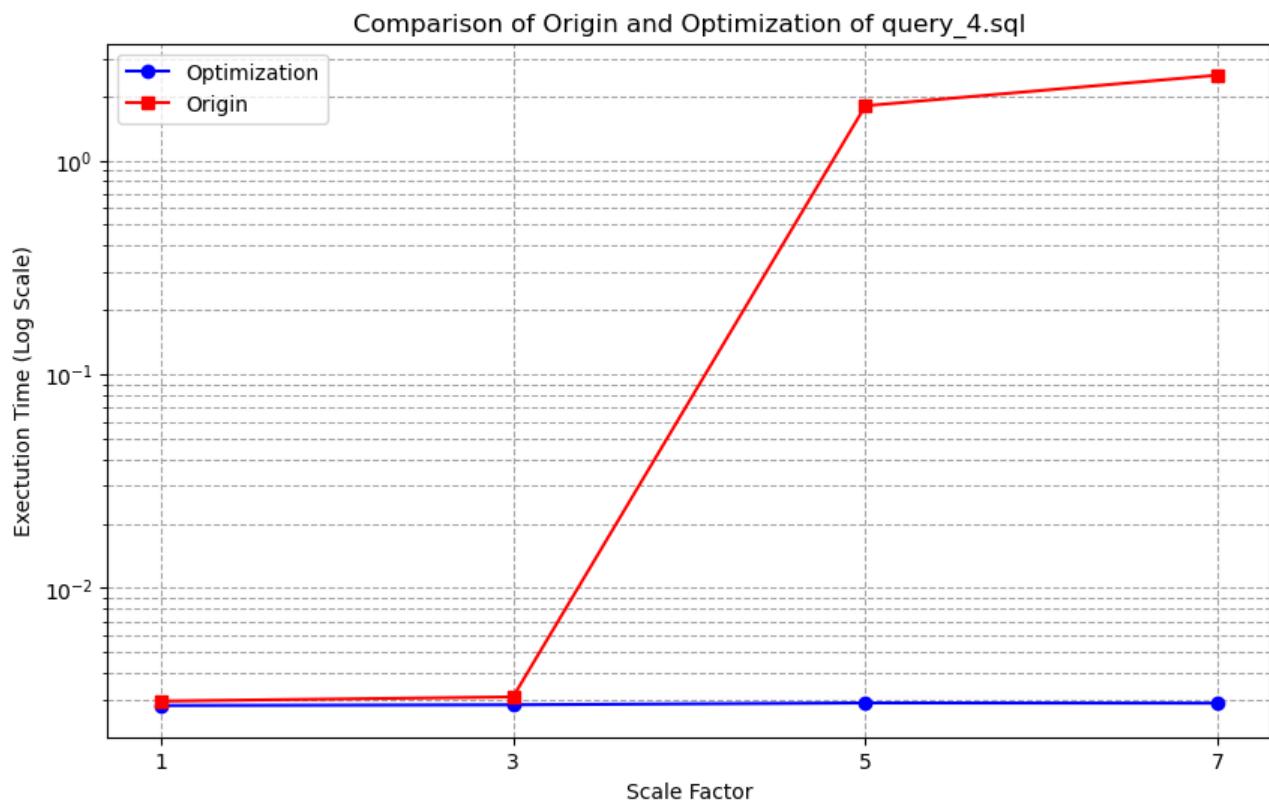


Figure 4.37: Optimization/Comparison of Origin and Optimization of query 4

The other optimized queries are in Appendix A section A

## Conclusion

### 5.1 Conclusion

In this project, We have systematically learned the basic theories and concepts of TPC-DS benchmark testing and implemented it on DuckDB, which is an open-source and particularly fast database. We finished the data and query generation. We conducted experiments for four scale factors on eight tests: load test, power test, throughput test 1,throughput test 2, data maintenance test 1,data maintenance test 2, data accessibility test, scalability test. It is evident that the execution time grows linearly with the size of data. We also made performance optimizations on some queries, which reduced their execution runtime. Then we presented and analyzed the results of our experiments in figures and charts with data visualization.

All the results are almost in line with our intuition about the size of the data, i.e., the larger the scale factor, the more resources are needed on the system. Runtimes vary from query to query, depending on the query complexity and the backend processes of the system. Also, we found that the PC equipment has some impact on the runtime, as well. If you're an analyst with a recent Macbook Pro and a small data warehouse, the laptop you use to write SQL queries might be faster than the data warehouse you run them on [Fra23].

It's worth noting that while TPC-DS is a comprehensive benchmark, no benchmark can capture all the nuances of every real-world workload. Therefore, while TPC-DS can provide valuable insights, it should be used in conjunction with other evaluation methodologies when making decisions about system architectures or purchasing decisions.

## References

- [TPC21] (TPC), Transaction Processing Performance Council (2021). “TPC Benchmark DS Standard Specification”. In.
- [Cod23] Code, Visual Studio (2023). “Visual Studio Code - Code Editing”. In: URL: <https://code.visualstudio.com/docs>.
- [Com23] Community, Dbeaver (2023). “Dbeaver Community Tool”. In: URL: <https://dbeaver.io/>.
- [Fra23] Fraser, George (2023). “Macbook Pro is impressive”. In: URL: <https://www.fivetran.com/blog/how-fast-is-duckdb-really>.
- [Ham23] Hamilton, Thomas (2023). “What is JMeter? Introduction & Uses”. In: URL: <https://www.guru99.com/introduction-to-jmeter.html>.
- [Koh+22] Kohn, André et al. (2022). “DuckDB-wasm: fast analytical processing for the web”. In: *Proceedings of the VLDB Endowment* 15.12, pp. 3574–3577.
- [Mah21] Mahmud, Md Hedayet (2021). “What is benchmarking and how can it be used to measure our current performance against an ideal or best practice model?” In: URL: <https://www.quora.com/What-is-benchmarking-and-how-can-it-be-used-to-measure-our-current-performance-against-an-ideal-or-best-practice-model>.
- [Pyt23] Python (2023). “General Python FAQ”. In: URL: <https://docs.python.org/3/faq/general.html/>.
- [RM19] Raasveldt, Mark and Hannes Mühleisen (2019). “DuckDB: an Embeddable Analytical Database”. In: *International Conference on Management of Data (SIGMOD '19)*. URL: <https://dl.acm.org/doi/abs/10.1145/3299869.3320212>.

- [Sam21] Samy Kabangu, John Ebden (2021). “Benchmarking Databases”. In.
- [VZ14] Vaisman, Alejandro and Esteban Zimányi (2014). “Data warehouse systems”. In: *Data-Centric Systems and Applications*.

# Appendices

## A Appendix A

### A.1 Optimized Queries

```
1 Optimize Query 11
2
3 -- Index for store_sales table
4 CREATE INDEX idx_ss_customer_sk ON store_sales(ss_customer_sk);
5 CREATE INDEX idx_ss_sold_date_sk ON store_sales(ss_sold_date_sk);
6
7 -- Index for customer table
8 CREATE INDEX idx_c_customer_sk ON customer(c_customer_sk);
9
10 -- Index for web_sales table
11 CREATE INDEX idx_ws_bill_customer_sk ON web_sales(ws_bill_customer_sk);
12 CREATE INDEX idx_ws_sold_date_sk ON web_sales(ws_sold_date_sk);
13
14 -- Index for date_dim table
15 CREATE INDEX idx_d_date_sk ON date_dim(d_date_sk);
16 CREATE INDEX idx_d_year ON date_dim(d_year);
17
18
19 with year_total as (
20 select c_customer_id customer_id
21     ,c_first_name customer_first_name
22     ,c_last_name customer_last_name
```

```
23      ,c_preferred_cust_flag customer_preferred_cust_flag
24      ,c_birth_country customer_birth_country
25      ,c_login customer_login
26      ,c_email_address customer_email_address
27      ,d_year dyear
28      ,sum(ss_ext_list_price-ss_ext_discount_amt) year_total
29      ,'S' sale_type
30
31 from customer
32      ,store_sales
33      ,date_dim
34 where c_customer_sk = ss_customer_sk
35 and ss_sold_date_sk = d_date_sk
36 group by c_customer_id
37      ,c_first_name
38      ,c_last_name
39      ,c_preferred_cust_flag
40      ,c_birth_country
41      ,c_login
42      ,c_email_address
43      ,d_year
44 union all
45 select c_customer_id customer_id
46      ,c_first_name customer_first_name
47      ,c_last_name customer_last_name
48      ,c_preferred_cust_flag customer_preferred_cust_flag
49      ,c_birth_country customer_birth_country
50      ,c_login customer_login
51      ,c_email_address customer_email_address
52      ,d_year dyear
53      ,sum(ws_ext_list_price-ws_ext_discount_amt) year_total
54      ,'W' sale_type
55
56 from customer
57      ,web_sales
58      ,date_dim
59 where c_customer_sk = ws_bill_customer_sk
60 and ws_sold_date_sk = d_date_sk
61 group by c_customer_id
62      ,c_first_name
63      ,c_last_name
64      ,c_preferred_cust_flag
```

```

63     ,c_birth_country
64     ,c_login
65     ,c_email_address
66     ,d_year
67   )
68
69 select
70   t_s_secyear.customer_id
71   ,t_s_secyear.customer_first_name
72   ,t_s_secyear.customer_last_name
73   ,t_s_secyear.customer_preferred_cust_flag
74 from year_total t_s_firstyear
75   ,year_total t_s_secyear
76   ,year_total t_w_firstyear
77   ,year_total t_w_secyear
78 where t_s_secyear.customer_id = t_s_firstyear.customer_id
79   and t_s_firstyear.customer_id = t_w_secyear.customer_id
80   and t_s_firstyear.customer_id = t_w_firstyear.customer_id
81   and t_s_firstyear.sale_type = 's'
82   and t_w_firstyear.sale_type = 'w'
83   and t_s_secyear.sale_type = 's'
84   and t_w_secyear.sale_type = 'w'
85   and t_s_firstyear.dyear = 2001
86   and t_s_secyear.dyear = 2001+1
87   and t_w_firstyear.dyear = 2001
88   and t_w_secyear.dyear = 2001+1
89   and t_s_firstyear.year_total > 0
90   and t_w_firstyear.year_total > 0
91   and case when t_w_firstyear.year_total > 0 then t_w_secyear.
92   year_total / t_w_firstyear.year_total else 0.0 end
93   > case when t_s_firstyear.year_total > 0 then t_s_secyear.
94   year_total / t_s_firstyear.year_total else 0.0 end
95 order by t_s_secyear.customer_id
96   ,t_s_secyear.customer_first_name
97   ,t_s_secyear.customer_last_name
98   ,t_s_secyear.customer_preferred_cust_flag
99 limit 100;
100
101 -----

```

Listing 1: Optimized Query 11.

```
1  
2 Optimized Query 14  
3  
4 CREATE INDEX idx_item_i_item_sk ON item (i_item_sk);  
5 CREATE INDEX idx_date_dim_d_date_sk ON date_dim (d_date_sk);  
6 CREATE INDEX idx_store_sales_ss_item_sk ON store_sales (ss_item_sk);  
7 CREATE INDEX idx_catalog_sales_cs_item_sk ON catalog_sales (cs_item_sk);  
8 CREATE INDEX idx_web_sales_ws_item_sk ON web_sales (ws_item_sk);  
9  
10  
11 CREATE INDEX idx_item_i_brand_id ON item (i_brand_id);  
12 CREATE INDEX idx_item_i_class_id ON item (i_class_id);  
13 CREATE INDEX idx_item_i_category_id ON item (i_category_id);  
14  
15  
16 with cross_items as  
17 (select i_item_sk ss_item_sk  
18 from item,  
19 (select iss.i_brand_id brand_id  
20 ,iss.i_class_id class_id  
21 ,iss.i_category_id category_id  
22 from store_sales  
23 ,item iss  
24 ,date_dim d1  
25 where ss_item_sk = iss.i_item_sk  
26 and ss_sold_date_sk = d1.d_date_sk  
27 and d1.d_year between 1999 AND 1999 + 2  
28 intersect  
29 select ics.i_brand_id  
30 ,ics.i_class_id  
31 ,ics.i_category_id  
32 from catalog_sales  
33 ,item ics  
34 ,date_dim d2  
35 where cs_item_sk = ics.i_item_sk  
36 and cs_sold_date_sk = d2.d_date_sk  
37 and d2.d_year between 1999 AND 1999 + 2  
38 intersect  
39 select iws.i_brand_id  
40 ,iws.i_class_id
```

```
41     ,iws.i_category_id
42 from web_sales
43     ,item iws
44     ,date_dim d3
45 where ws_item_sk = iws.i_item_sk
46   and ws_sold_date_sk = d3.d_date_sk
47   and d3.d_year between 1999 AND 1999 + 2) x
48 where i_brand_id = brand_id
49   and i_class_id = class_id
50   and i_category_id = category_id
51 ),
52 avg_sales as
53 (select avg(quantity*list_price) average_sales
54  from (select ss_quantity quantity
55            ,ss_list_price list_price
56       from store_sales
57            ,date_dim
58      where ss_sold_date_sk = d_date_sk
59        and d_year between 1999 and 1999 + 2
60    union all
61      select cs_quantity quantity
62            ,cs_list_price list_price
63       from catalog_sales
64            ,date_dim
65      where cs_sold_date_sk = d_date_sk
66        and d_year between 1999 and 1999 + 2
67    union all
68      select ws_quantity quantity
69            ,ws_list_price list_price
70       from web_sales
71            ,date_dim
72      where ws_sold_date_sk = d_date_sk
73        and d_year between 1999 and 1999 + 2) x)
74 select this_year.channel ty_channel
75           ,this_year.i_brand_id ty_brand
76           ,this_year.i_class_id ty_class
77           ,this_year.i_category_id ty_category
78           ,this_year.sales ty_sales
79           ,this_year.number_sales ty_number_sales
80           ,last_year.channel ly_channel
```

```
81             ,last_year.i_brand_id ly_brand
82             ,last_year.i_class_id ly_class
83             ,last_year.i_category_id ly_category
84             ,last_year.sales ly_sales
85             ,last_year.number_sales ly_number_sales
86
87 from
88 (select 'store' channel, i_brand_id,i_class_id,i_category_id
89         ,sum(ss_quantity*ss_list_price) sales, count(*) number_sales
90 from store_sales
91         ,item
92         ,date_dim
93 where ss_item_sk in (select ss_item_sk from cross_items)
94     and ss_item_sk = i_item_sk
95     and ss_sold_date_sk = d_date_sk
96     and d_week_seq = (select d_week_seq
97                         from date_dim
98                         where d_year = 1999 + 1
99                             and d_moy = 12
100                            and d_dom = 3)
101 group by i_brand_id,i_class_id,i_category_id
102 having sum(ss_quantity*ss_list_price) > (select average_sales from avg_sales
103 )) this_year,
104 (select 'store' channel, i_brand_id,i_class_id
105         ,i_category_id, sum(ss_quantity*ss_list_price) sales, count(*)
106         number_sales
107 from store_sales
108         ,item
109         ,date_dim
110 where ss_item_sk in (select ss_item_sk from cross_items)
111     and ss_item_sk = i_item_sk
112     and ss_sold_date_sk = d_date_sk
113     and d_week_seq = (select d_week_seq
114                         from date_dim
115                         where d_year = 1999
116                             and d_moy = 12
117                             and d_dom = 3)
118 group by i_brand_id,i_class_id,i_category_id
119 having sum(ss_quantity*ss_list_price) > (select average_sales from avg_sales
120 )) last_year
121 where this_year.i_brand_id= last_year.i_brand_id
```

```

118     and this_year.i_class_id = last_year.i_class_id
119     and this_year.i_category_id = last_year.i_category_id
120 order by this_year.channel, this_year.i_brand_id, this_year.i_class_id,
121         this_year.i_category_id
122 limit 100;
123 -----

```

Listing 2: Optimized Query 14.

```

1 --Optimized Query 14b
2
3 CREATE INDEX idx_item_i_item_sk ON item (i_item_sk);
4 CREATE INDEX idx_date_dim_d_date_sk ON date_dim (d_date_sk);
5 CREATE INDEX idx_store_sales_ss_item_sk ON store_sales (ss_item_sk);
6 CREATE INDEX idx_catalog_sales_cs_item_sk ON catalog_sales (cs_item_sk);
7 CREATE INDEX idx_web_sales_ws_item_sk ON web_sales (ws_item_sk);
8
9
10 CREATE INDEX idx_item_i_brand_id ON item (i_brand_id);
11 CREATE INDEX idx_item_i_class_id ON item (i_class_id);
12 CREATE INDEX idx_item_i_category_id ON item (i_category_id);
13
14
15
16 with cross_items as
17 (select i_item_sk ss_item_sk
18 from item,
19 (select iss.i_brand_id brand_id
20 ,iss.i_class_id class_id
21 ,iss.i_category_id category_id
22 from store_sales
23 ,item iss
24 ,date_dim d1
25 where ss_item_sk = iss.i_item_sk
26 and ss_sold_date_sk = d1.d_date_sk
27 and d1.d_year between 1999 AND 1999 + 2
28 intersect
29 select ics.i_brand_id
30 ,ics.i_class_id

```

```
31     ,ics.i_category_id
32 from catalog_sales
33     ,item ics
34     ,date_dim d2
35 where cs_item_sk = ics.i_item_sk
36   and cs_sold_date_sk = d2.d_date_sk
37   and d2.d_year between 1999 AND 1999 + 2
38 intersect
39 select iws.i_brand_id
40       ,iws.i_class_id
41       ,iws.i_category_id
42 from web_sales
43       ,item iws
44       ,date_dim d3
45 where ws_item_sk = iws.i_item_sk
46   and ws_sold_date_sk = d3.d_date_sk
47   and d3.d_year between 1999 AND 1999 + 2)
48 where i_brand_id = brand_id
49       and i_class_id = class_id
50       and i_category_id = category_id
51 ),
52 avg_sales as
53 (select avg(quantity*list_price) average_sales
54   from (select ss_quantity quantity
55             ,ss_list_price list_price
56         from store_sales
57             ,date_dim
58        where ss_sold_date_sk = d_date_sk
59        and d_year between 1999 and 1999 + 2
60  union all
61      select cs_quantity quantity
62             ,cs_list_price list_price
63        from catalog_sales
64             ,date_dim
65        where cs_sold_date_sk = d_date_sk
66        and d_year between 1999 and 1999 + 2
67  union all
68      select ws_quantity quantity
69             ,ws_list_price list_price
70        from web_sales
```

```
71      ,date_dim
72      where ws_sold_date_sk = d_date_sk
73          and d_year between 1999 and 1999 + 2) x)
74  select  channel, i_brand_id,i_class_id,i_category_id,sum(sales), sum(
75      number_sales)
76 from(
77     select 'store' channel, i_brand_id,i_class_id
78         ,i_category_id,sum(ss_quantity*ss_list_price) sales
79         , count(*) number_sales
80   from store_sales
81     ,item
82     ,date_dim
83   where ss_item_sk in (select ss_item_sk from cross_items)
84       and ss_item_sk = i_item_sk
85       and ss_sold_date_sk = d_date_sk
86       and d_year = 1999+2
87       and d_moy = 11
88   group by i_brand_id,i_class_id,i_category_id
89   having sum(ss_quantity*ss_list_price) > (select average_sales from
90 avg_sales)
91   union all
92     select 'catalog' channel, i_brand_id,i_class_id,i_category_id , sum(
93      cs_quantity*cs_list_price) sales, count(*) number_sales
94   from catalog_sales
95     ,item
96     ,date_dim
97   where cs_item_sk in (select ss_item_sk from cross_items)
98       and cs_item_sk = i_item_sk
99       and cs_sold_date_sk = d_date_sk
100      and d_year = 1999+2
101      and d_moy = 11
102  group by i_brand_id,i_class_id,i_category_id
103  having sum(cs_quantity*cs_list_price) > (select average_sales from
104 avg_sales)
105  union all
106    select 'web' channel, i_brand_id,i_class_id,i_category_id , sum(
107      ws_quantity*ws_list_price) sales , count(*) number_sales
108   from web_sales
109     ,item
110     ,date_dim
```

```

106     where ws_item_sk in (select ss_item_sk from cross_items)
107         and ws_item_sk = i_item_sk
108         and ws_sold_date_sk = d_date_sk
109         and d_year = 1999+2
110         and d_moy = 11
111     group by i_brand_id,i_class_id,i_category_id
112     having sum(ws_quantity*ws_list_price) > (select average_sales from
113         avg_sales)
114 ) y
115 group by rollup (channel, i_brand_id,i_class_id,i_category_id)
116 --order by channel,i_brand_id,i_class_id,i_category_id
117 ORDER BY channel NULLS FIRST,
118         i_brand_id NULLS FIRST,
119         i_class_id NULLS FIRST,
120         i_category_id NULLS FIRST
121 limit 100;
122 -----

```

Listing 3: Optimized Query 14b.

```

1 CREATE TEMP TABLE tmp_avg_qoh AS
2 SELECT i_product_name, i_brand, i_class, i_category, AVG(
3     inv_quantity_on_hand) AS qoh
4 FROM inventory
5 JOIN date_dim ON inv_date_sk = d_date_sk
6 JOIN item ON inv_item_sk = i_item_sk
7 WHERE d_month_seq BETWEEN 1200 AND 1200 + 11
8 GROUP BY i_product_name, i_brand, i_class, i_category;
9
10 SELECT i_product_name, i_brand, i_class, i_category, qoh
11 FROM tmp_avg_qoh
12 ORDER BY roqoh, i_product_name, i_brand, i_class, i_category
13 LIMIT 100;
14
15 OR THIS
16
17 CREATE INDEX idx_inv_date_sk ON inventory (inv_date_sk);
18 CREATE INDEX idx_inv_item_sk ON inventory (inv_item_sk);

```

```

19 select    i_product_name
20          ,i_brand
21          ,i_class
22          ,i_category
23          ,avg(inv_quantity_on_hand) qoh
24 from inventory
25          ,date_dim
26          ,item
27 where inv_date_sk=d_date_sk
28         and inv_item_sk=i_item_sk
29         and d_month_seq between 1200 and 1200 + 11
30 group by rollup(i_product_name
31                  ,i_brand
32                  ,i_class
33                  ,i_category)
34 order by qoh, i_product_name, i_brand, i_class, i_category
35 limit 100;

```

Listing 4: Optimized Query 22.

```

1 CREATE TEMPORARY TABLE frequent_ss_items AS
2 SELECT
3     SUBSTR(i_item_desc, 1, 30) AS itemdesc,
4     i_item_sk AS item_sk,
5     d_date AS solddate,
6     COUNT(*) AS cnt
7 FROM
8     store_sales
9 JOIN
10    date_dim ON ss_sold_date_sk = d_date_sk
11 JOIN
12    item ON ss_item_sk = i_item_sk
13 WHERE
14     d_year IN (2000, 2001, 2002, 2003)
15 GROUP BY
16     SUBSTR(i_item_desc, 1, 30), i_item_sk, d_date
17 HAVING COUNT(*) > 4;
18
19 -- Create an index on the temporary table
20 CREATE INDEX idx_frequent_items ON frequent_ss_items (item_sk);
21

```

```
22 -- Create a temporary table for maximum store sales
23 CREATE TEMPORARY TABLE max_store_sales AS
24 SELECT MAX(csales) AS tpcds_cmax
25 FROM (
26     SELECT
27         c_customer_sk ,
28         SUM(ss_quantity * ss_sales_price) AS csales
29     FROM
30         store_sales
31     JOIN
32         customer ON ss_customer_sk = c_customer_sk
33     JOIN
34         date_dim ON ss_sold_date_sk = d_date_sk
35     WHERE
36         d_year IN (2000, 2001, 2002, 2003)
37     GROUP BY
38         c_customer_sk
39 );
40
41 -- Create a temporary table for the best store sales customers
42 CREATE TEMPORARY TABLE best_ss_customer AS
43 SELECT
44     c_customer_sk ,
45     SUM(ss_quantity * ss_sales_price) AS ssales
46 FROM
47     store_sales
48 JOIN
49     customer ON ss_customer_sk = c_customer_sk
50 GROUP BY
51     c_customer_sk
52 HAVING
53     SUM(ss_quantity * ss_sales_price) > (50/100.0) * (SELECT tpcds_cmax FROM
54     max_store_sales);
55
56 SELECT
57     c_last_name ,
58     c_first_name ,
59     sales
60 FROM (
61     SELECT
```

```
61      c_last_name ,
62      c_first_name ,
63      SUM(cs_quantity * cs_list_price) AS sales
64
65      FROM
66
67      catalog_sales
68
69      JOIN
70
71          customer ON cs_bill_customer_sk = c_customer_sk
72
73      JOIN
74
75          date_dim ON cs_sold_date_sk = d_date_sk
76
77      WHERE
78
79          d_year = 2000
80
81      AND
82
83          d_moy = 2
84
85      AND
86
87          cs_item_sk IN (SELECT item_sk FROM frequent_ss_items)
88
89      AND
90
91          cs_bill_customer_sk IN (SELECT c_customer_sk FROM best_ss_customer)
92
93      AND
94
95          cs_bill_customer_sk = c_customer_sk
96
97      GROUP BY
98
99          c_last_name , c_first_name
100
101     UNION ALL
102
103     SELECT
104
105         c_last_name ,
106
107         c_first_name ,
108
109         SUM(ws_quantity * ws_list_price) AS sales
110
111     FROM
112
113         web_sales
114
115     JOIN
116
117         customer ON ws_bill_customer_sk = c_customer_sk
118
119     JOIN
120
121         date_dim ON ws_sold_date_sk = d_date_sk
122
123     WHERE
124
125         d_year = 2000
126
127     AND
128
129         d_moy = 2
130
131     AND
132
133         ws_item_sk IN (SELECT item_sk FROM frequent_ss_items)
134
135     AND
136
137         ws_bill_customer_sk IN (SELECT c_customer_sk FROM best_ss_customer)
```

```

101    AND
102        ws_bill_customer_sk = c_customer_sk
103    GROUP BY
104        c_last_name, c_first_name
105 ) AS final_results
106 ORDER BY
107        c_last_name, c_first_name, sales
108 LIMIT 100;
109 -----

```

Listing 5: Optimized Query 23.

```

1 -- Optimized Query 47
2
3 -- Create a temporary table for CTE v1
4 CREATE TEMP TABLE temp_v1 AS
5 SELECT i_category, i_brand,
6       s_store_name, s_company_name,
7       d_year, d_moy,
8       SUM(ss_sales_price) sum_sales,
9       AVG(SUM(ss_sales_price)) OVER
10      (PARTITION BY i_category, i_brand,
11           s_store_name, s_company_name, d_year)
12      avg_monthly_sales,
13      RANK() OVER/
14      (PARTITION BY i_category, i_brand,
15           s_store_name, s_company_name
16           ORDER BY d_year, d_moy) rn
17 FROM item, store_sales, date_dim, store
18 WHERE ss_item_sk = i_item_sk AND
19       ss_sold_date_sk = d_date_sk AND
20       ss_store_sk = s_store_sk AND
21       (
22         d_year = 1999 or
23         (d_year = 1999 - 1 and d_moy = 12) or
24         (d_year = 1999 + 1 and d_moy = 1)
25       )
26 GROUP BY i_category, i_brand,
27       s_store_name, s_company_name,
28       d_year, d_moy;

```

```

29
30 -- Create a temporary table for CTE v2
31 CREATE TEMP TABLE temp_v2 AS
32 SELECT v1.i_category, v1.i_brand, v1.s_store_name, v1.s_company_name, v1.
33     s_store_name, v1.s_company_name
34     ,v1.d_year
35     ,v1.avg_monthly_sales
36     ,v1.sum_sales, v1_lag.sum_sales psum, v1_lead.sum_sales nsum,v1.
37     d_year, v1.d_moy
38 FROM temp_v1 v1
39 JOIN temp_v1 v1_lag ON v1.i_category = v1_lag.i_category and
40                     v1.i_brand = v1_lag.i_brand and
41                     v1.s_store_name = v1_lag.s_store_name and
42                     v1.s_company_name = v1_lag.s_company_name and
43                     v1.rn = v1_lag.rn + 1
44 JOIN temp_v1 v1_lead ON v1.i_category = v1_lead.i_category and
45                     v1.i_brand = v1_lead.i_brand and
46                     v1.s_store_name = v1_lead.s_store_name and
47                     v1.s_company_name = v1_lead.s_company_name and
48                     v1.rn = v1_lead.rn - 1;
49
50
51
52
53
54
55
56
57
58 -----

```

Listing 6: Optimized Query 47.

```

1 -- Optimized Query 57
2
3
4 -- Create a temporary table for CTE v1

```

```
5 CREATE TEMP TABLE temp_v11 AS
6 SELECT i_category, i_brand, cc_name, d_year, d_moy,
7     SUM(cs_sales_price) sum_sales,
8     AVG(SUM(cs_sales_price)) OVER
9         (PARTITION BY i_category, i_brand, cc_name, d_year)
10        avg_monthly_sales,
11        RANK() OVER
12            (PARTITION BY i_category, i_brand, cc_name
13             ORDER BY d_year, d_moy) rn
14 FROM item, catalog_sales, date_dim, call_center
15 WHERE cs_item_sk = i_item_sk AND
16     cs_sold_date_sk = d_date_sk AND
17     cc_call_center_sk = cs_call_center_sk AND
18 (
19     d_year = 1999 or
20     (d_year = 1999 - 1 and d_moy = 12) or
21     (d_year = 1999 + 1 and d_moy = 1)
22 )
23 GROUP BY i_category, i_brand, cc_name, d_year, d_moy;
24
25 --Create a temporary table for CTE v2
26 CREATE TEMP TABLE temp_v22 AS
27 SELECT v11.i_category, v11.i_brand, v11.cc_name
28     ,v11.d_year
29     ,v11.avg_monthly_sales
30     ,v11.sum_sales, v1_lag.sum_sales psum, v1_lead.sum_sales nsum,
31     v11.i_category, v11.i_brand, v11.cc_name, v11.d_year, v11.d_moy
32 FROM temp_v11 v11
33 JOIN temp_v11 v1_lag ON v11.i_category = v1_lag.i_category and
34                 v11.i_brand = v1_lag.i_brand and
35                 v11.cc_name = v1_lag.cc_name and
36                 v11.rn = v1_lag.rn + 1
37 JOIN temp_v11 v1_lead ON v11.i_category = v1_lead.i_category and
38                 v11.i_brand = v1_lead.i_brand and
39                 v11.cc_name = v1_lead.cc_name and
40                 v11.rn = v1_lead.rn - 1;
41
42 --Retrieve the final result
43 SELECT *
44 FROM temp_v22
```

```
45 WHERE d_year = 1999 and
46     avg_monthly_sales > 0 and
47     CASE WHEN avg_monthly_sales > 0 THEN ABS(sum_sales - avg_monthly_sales)
48         / avg_monthly_sales ELSE NULL END > 0.1
49 ORDER BY cc_name
50
51
```

**Listing 7:** Optimized Query 57.

```
1 --Optimized Query 64
2
3 Indexes for Join Conditions
4 CREATE INDEX idx_cs_item_sk ON catalog_sales (cs_item_sk);
5 CREATE INDEX idx_cr_item_sk ON catalog_returns (cr_item_sk);
6
7 CREATE INDEX idx_customer_join ON customer (c_customer_sk,
8     c_current_cdemo_sk, c_current_hdemo_sk, c_current_addr_sk);
9
10 --CREATE INDEX idx_ca_address_sk ON customer_address (ca_address_sk);
11
12 CREATE INDEX idx_promotion_promo_sk ON promotion (p_promo_sk);
13
14 CREATE INDEX idx_hd_income_band_sk ON household_demographics (
15     hd_income_band_sk);
16
17 CREATE INDEX idx_ib_income_band_sk ON income_band (ib_income_band_sk);
18
19 CREATE INDEX idx_item_join ON item (i_item_sk, i_color, i_current_price);
20
21 -- Indexes for Aggregation and Filtering
22
23
24
25 with cs_ui as
26 (select cs_item_sk
```

```
27      ,sum(cs_ext_list_price) as sale,sum(cr_refunded_cash+
28      cr_reversed_charge+cr_store_credit) as refund
29
30 from catalog_sales
31      ,catalog_returns
32 where cs_item_sk = cr_item_sk
33      and cs_order_number = cr_order_number
34 group by cs_item_sk
35 having sum(cs_ext_list_price)>2*sum(cr_refunded_cash+cr_reversed_charge+
36      cr_store_credit)),
37 cross_sales as
38 (select i_product_name product_name
39      ,i_item_sk item_sk
40      ,s_store_name store_name
41      ,s_zip store_zip
42      ,ad1.ca_street_number b_street_number
43      ,ad1.ca_street_name b_street_name
44      ,ad1.ca_city b_city
45      ,ad1.ca_zip b_zip
46      ,ad2.ca_street_number c_street_number
47      ,ad2.ca_street_name c_street_name
48      ,ad2.ca_city c_city
49      ,ad2.ca_zip c_zip
50      ,d1.d_year as syear
51      ,d2.d_year as fsyear
52      ,d3.d_year s2year
53      ,count(*) cnt
54      ,sum(ss_wholesale_cost) s1
55      ,sum(ss_list_price) s2
56      ,sum(ss_coupon_amt) s3
57
58 FROM store_sales
59      ,store_returns
60      ,cs_ui
61      ,date_dim d1
62      ,date_dim d2
63      ,date_dim d3
64      ,store
65      ,customer
66      ,customer_demographics cd1
67      ,customer_demographics cd2
68      ,promotion
```

```
65      ,household_demographics hd1
66      ,household_demographics hd2
67      ,customer_address ad1
68      ,customer_address ad2
69      ,income_band ib1
70      ,income_band ib2
71      ,item
72 WHERE ss_store_sk = s_store_sk AND
73      ss_sold_date_sk = d1.d_date_sk AND
74      ss_customer_sk = c_customer_sk AND
75      ss_cdemo_sk= cd1.cd_demo_sk AND
76      ss_hdemo_sk = hd1.hd_demo_sk AND
77      ss_addr_sk = ad1.ca_address_sk and
78      ss_item_sk = i_item_sk and
79      ss_item_sk = sr_item_sk and
80      ss_ticket_number = sr_ticket_number and
81      ss_item_sk = cs_ui.cs_item_sk and
82      c_current_cdemo_sk = cd2.cd_demo_sk AND
83      c_current_hdemo_sk = hd2.hd_demo_sk AND
84      c_current_addr_sk = ad2.ca_address_sk and
85      c_first_sales_date_sk = d2.d_date_sk and
86      c_first_shipto_date_sk = d3.d_date_sk and
87      ss_promo_sk = p_promo_sk and
88      hd1.hd_income_band_sk = ib1.ib_income_band_sk and
89      hd2.hd_income_band_sk = ib2.ib_income_band_sk and
90      cd1.cd_marital_status <> cd2.cd_marital_status and
91      i_color in ('purple','burlywood','indian','spring','floral','medium',
92 ) and
93      i_current_price between 64 and 64 + 10 and
94      i_current_price between 64 + 1 and 64 + 15
95 group by i_product_name
96      ,i_item_sk
97      ,s_store_name
98      ,s_zip
99      ,ad1.ca_street_number
100     ,ad1.ca_street_name
101     ,ad1.ca_city
102     ,ad1.ca_zip
103     ,ad2.ca_street_number
104     ,ad2.ca_street_name
```

```
104     ,ad2.ca_city
105     ,ad2.ca_zip
106     ,d1.d_year
107     ,d2.d_year
108     ,d3.d_year
109 )
110 select cs1.product_name
111   ,cs1.store_name
112   ,cs1.store_zip
113   ,cs1.b_street_number
114   ,cs1.b_street_name
115   ,cs1.b_city
116   ,cs1.b_zip
117   ,cs1.c_street_number
118   ,cs1.c_street_name
119   ,cs1.c_city
120   ,cs1.c_zip
121   ,cs1.syear
122   ,cs1.cnt
123   ,cs1.s1 as s11
124   ,cs1.s2 as s21
125   ,cs1.s3 as s31
126   ,cs2.s1 as s12
127   ,cs2.s2 as s22
128   ,cs2.s3 as s32
129   ,cs2.syear
130   ,cs2.cnt
131 from cross_sales cs1,cross_sales cs2
132 where cs1.item_sk=cs2.item_sk and
133       cs1.syear = 1999 and
134       cs2.syear = 1999 + 1 and
135       cs2.cnt <= cs1.cnt and
136       cs1.store_name = cs2.store_name and
137       cs1.store_zip = cs2.store_zip
138 order by cs1.product_name
139   ,cs1.store_name
140   ,cs2.cnt
141   ,cs1.s1
142   ,cs2.s1;
143
```

Listing 8: Optimized Query 64.